



Advanced Deep Learning 2025

GANs & GAN editing



Vicky Kalogeiton

Lecture 1: CSC_52087_EP

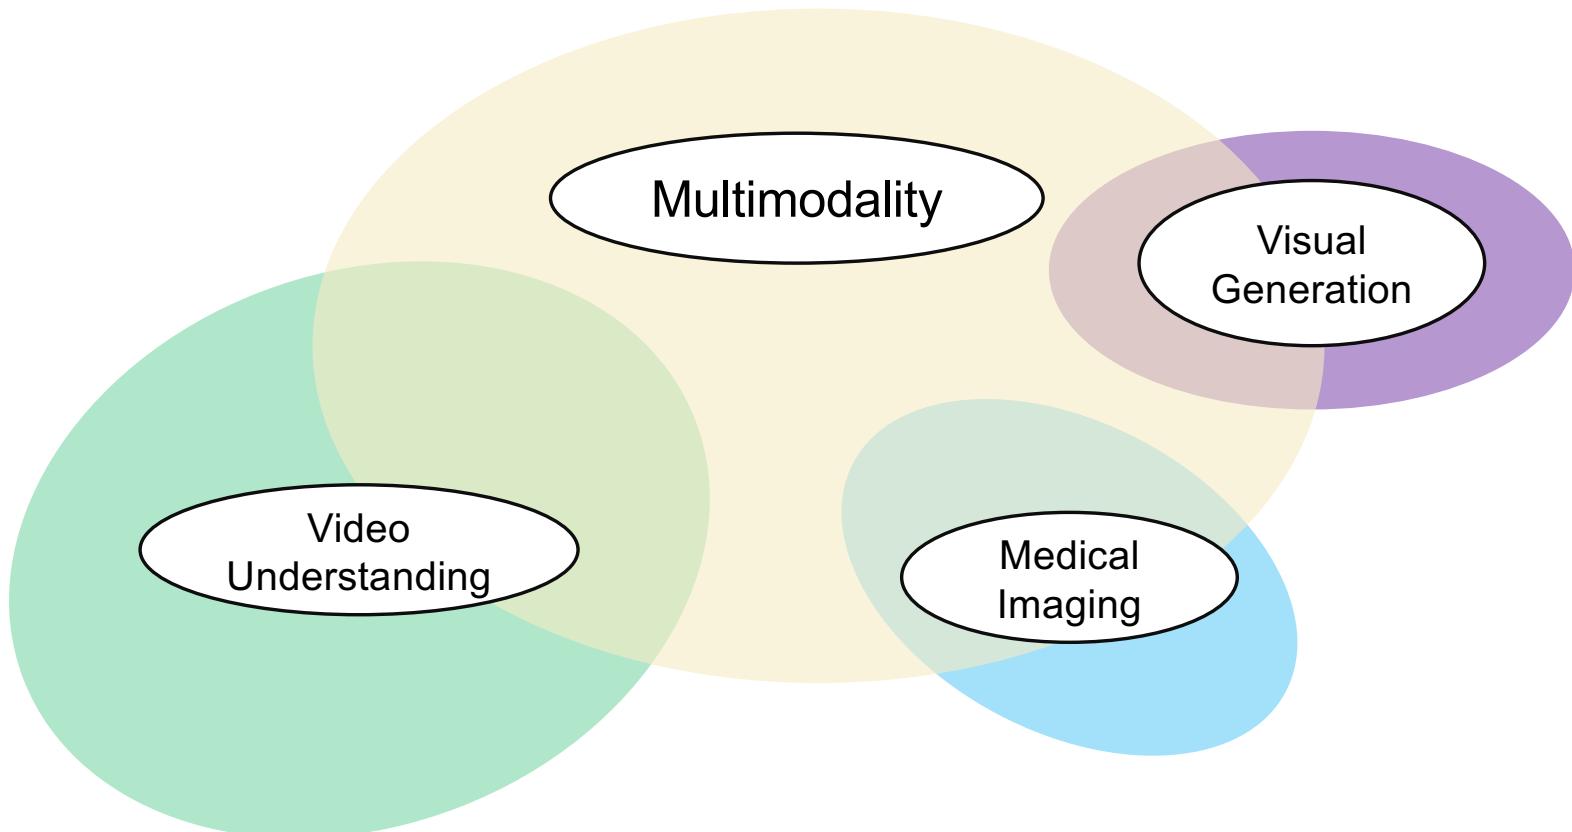


About me

- *Assistant Professor*, 2020 –
 - VISTA Group, Ecole Polytechnique, France
 - Main CV and genAI researcher at Polytechnique
- *Research Fellow*, 2019 – 2021
- *Post-doc*, 2018 – 2020
 - Visual Geometry Group, University of Oxford, UK
 - Andrew Zisserman
- *PhD*, 2013 – 2017
 - University of Edinburgh, UK, INRIA, Grenoble, France
 - Vittorio Ferrari, Cordelia Schmid



Research agenda



Today's lecture

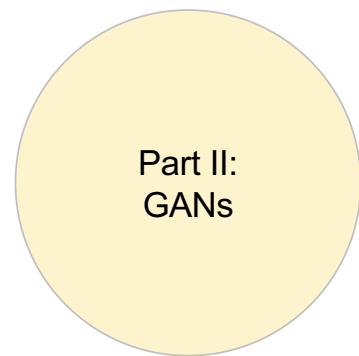
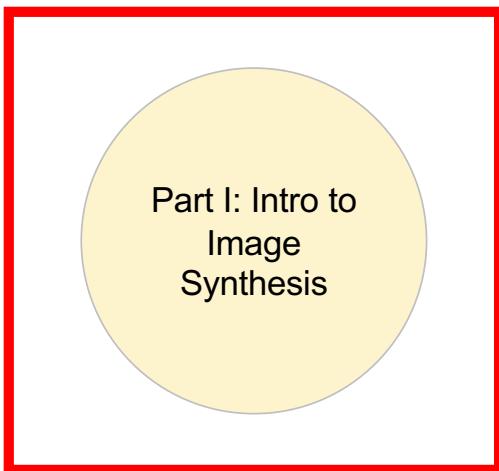
Part I: Intro to
Image
Synthesis

Part II:
GANs

Part III:
GAN editing

Slides adapted from many resources:
[Fei-Fei Li & Andrej Karpathy & Justin Johnson & Ross Girshick & VGG]

Today's lecture



Slides adapted from many resources:
[Fei-Fei Li & Andrej Karpathy & Justin Johnson & Ross Girshick & VGG]

How do we understand the world?

- Our brains are like powerful computers, constantly predicting what will happen next based on past experiences
- Just like we use past knowledge to guess future events, generative AI attempts to forecast and create based on what it has learned



The Future of AI with Generative Models

- Generative AI could be a critical part of future AI advancements
- It combines creating new things and understanding causality—how one action leads to another
- Like us, these AI systems can simulate different scenarios and outcomes, helping to make better decisions



Precision Meets Creativity: A new era in AI

- Predict
- Classify
- Recommend

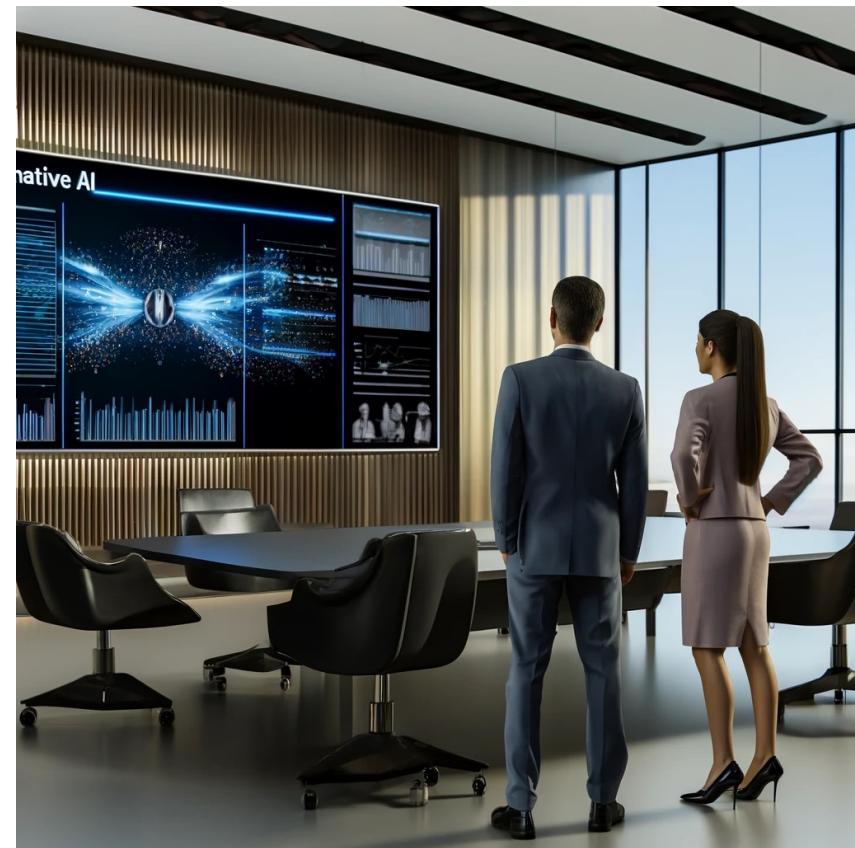


- Imagine
- Augment
- Create

AGI? Artificial General Intelligence

Transforming Business with Generative AI

- Innovating product design
 - Generative Games, word synthesis
 - Shoes, jewelry, clothing, ...
 - Rapid Prototyping
 - Material Innovation
- Enhancing customer interaction
 - Personalized recommendations
 - Automated customer support
 - Interactive virtual assistant
 - Sentiment Analysis
- Streamlining content creation
 - 3D house models
 - AI-assisted Video Production, Choreography
 - Dynamic Web Content
 - Marketing Material Creation



Creative palette

- Images
- Audio
- Text
- 3D models
-



Information

Introduction



2014



2015



2016



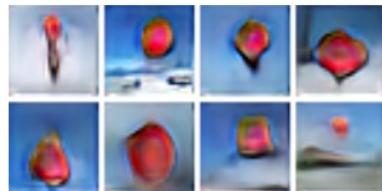
2017

2014 GAN

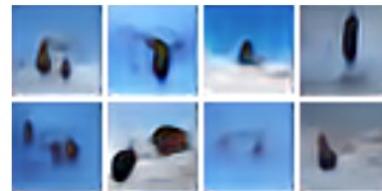
2017 PGGAN



Introduction



A stop sign is flying in blue skies.



A herd of elephants flying in the blue skies.



A toilet seat sits open in the grass field.



A person skiing on sand clad vast desert.



Introduction

This small bird has a yellow crown and a white belly.



This bird has a blue crown with white throat and brown secondaries.



This bird has a red head, throat and chest, with a white belly.



A primarily black bird with streaks of white and yellow and a medium sized beak.



People at the park flying kites and walking.



The bathroom with the white tile has been cleaned.



Multiple people are standing on the beach at the edge of the water.



A clock that is on the side of a tower.



2014 GAN

2017 PGGAN

2019 DM-GAN

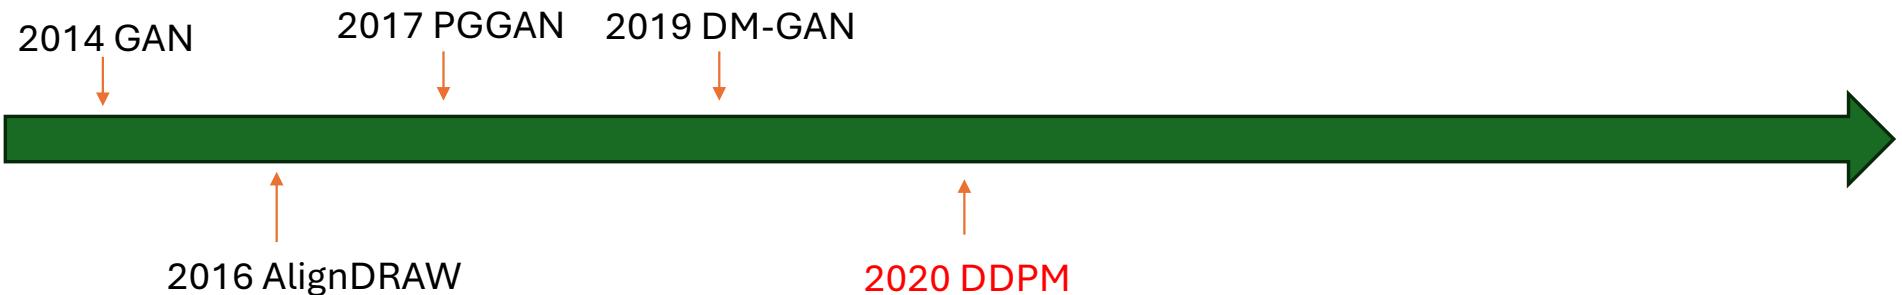
2016 AlignDRAW



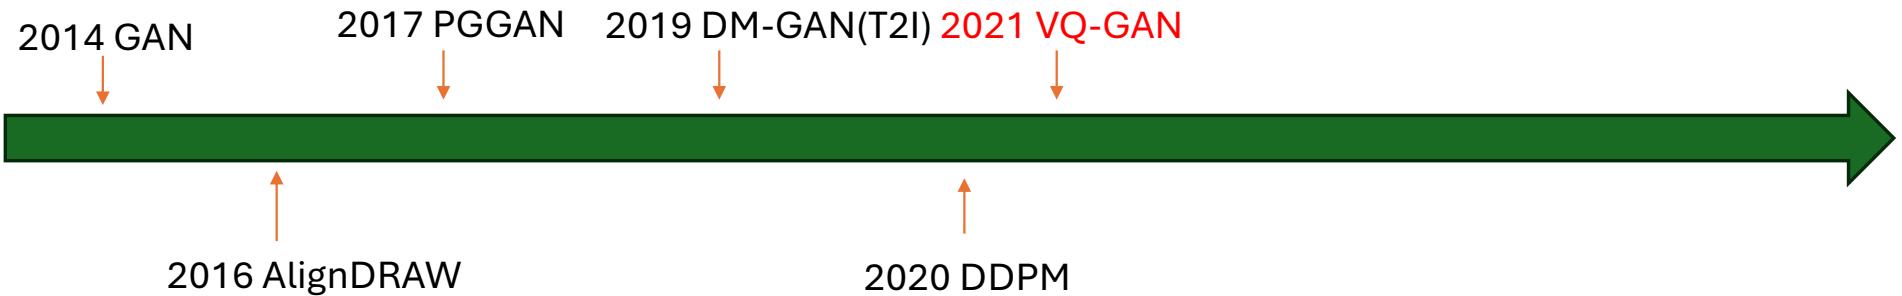
Introduction



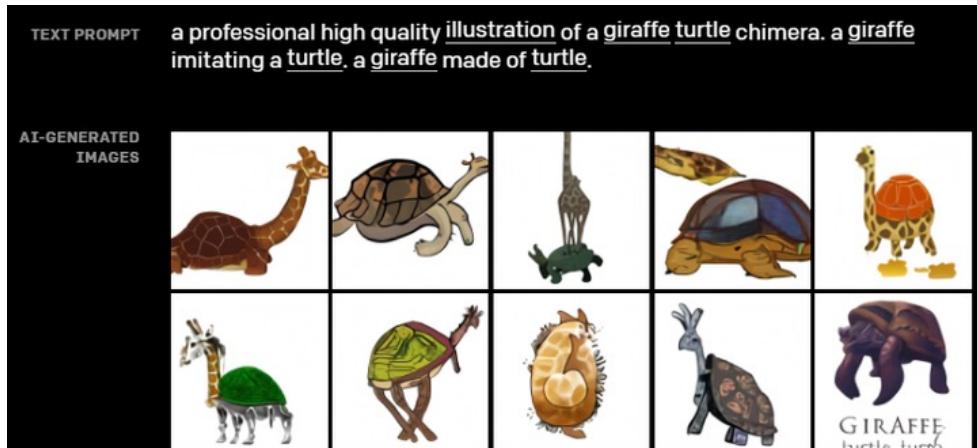
Figure 1: Generated samples on CelebA-HQ 256×256 (left) and unconditional CIFAR10 (right)



Introduction



Introduction



2014 GAN

2017 PGGAN

2019 DM-GAN(T2I)

2021 VQ-GAN

2022 DALLE, GLIDE

2016 AlignDRAW

2020 DDPM

Vicky Kalogeiton

Lecture 1: CSC_52087_EP

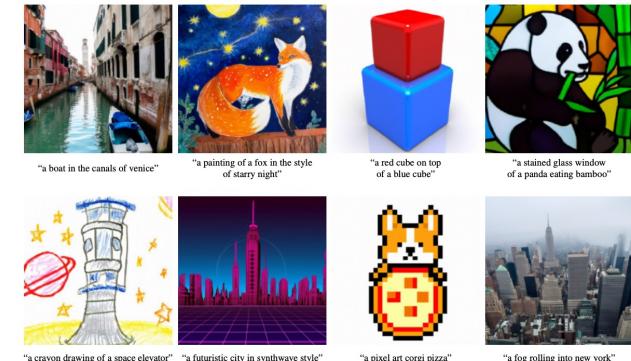


Figure 1. Selected samples from GLIDE using classifier-free guidance. We observe that our model can produce photorealistic images with shadows and reflections, can compose multiple concepts in the correct way, and can produce artistic renderings of novel concepts. For random sample grids, see Figure 17 and 18.

GLIDE

Examples



<http://www.youtube.com/watch?v= x9AwxfjxvE&t=1m45s>

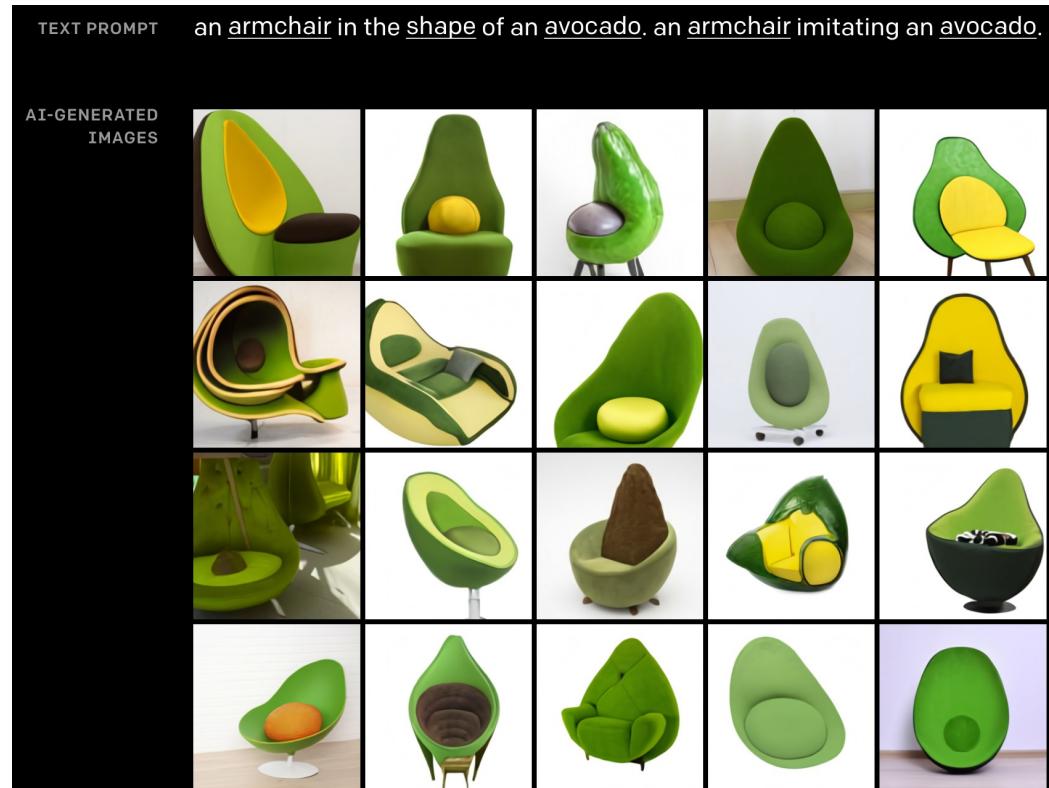
Vicky Kalogeiton

Lecture 1: CSC_52087_EP

5 January 2021



DALLE



<https://openai.com/blog/dall-e/>

Vicky Kalogeiton

Lecture 1: CSC_52087_EP

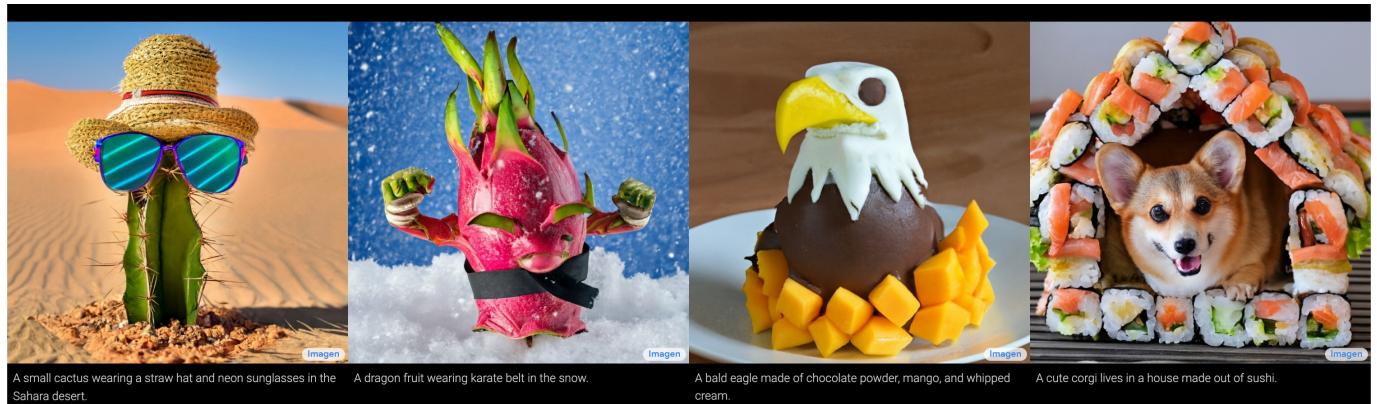
23 May 2022



Imagen by Google AI



A photo of a Corgi dog riding a bike in Times Square. It is wearing sunglasses and a beach hat.



<https://imagen.research.google/>

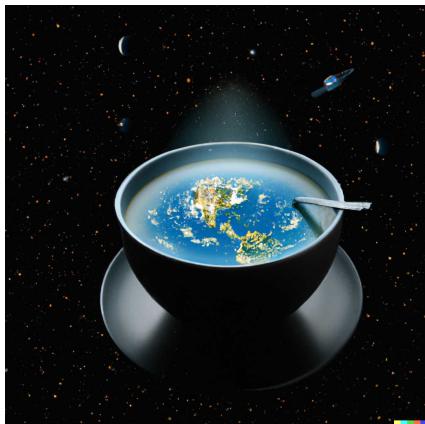
Vicky Kalogeiton

Lecture 1: CSC_52087_EP

13 April 2022



Dalle-2 (Text-to-Image)



A bowl of soup as a planet in the universe



An astronaut riding a horse in a photorealistic style



Teddy bears mixing sparkling chemicals as mad scientists

Diffusion Models



OpenAI: DALL-E3



Midjourney

music, audio, animation, video, physical etc....

Stable Diffusion 3

February 2024



Vicky Kalogeiton

Lecture 1: CSC_52087_EP

Make-A-Video (Text-to-Video)



A confused grizzly bear
in a calculus class



A golden retriever eating ice
cream on a beautiful tropical
beach at sunset, high resolution



A panda playing on a
swing set

SORA (Text-to-Video)

February 2024



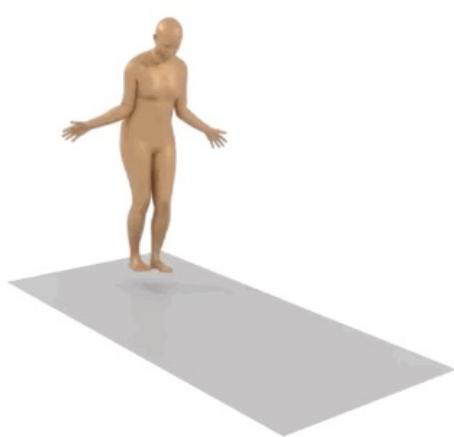
Vicky Kalogeiton

Lecture 1: CSC_52087_EP

Human Motion Diffusion (Text-to-Motion)



“A person punches in a manner consistent with martial arts.”

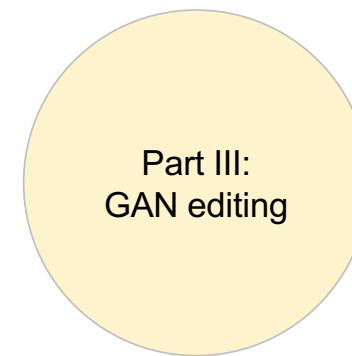
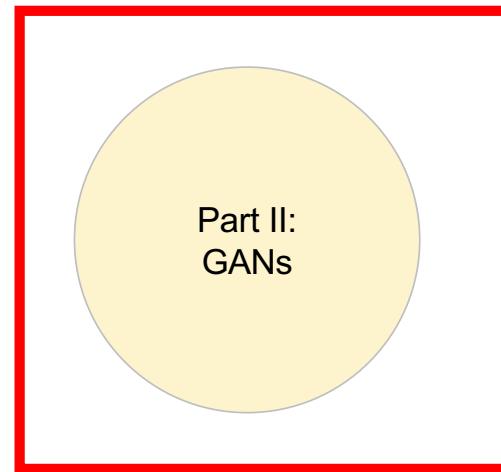
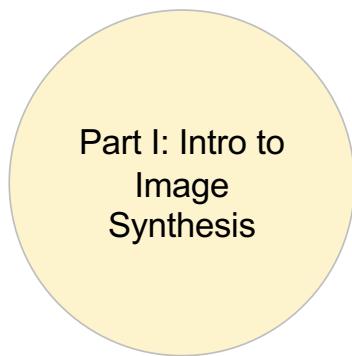


“A person is skipping rope.”



“a man kicks with something or someone with his left leg.”

Today's lecture



Slides adapted from many resources:
[Fei-Fei Li & Andrej Karpathy & Justin Johnson & Ross Girshick & VGG]

Part II: GANs

- Relation among models
- Training objective
- Training GANs
- DCGAN
- Tricks for improving performance
- StyleGAN
- Metrics
- Conditional GANs
- Image Translation



Relation among models

Relation among all models

Bayes' Rule:

$$P(x|y) = \frac{P(y|x)}{P(y)} P(x)$$

Relation among all models

Bayes' Rule:

$$P(x|y) = \frac{P(y|x)}{P(y)}$$

\$P(x|y)\$ \$P(y|x)\$ \$P(y)\$
 Conditional Generative Model Discriminative model Prior over labels
\$P(x)\$ (unconditional)
Generative model

What can we do with them?

- **Discriminative Model:** Learn a probability distribution $p(y/x)$
 - Assign labels to data
 - Feature Learning using labels
- **Generative Model:** Learn a probability distribution $p(x)$
- **Conditional Generative Model:** Learn $p(x/y)$

What can we do with them?

- **Discriminative Model:** Learn a probability distribution $p(y/x)$
 - Assign labels to data
 - Feature Learning using labels
- **Generative Model:** Learn a probability distribution $p(x)$
 - Detect Outliers
 - Unsupervised Feature Learning
- **Conditional Generative Model:** Learn $p(x/y)$

What can we do with them?

- **Discriminative Model:** Learn a probability distribution $p(y/x)$
 - Assign labels to data
 - Feature Learning using labels
- **Generative Model:** Learn a probability distribution $p(x)$
 - Detect Outliers
 - Unsupervised Feature Learning
 - Sample to **generate** new data
- **Conditional Generative Model:** Learn $p(x/y)$

What can we do with them?

- **Discriminative Model:** Learn a probability distribution $p(y/x)$
 - Assign labels to data
 - Feature Learning using labels
- **Generative Model:** Learn a probability distribution $p(x)$
 - Detect Outliers
 - Unsupervised Feature Learning
 - Sample to **generate** new data
- **Conditional Generative Model:** Learn $p(x/y)$
 - Generate new data conditioned on input labels



Generative Adversarial Networks

Generative Adversarial Networks

- Assume data x_i drawn from distribution $p_{\text{data}}(x)$.
 - Want to sample from p_{data}

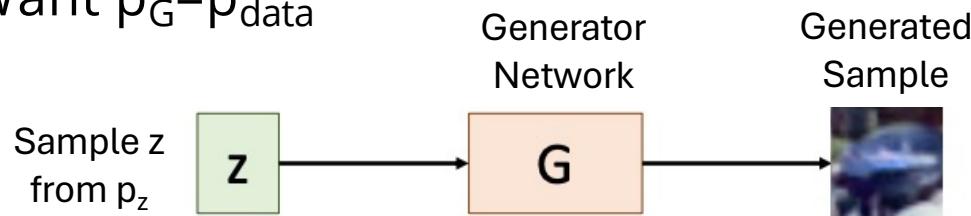
Introduced in 2014 by Ian Goodfellow

Generative Adversarial Networks

- Assume data x_i drawn from distribution $p_{\text{data}}(x)$.
 - Want to sample from p_{data}
- **Idea:** Introduce a latent variable z with simple prior $p(z)$
 - Sample $z \sim p(z)$ and pass to a Generator Network $x = G(z)$
 - X is a sample from the Generator distribution p_G .
 - Want $p_G = p_{\text{data}}$

Generative Adversarial Networks

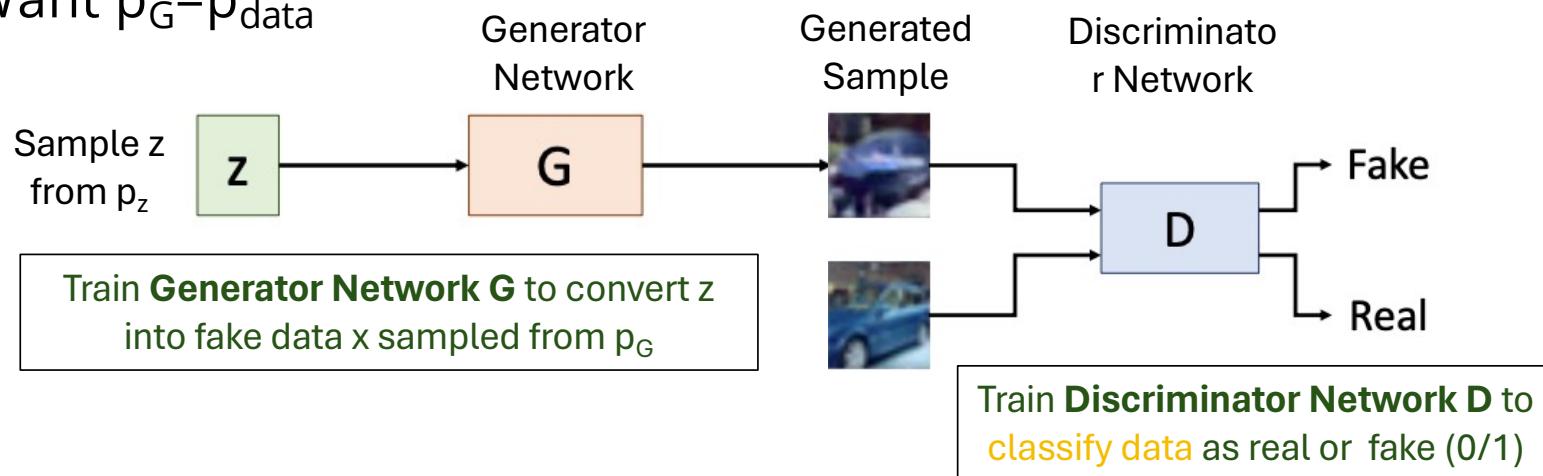
- Assume data x_i drawn from distribution $p_{\text{data}}(x)$.
 - Want to sample from p_{data}
- **Idea:** Introduce a latent variable z with simple prior $p(z)$
 - Sample $z \sim p(z)$ and pass to a Generator Network $x=G(z)$
 - X is a sample from the Generator distribution p_G .
 - Want $p_G = p_{\text{data}}$



Train **Generator Network G** to convert z
 into fake data x sampled from p_G

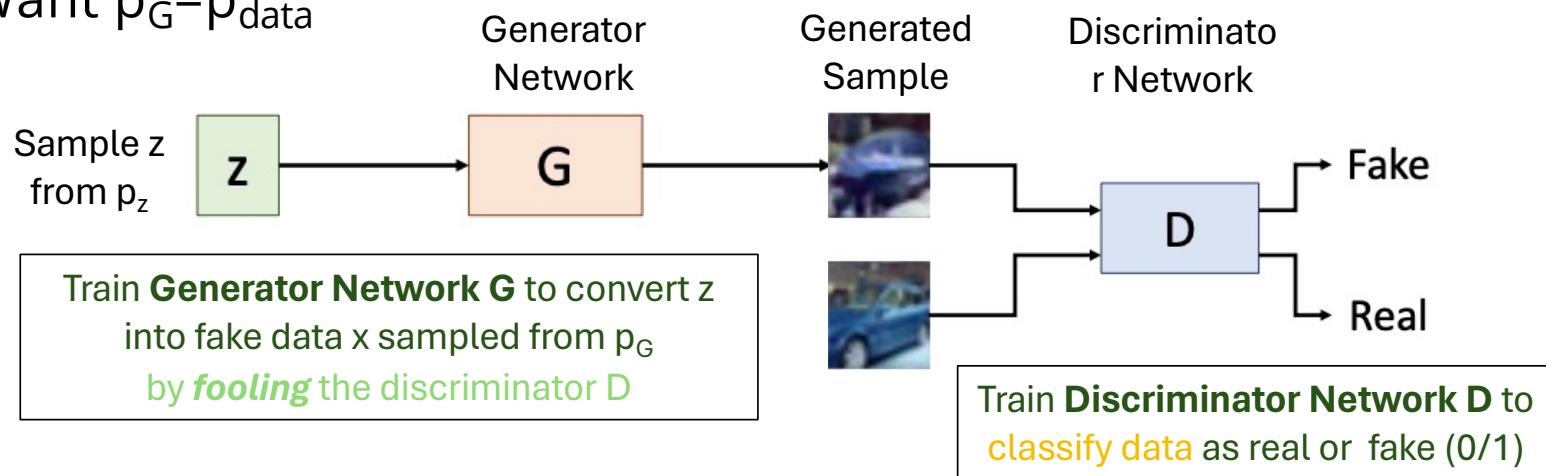
Generative Adversarial Networks

- Assume data x_i drawn from distribution $p_{\text{data}}(x)$.
 - Want to sample from p_{data}
- **Idea:** Introduce a latent variable z with simple prior $p(z)$
 - Sample $z \sim p(z)$ and pass to a Generator Network $x=G(z)$
 - X is a sample from the Generator distribution p_G .
 - Want $p_G=p_{\text{data}}$



Generative Adversarial Networks

- Assume data x_i drawn from distribution $p_{\text{data}}(x)$.
 - Want to sample from p_{data}
- **Idea:** Introduce a latent variable z with simple prior $p(z)$
 - Sample $z \sim p(z)$ and pass to a Generator Network $x=G(z)$
 - X is a sample from the Generator distribution p_G .
 - Want $p_G=p_{\text{data}}$





GANs: Training Objective

GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

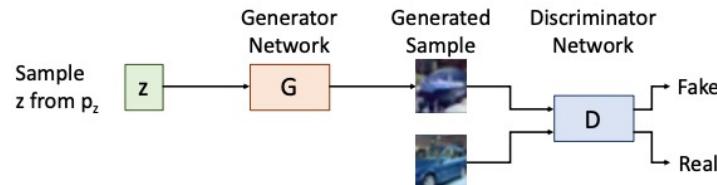
$$\min_{\mathbf{G}} \max_{\mathbf{D}} \left(E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - \mathbf{D}(\mathbf{G}(\mathbf{z})))] \right)$$

GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

Discriminator
wants
 $D(x) = 1$ for real
data

$$\min_G \max_D \left(E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

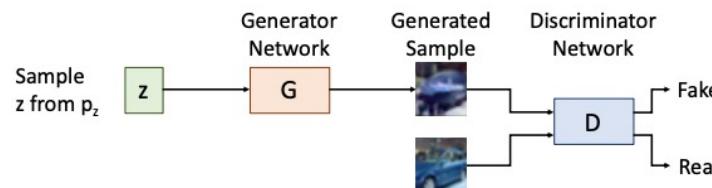


GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

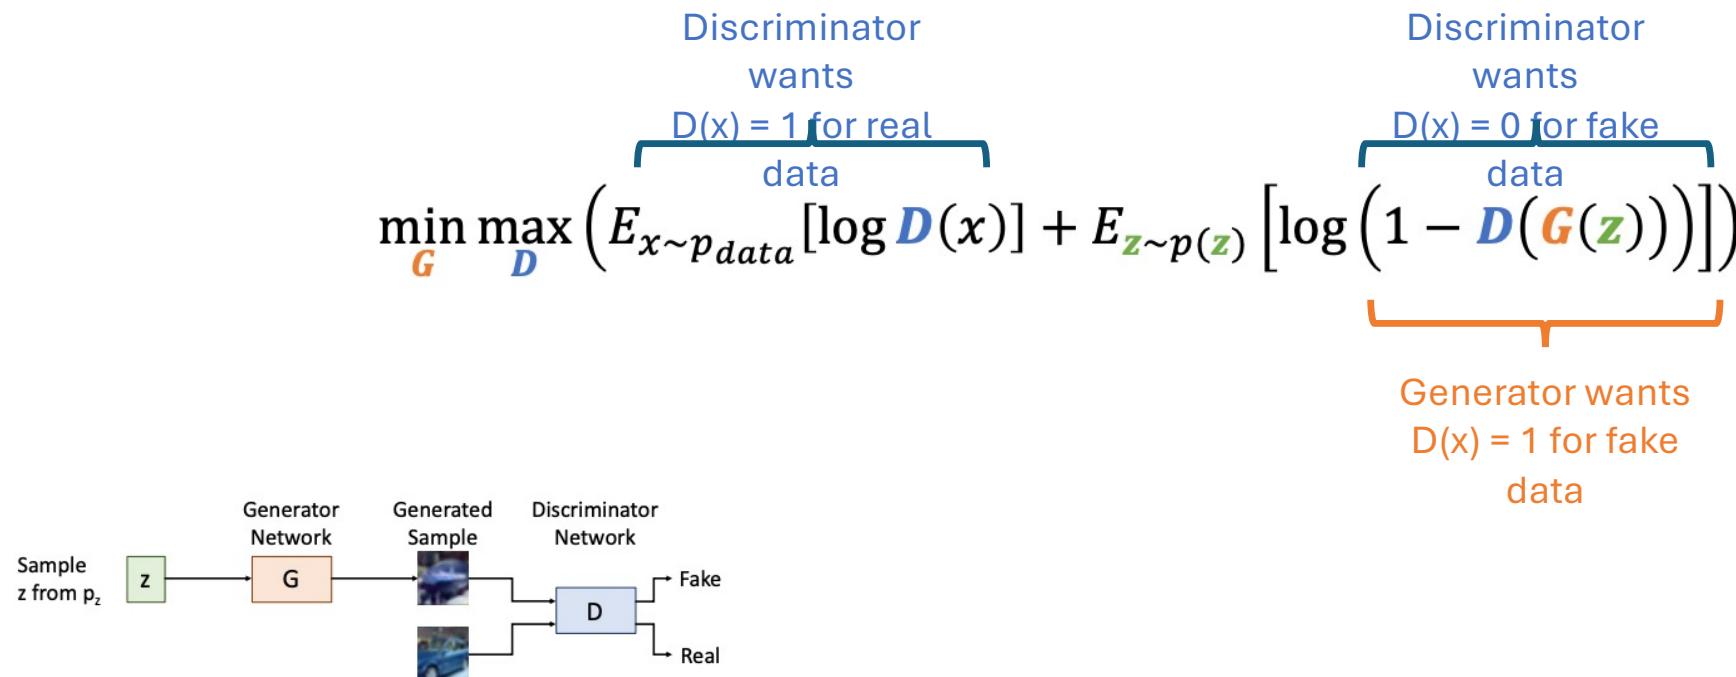
$$\min_G \max_D \left(E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} \left[\log (1 - D(G(z))) \right] \right)$$

Discriminator wants
 $D(x) = 1$ for real data
 $D(x) = 0$ for fake data



GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game



GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

Train G and D using alternating gradient updates

$$\min_{\mathbf{G}} \max_{\mathbf{D}} \left(E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - \mathbf{D}(\mathbf{G}(\mathbf{z})))] \right)$$

GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

Train G and D using alternating gradient updates

$$\begin{aligned} & \min_{\mathbf{G}} \max_{\mathbf{D}} \left(E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - \mathbf{D}(\mathbf{G}(\mathbf{z})))] \right) \\ &= \min_{\mathbf{G}} \max_{\mathbf{D}} \mathbf{V}(\mathbf{G}, \mathbf{D}) \end{aligned}$$

GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

Train G and D using alternating gradient updates

$$\begin{aligned}
 & \min_{\mathbf{G}} \max_{\mathbf{D}} \left(E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} \left[\log \left(1 - \mathbf{D}(\mathbf{G}(\mathbf{z})) \right) \right] \right) \\
 &= \min_{\mathbf{G}} \max_{\mathbf{D}} \mathbf{V}(\mathbf{G}, \mathbf{D}) \quad \text{For t in 1, ... T:} \\
 &\quad 1. \text{ (Update D)} \quad \mathbf{D} = \mathbf{D} + \alpha_{\mathbf{D}} \frac{\partial \mathbf{V}}{\partial \mathbf{D}} \\
 &\quad 2. \text{ (Update G)} \quad \mathbf{G} = \mathbf{G} - \alpha_{\mathbf{G}} \frac{\partial \mathbf{V}}{\partial \mathbf{G}}
 \end{aligned}$$

GANs: Training Objective

- Jointly train generator G and discriminator D with a minimax game

Train G and D using alternating gradient updates

$$\min_{\mathbf{G}} \max_{\mathbf{D}} \left(E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} \left[\log \left(1 - \mathbf{D}(\mathbf{G}(\mathbf{z})) \right) \right] \right)$$

$$= \min_{\mathbf{G}} \max_{\mathbf{D}} \mathcal{V}(\mathbf{G}, \mathbf{D})$$

We are not minimizing any overall loss! No training curves to look at!

For t in 1, ... T:

- (Update D) $\mathbf{D} = \mathbf{D} + \alpha_{\mathbf{D}} \frac{\partial \mathcal{V}}{\partial \mathbf{D}}$
- (Update G) $\mathbf{G} = \mathbf{G} - \alpha_{\mathbf{G}} \frac{\partial \mathcal{V}}{\partial \mathbf{G}}$

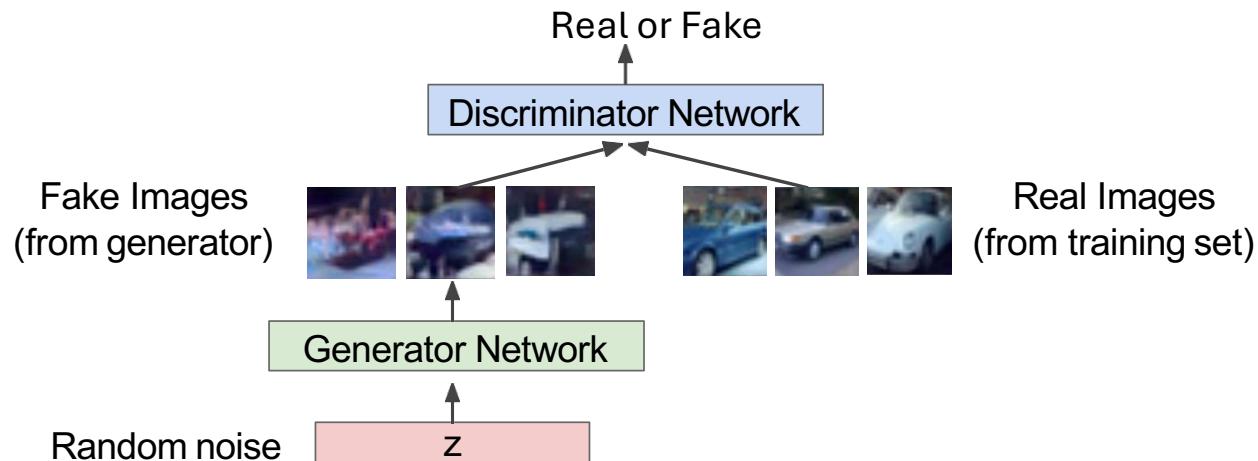


Training GANs

Training GANs: Two-player game

Generator network: try to fool the discriminator by generating real-looking images

Discriminator network: try to distinguish between real and fake images

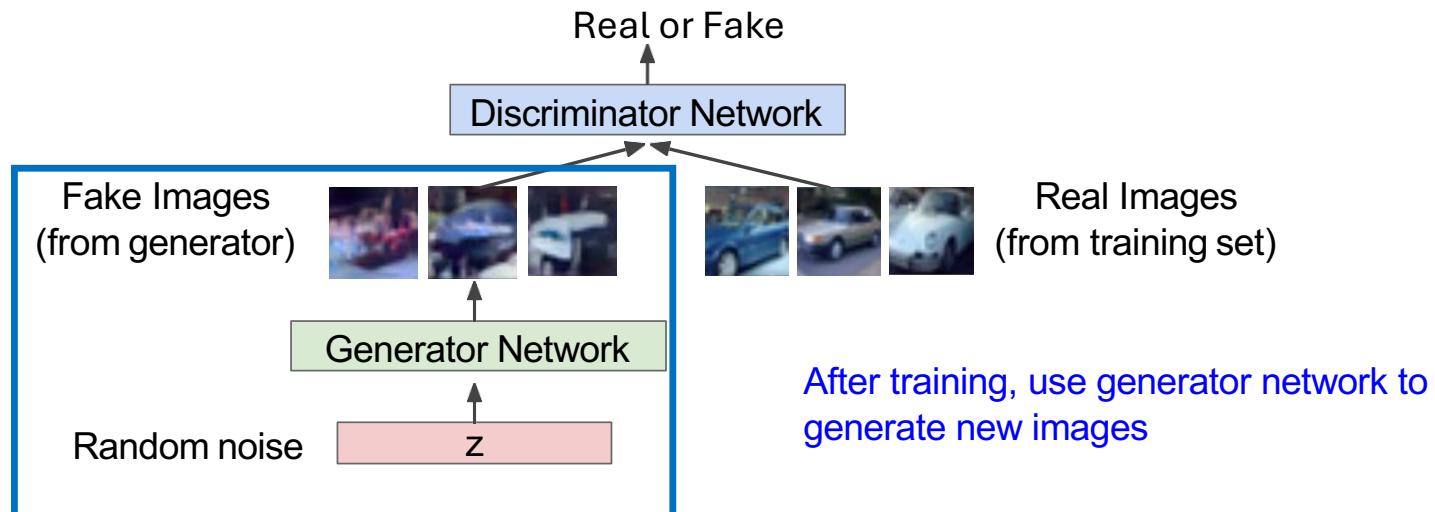


Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

Training GANs: Two-player game

Generator network: try to fool the discriminator by generating real-looking images

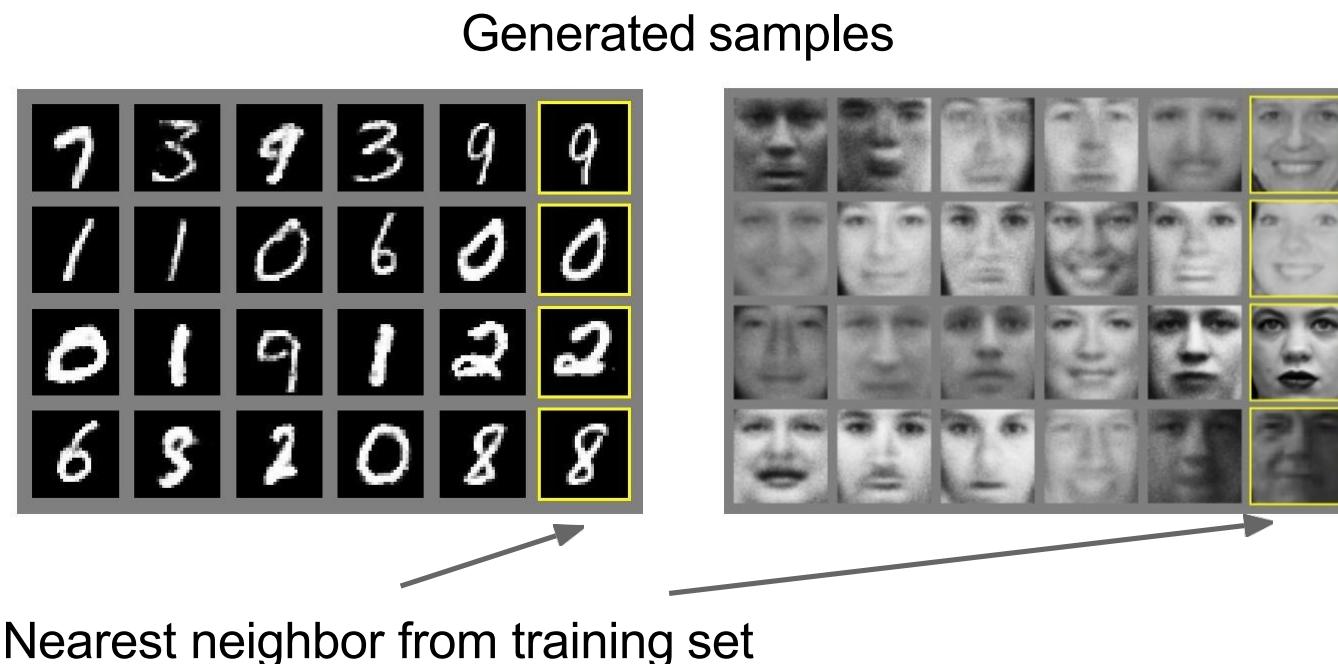
Discriminator network: try to distinguish between real and fake images



Ian Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014

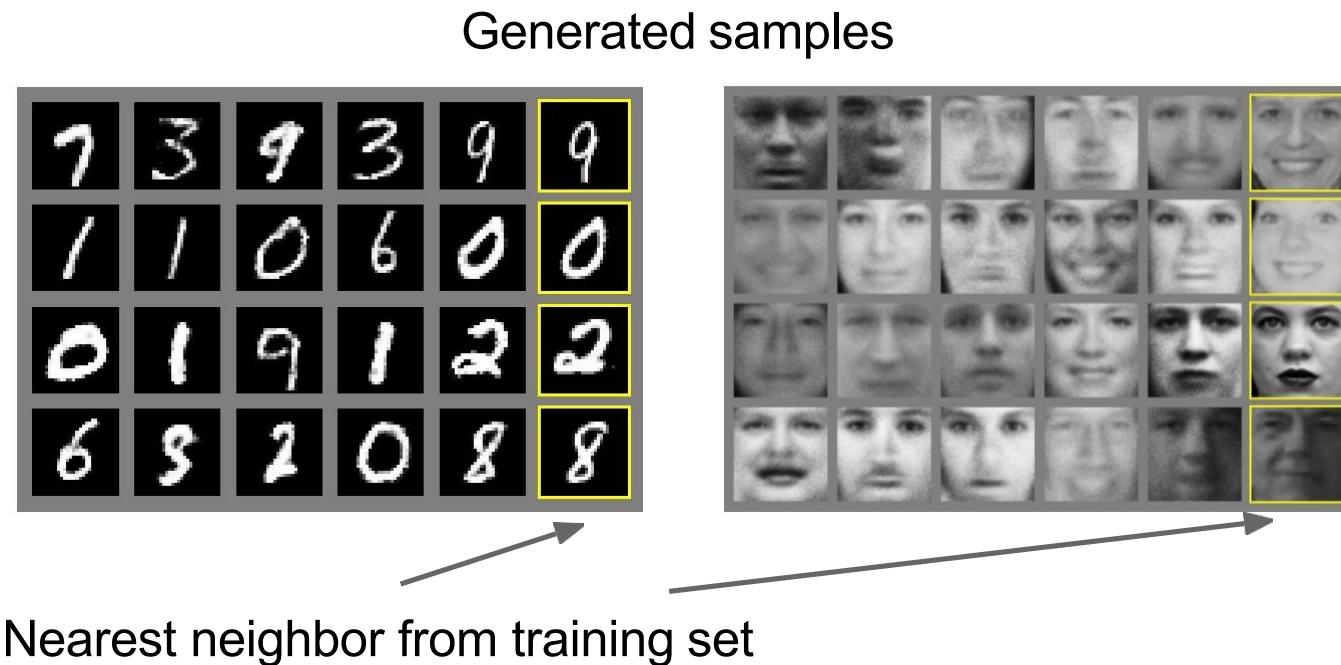
Examples of GANs!

Ian Goodfellow et al., "Generative Adversarial Nets", NIPS 2014



Training GANS is hard! Mode-collapse

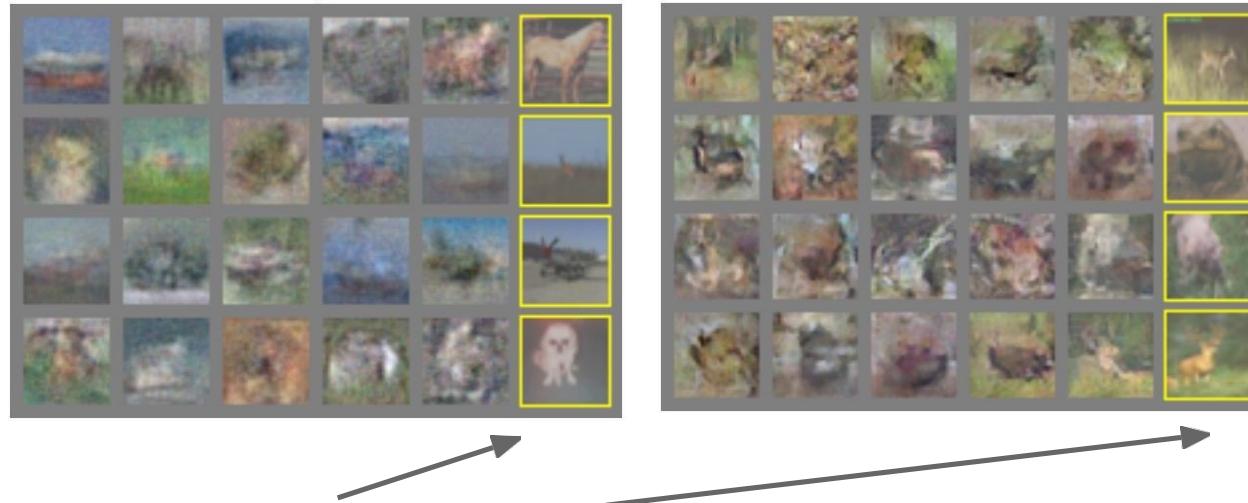
Generator can “memorize” real images



Training GANS is hard!

Generator can produce “fuzzy” images

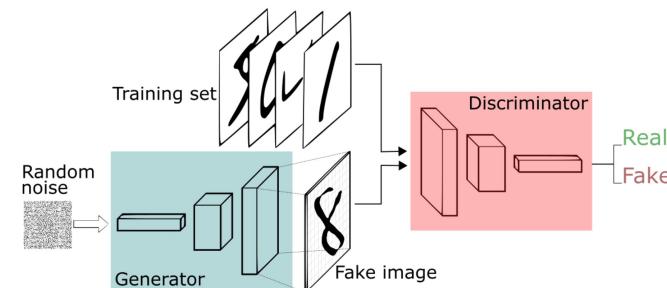
Generated samples (CIFAR-10)



Nearest neighbor from training set

Training GANS is hard!

- Training GANs is still very hard
 - Many problems exist
 - Non-convergence
 - The models never converge and worse they become unstable
 - Mode collapse
 - The generator produces a single or limited modes
 - i.e. the images are not as diverse as the true data
- Many tricks exist





Deep Convolutional (DC) GAN

Original GAN – Goodfellow 2014

- Original GAN paper
- Uses only fully connected layers
 - Limited to generating small images
- Discriminator
 - Binary cross entropy loss



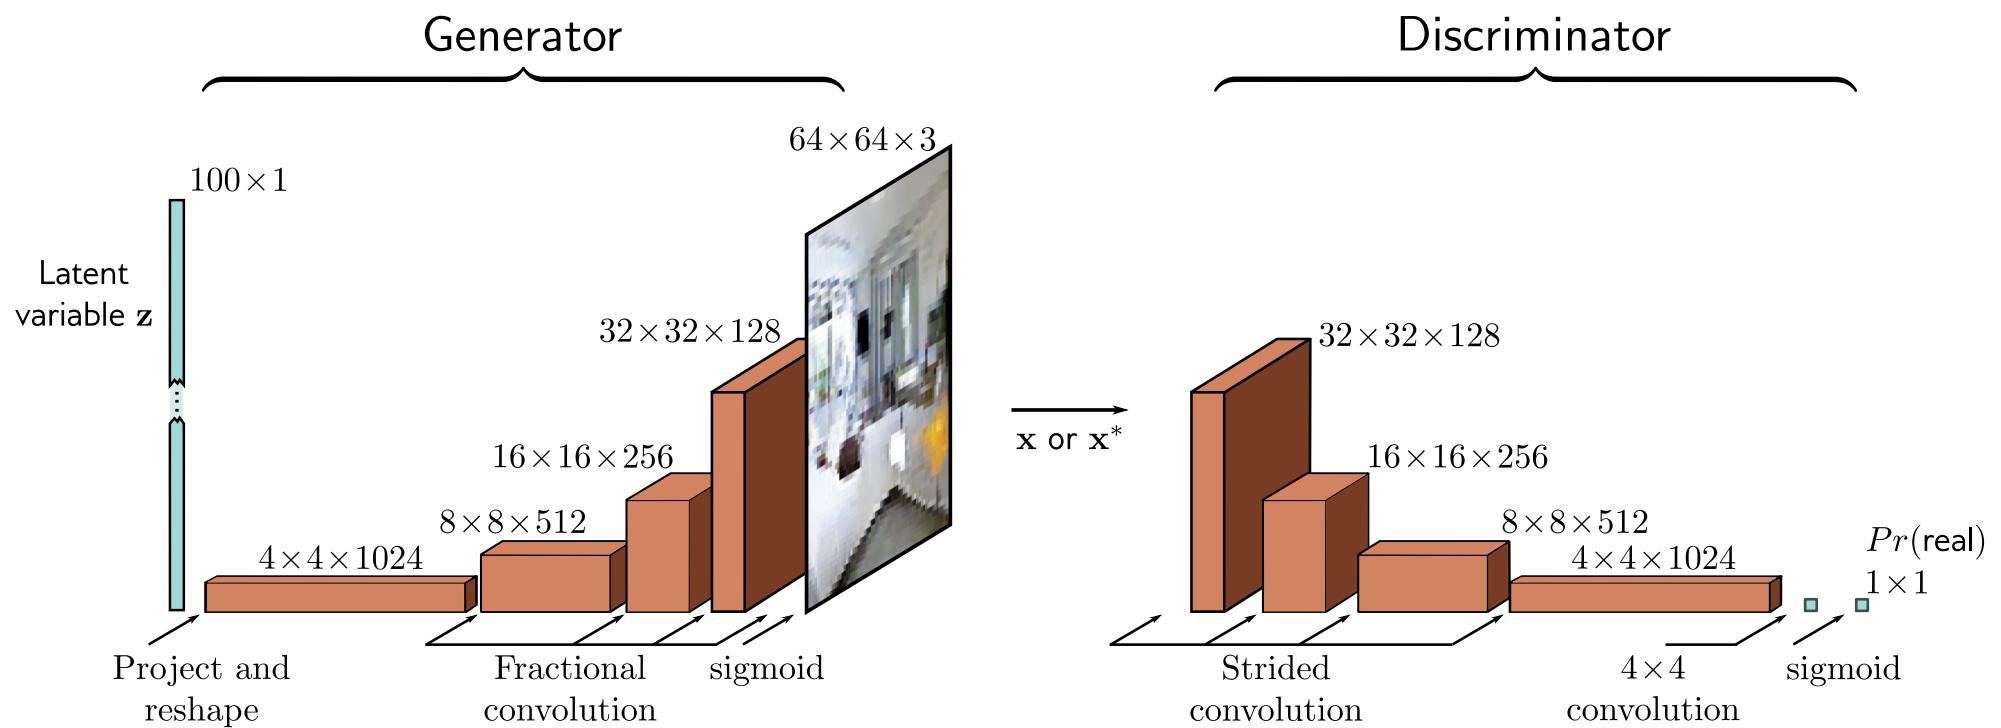
DCGAN 2015/2016

How to do GANs with convolutional layers?

- Replace any pooling layers with strided convolutions (discriminator) and transposed-strided convolutions (generator)
- Use batchnorm in both the generator and the discriminator
- Use LeakyReLU activation in the discriminator for all layers
- Use ReLU activation in generator for all layers except for the output, which uses Tanh
 - Later people recommend using LeakyReLU in both G and D

[Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks.", ICLR 2016]

Deep Convolutional (DC) GAN



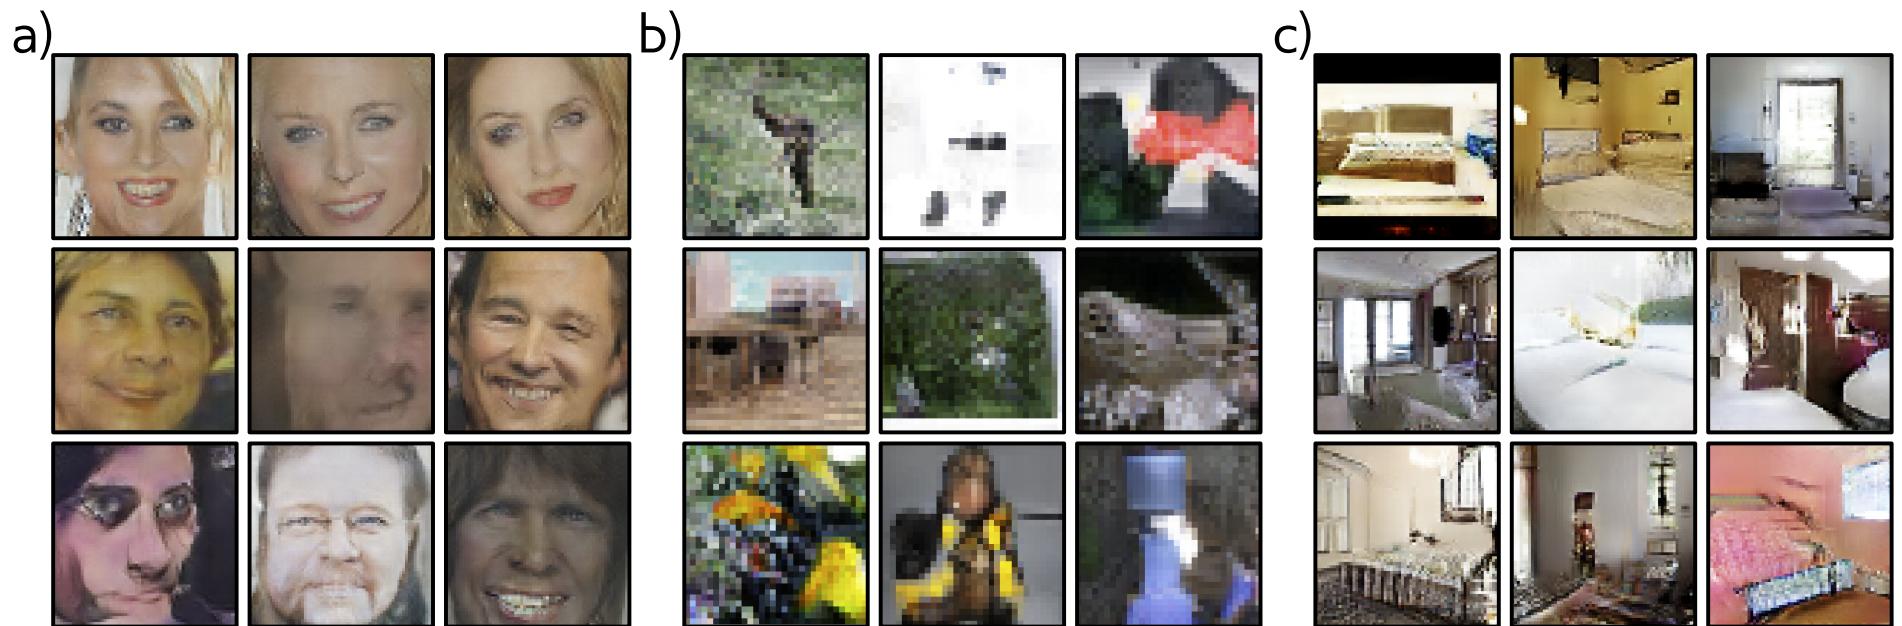
DC GAN Results

Samples from the model look much better!



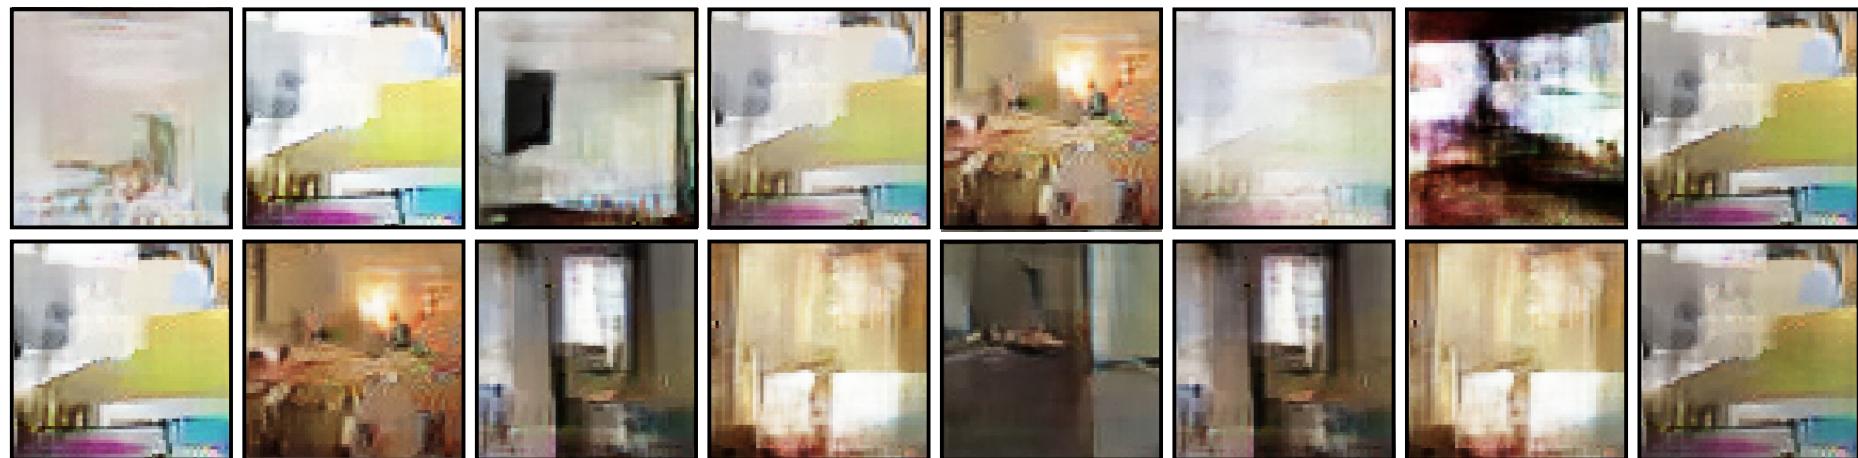
Radford et al, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks", ICLR 2016

DC GAN Results

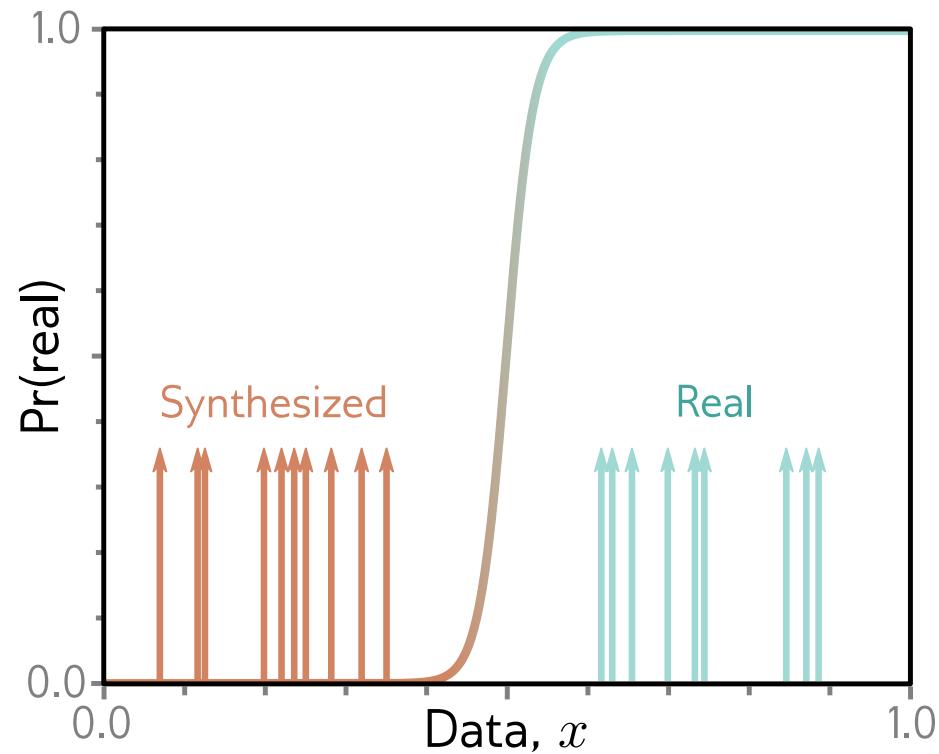


Mode collapse

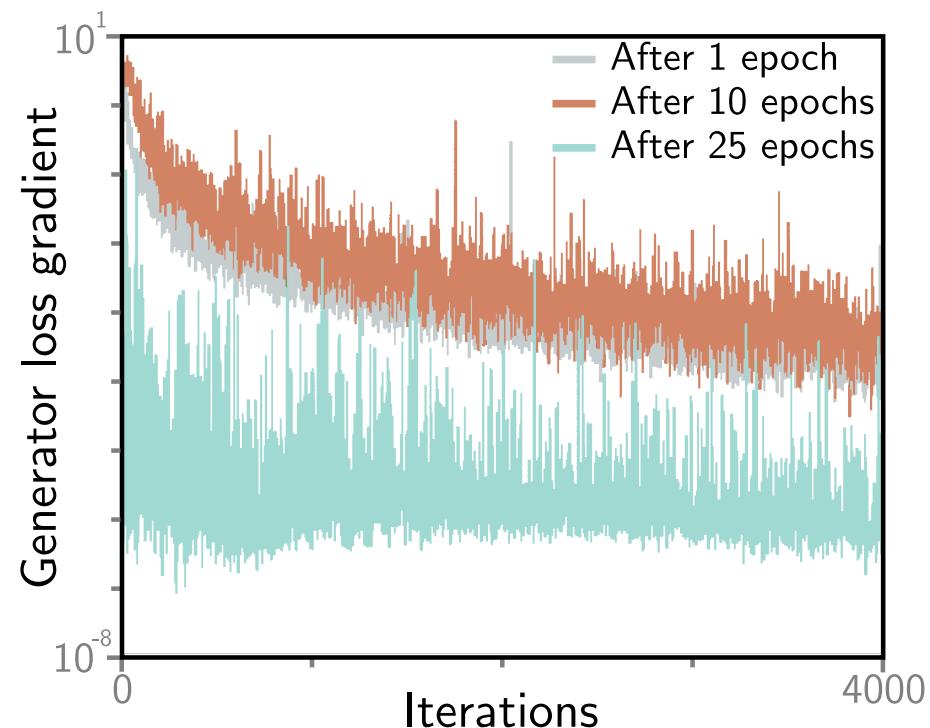
Training GANs is difficult. Often fails.



GAN training instability



GAN training instability



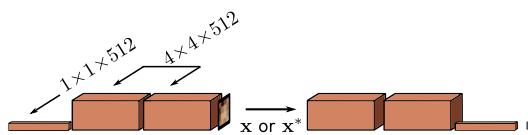
Fix generator and continue training discriminator, then generator gradients disappear.



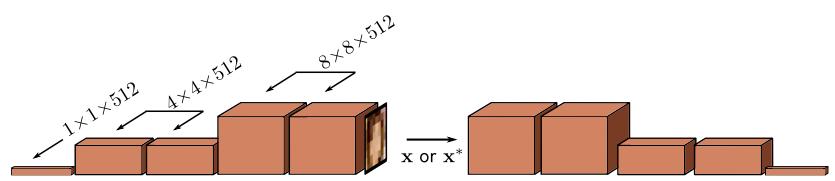
Tricks for improving performance

Trick 1: Progressive growing

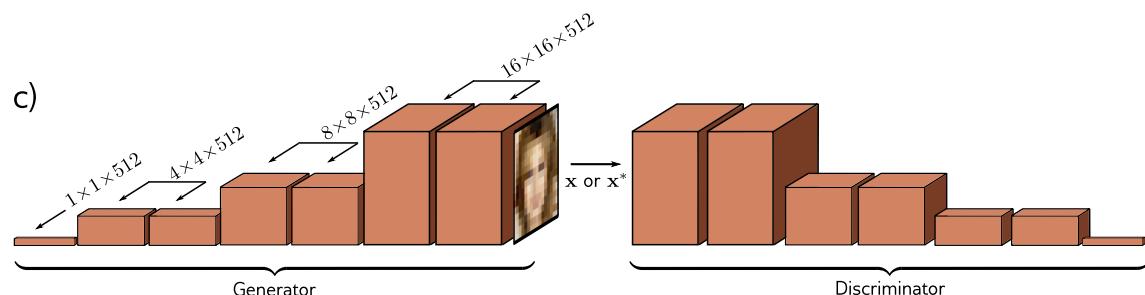
a)



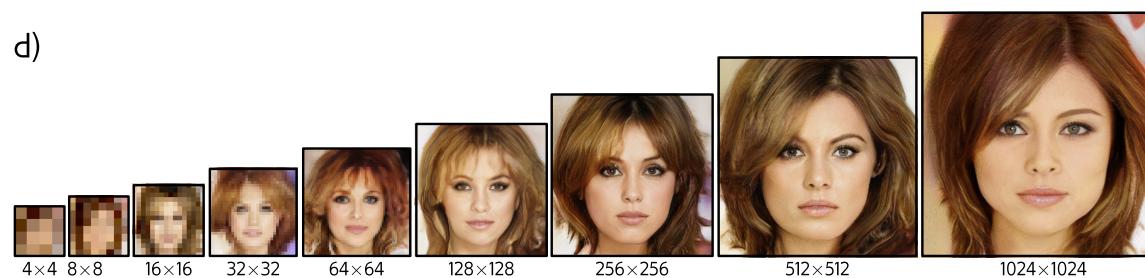
b)



c)



d)



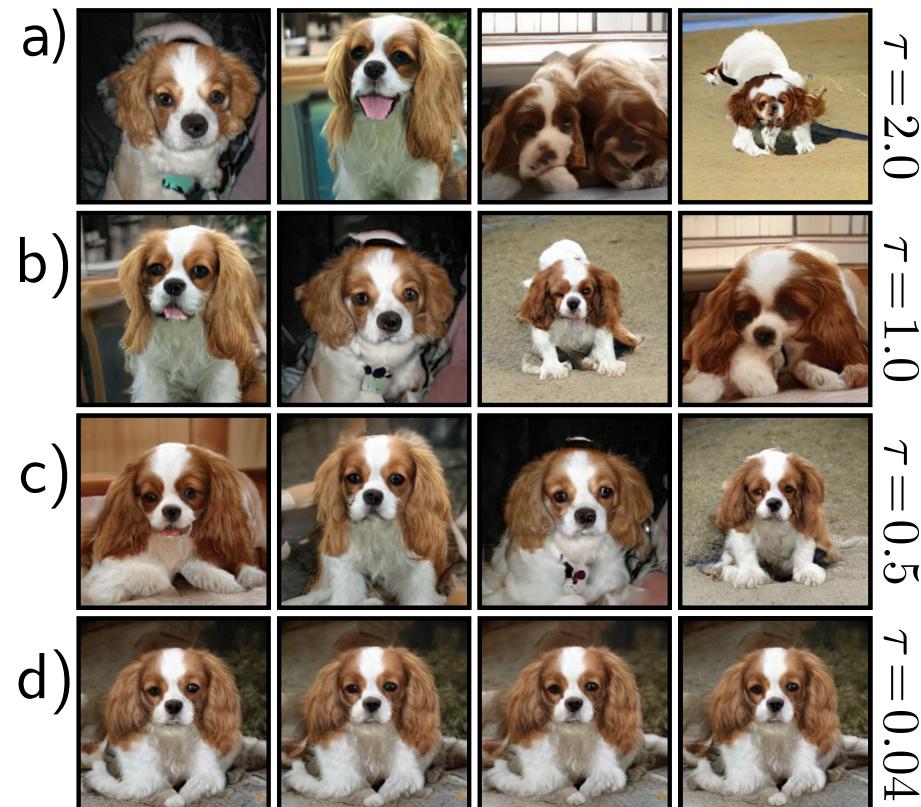


Trick 2: Minibatch discrimination

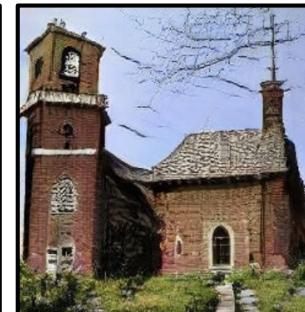
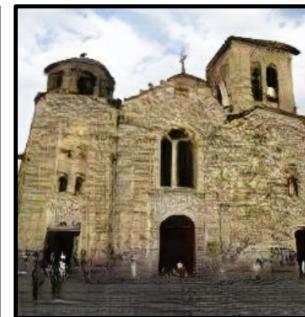
- Add in statistics across minibatch as input to discriminator
- Forces generated batch to have similar variation to real batch

Trick 3: Truncation

Only choose random values of latent variables that are less than a threshold τ .



Example results



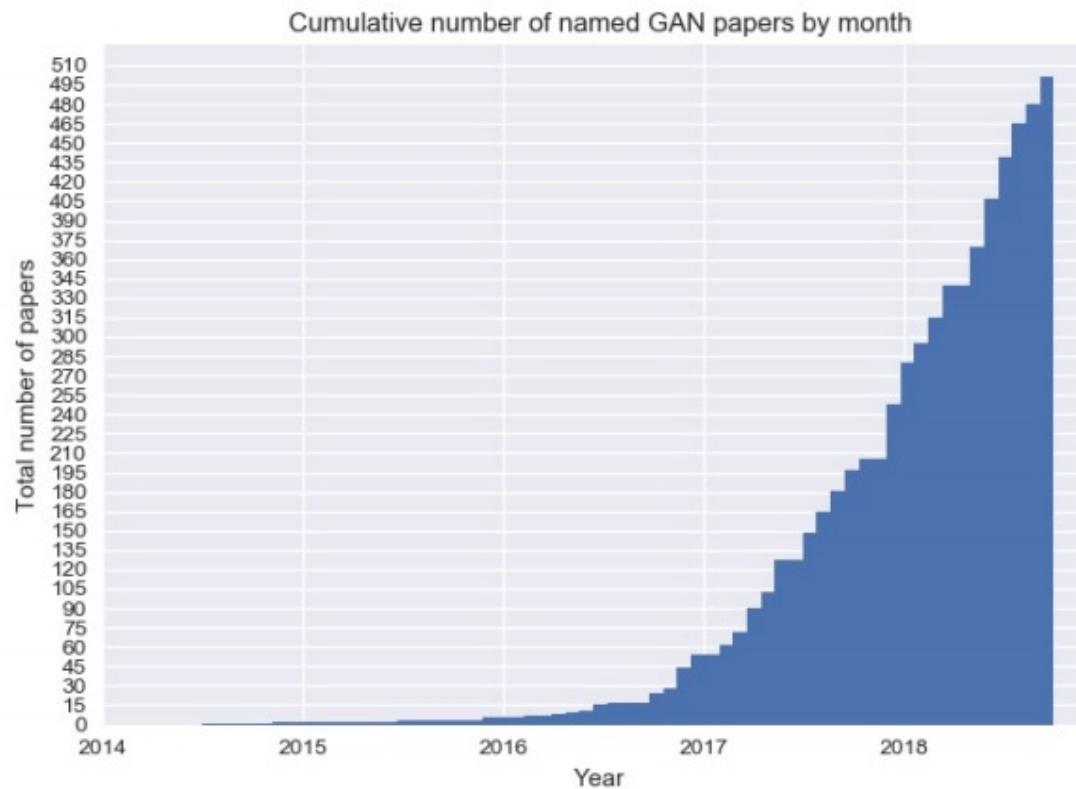
Interpolation





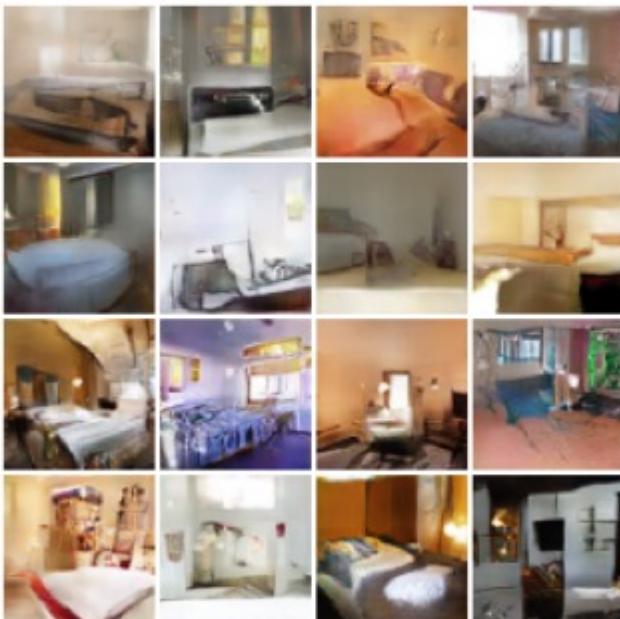
StyleGAN

Generative Adversarial Networks



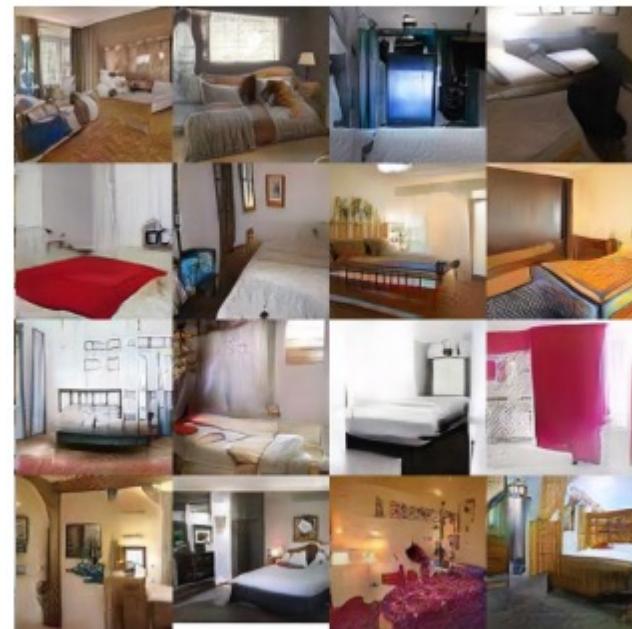
GAN improvements – Wasserstein Loss

Wasserstein GAN (WGAN)



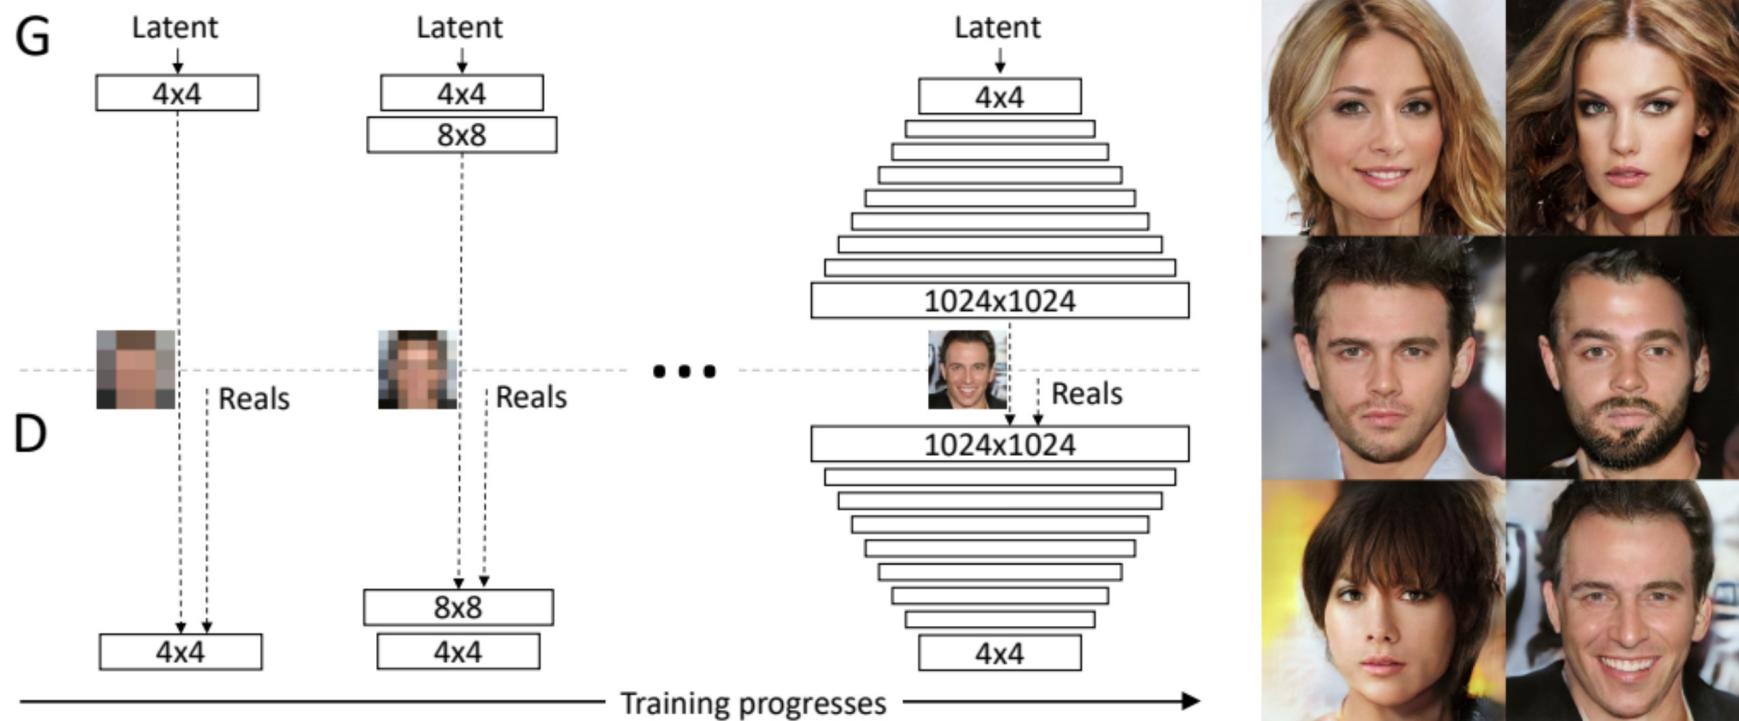
Arjovsky, Chintala, and Bottou, "Wasserstein GAN", 2017

WGAN with Gradient Penalty
(WGAN-GP)



Gulrajani et al, "Improved Training of
Wasserstein GANs", NeurIPS 2017

Progressive GAN, Karras et al ICLR 2018



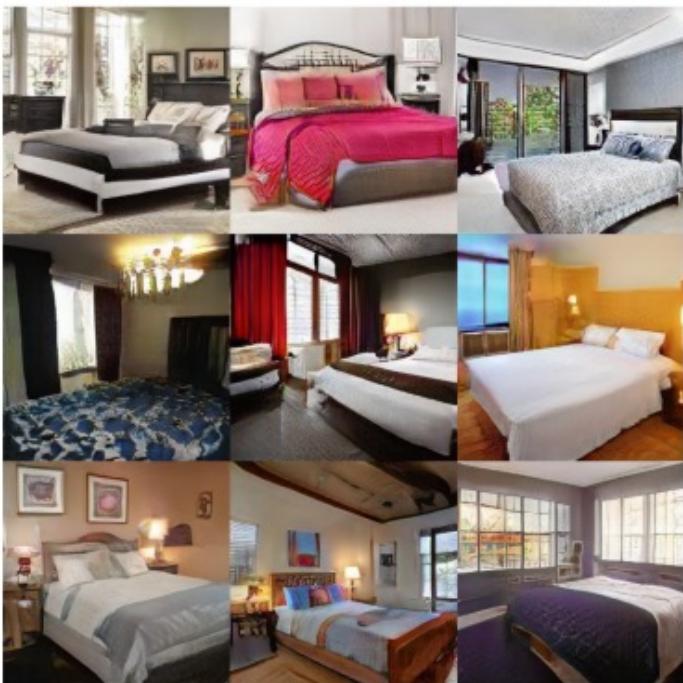
Karras et al, "Progressive Growing of GANs for Improved Quality, Stability, and Variation", ICLR 2018

Vicky Kalogeiton

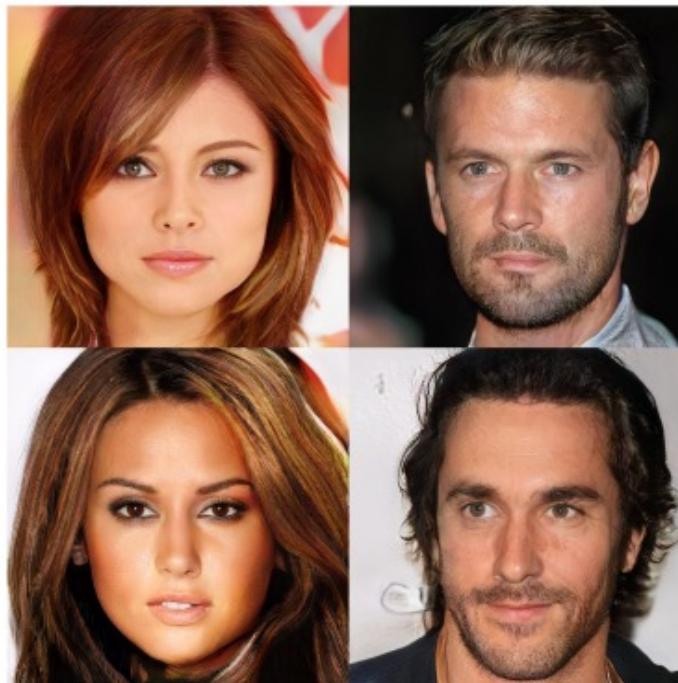
Lecture 1: CSC_52087_EP

GAN improvements – High Resolution Progressive GAN, Karras et al ICLR 2018

256 x 256 bedrooms



1024 x 1024 faces



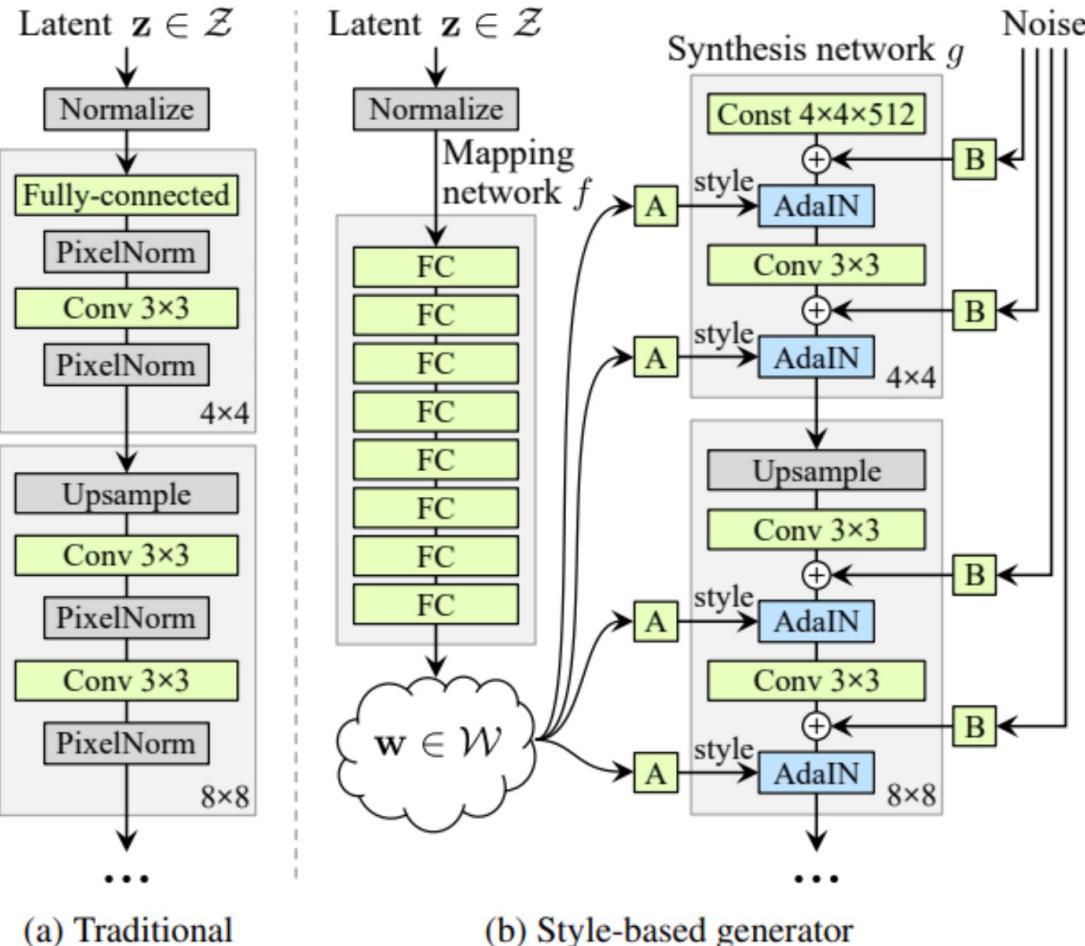
Karras et al, “Progressive Growing of GANs for Improved Quality, Stability, and Variation”, ICLR 2018

Vicky Kalogeiton

Lecture 1: CSC_52087_EP

80

StyleGAN, Karras et al CVPR 2019



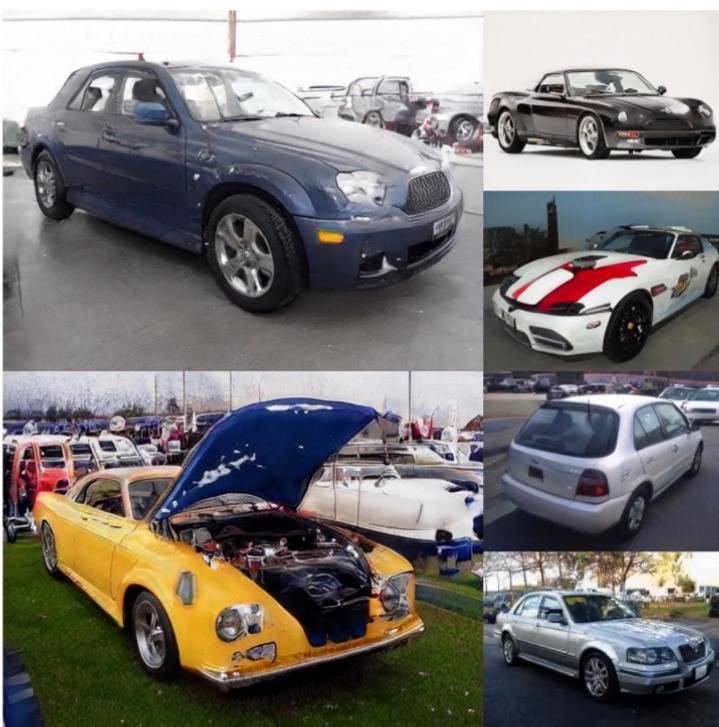
$$\text{AdaIN}(\mathbf{x}_i, \mathbf{y}) = \mathbf{y}_{s,i} \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{b,i}$$

AdaIN: [Huang and Belongie, ICCV 2017]
normalization followed by a modulation

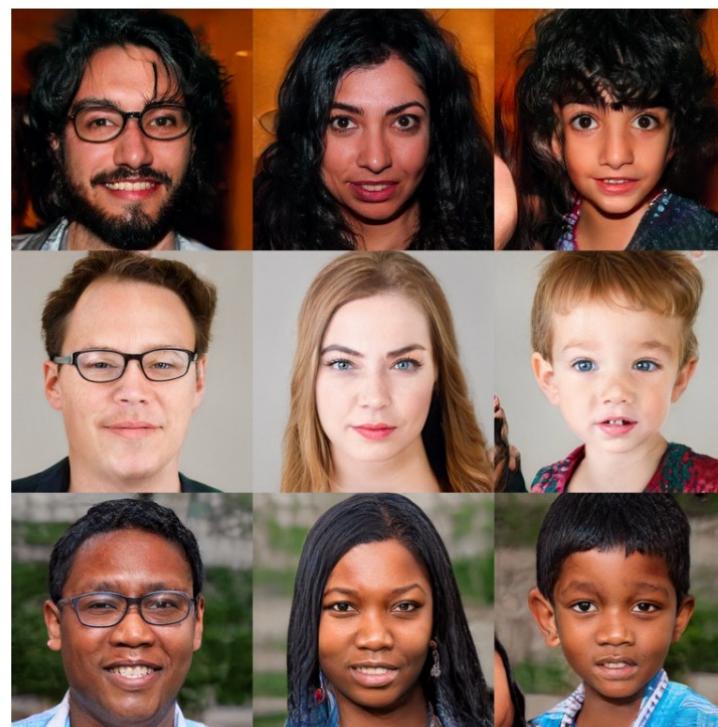
Figure 1. While a traditional generator [30] feeds the latent code through the input layer only, we first map the input to an intermediate latent space \mathcal{W} , which then controls the generator through adaptive instance normalization (AdaIN) at each convolution layer. Gaussian noise is added after each convolution, before evaluating the nonlinearity. Here “A” stands for a learned affine transform, and “B” applies learned per-channel scaling factors to the noise input. The mapping network f consists of 8 layers and the synthesis network g consists of 18 layers — two for each resolution ($4^2 - 1024^2$). The output of the last layer is converted to RGB using a separate 1×1 convolution, similar to Karras et al. [30]. Our generator has a total of 26.2M trainable parameters, compared to 23.1M in the traditional generator.

GAN improvements – High Resolution StyleGAN, Karras et al CVPR 2019

512 x 384 cars



1024 x 1024 faces



Karras et al, "A Style-Based Generator Architecture for Generative Adversarial Networks", CVPR 2019

[Images](#) are licensed under [CC BY-NC 4.0](#)

StyleGAN Interpolations

Karras et al CVPR 2019



StyleGAN artifacts



Figure 1. Instance normalization causes water droplet -like artifacts in StyleGAN images. These are not always obvious in the generated images, but if we look at the activations inside the generator network, the problem is always there, in all feature maps starting from the 64x64 resolution. It is a systemic problem that plagues all StyleGAN images.

StyleGAN2, Karras et al CVPR 2020

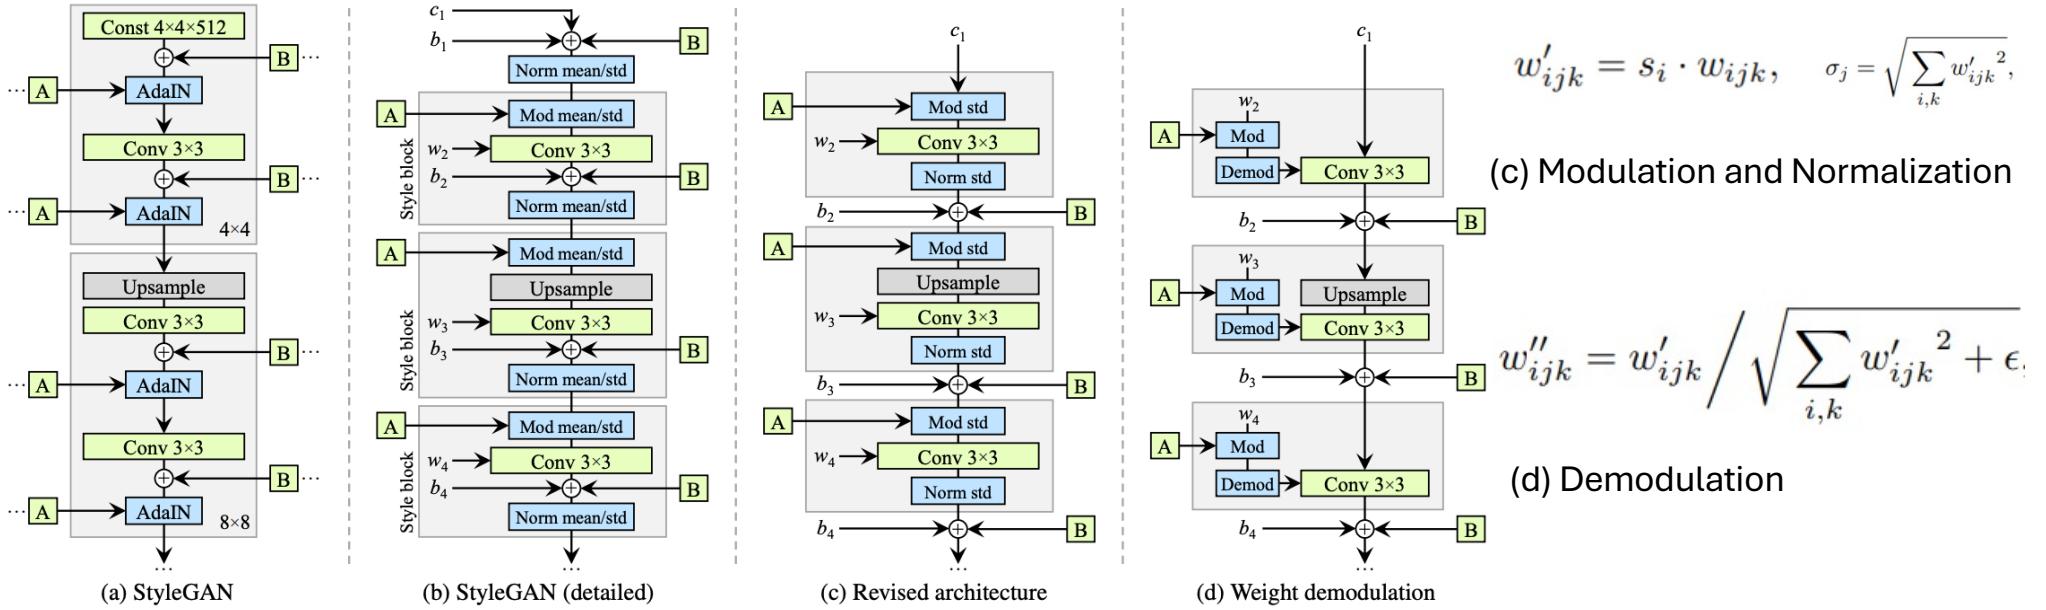


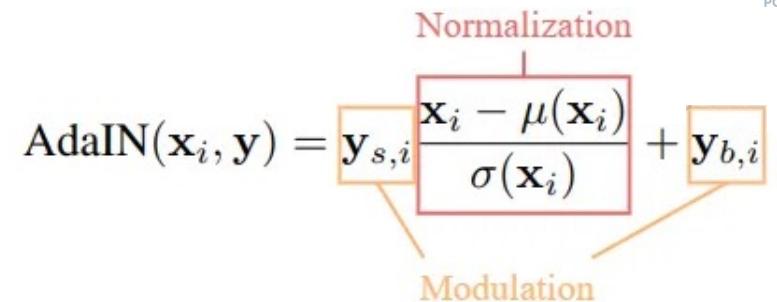
Figure 2. We redesign the architecture of the StyleGAN synthesis network. (a) The original StyleGAN, where \boxed{A} denotes a learned affine transform from \mathcal{W} that produces a style and \boxed{B} is a noise broadcast operation. (b) The same diagram with full detail. Here we have broken the AdaIN to explicit normalization followed by modulation, both operating on the mean and standard deviation per feature map. We have also annotated the learned weights (w), biases (b), and constant input (c), and redrawn the gray boxes so that one style is active per box. The activation function (leaky ReLU) is always applied right after adding the bias. (c) We make several changes to the original architecture that are justified in the main text. We remove some redundant operations at the beginning, move the addition of b and \boxed{B} to be outside active area of a style, and adjust only the standard deviation per feature map. (d) The revised architecture enables us to replace instance normalization with a “demodulation” operation, which we apply to the weights associated with each convolution layer.

StyleGAN2, Karras et al CVPR 2020

$$\text{AdaIN}(\mathbf{x}_i, \mathbf{y}) = \mathbf{y}_{s,i} \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{b,i}$$

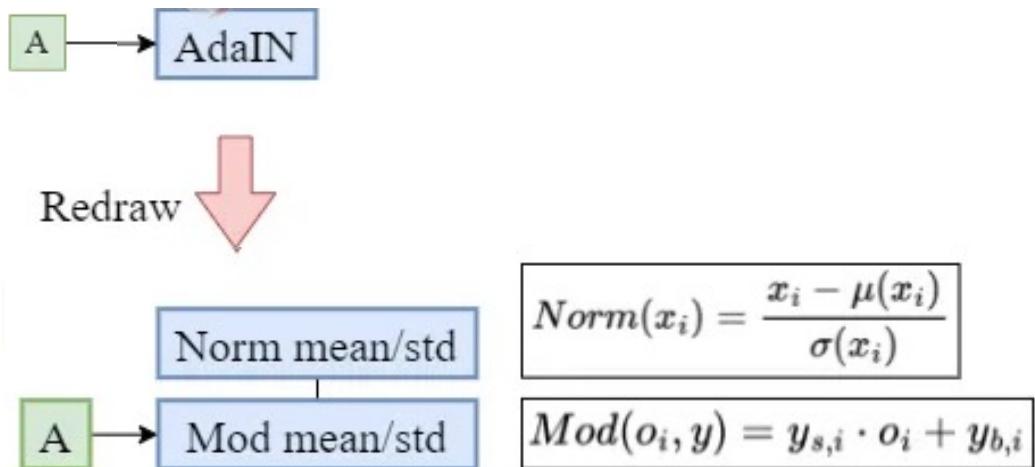
Normalization

Modulation



The diagram illustrates the AdaIN equation. It shows the formula $\text{AdaIN}(\mathbf{x}_i, \mathbf{y}) = \mathbf{y}_{s,i} \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{b,i}$. A red box encloses the term $\frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)}$, which is labeled 'Normalization'. Two orange boxes enclose the terms $\mathbf{y}_{s,i}$ and $\mathbf{y}_{b,i}$, which are labeled 'Modulation'.

StyleGAN2, Karras et al CVPR 2020



Normalization

$$\text{AdaIN}(\mathbf{x}_i, \mathbf{y}) = \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{s,i}$$

Modulation

$$\text{Mod}(o_i, \mathbf{y}) = y_{s,i} \cdot o_i + y_{b,i}$$



Norm std

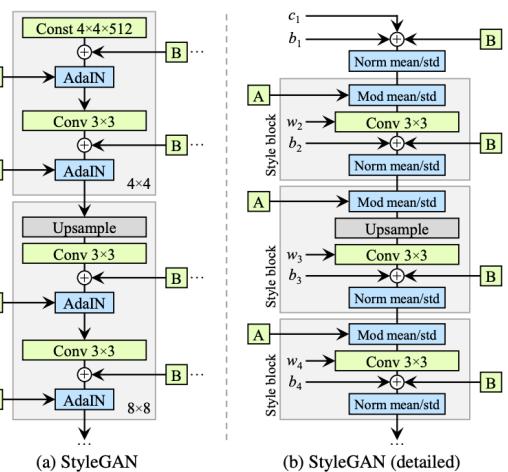
$$\text{Norm}(x_i) = \frac{x_i}{\sigma(x_i)}$$

Mod std

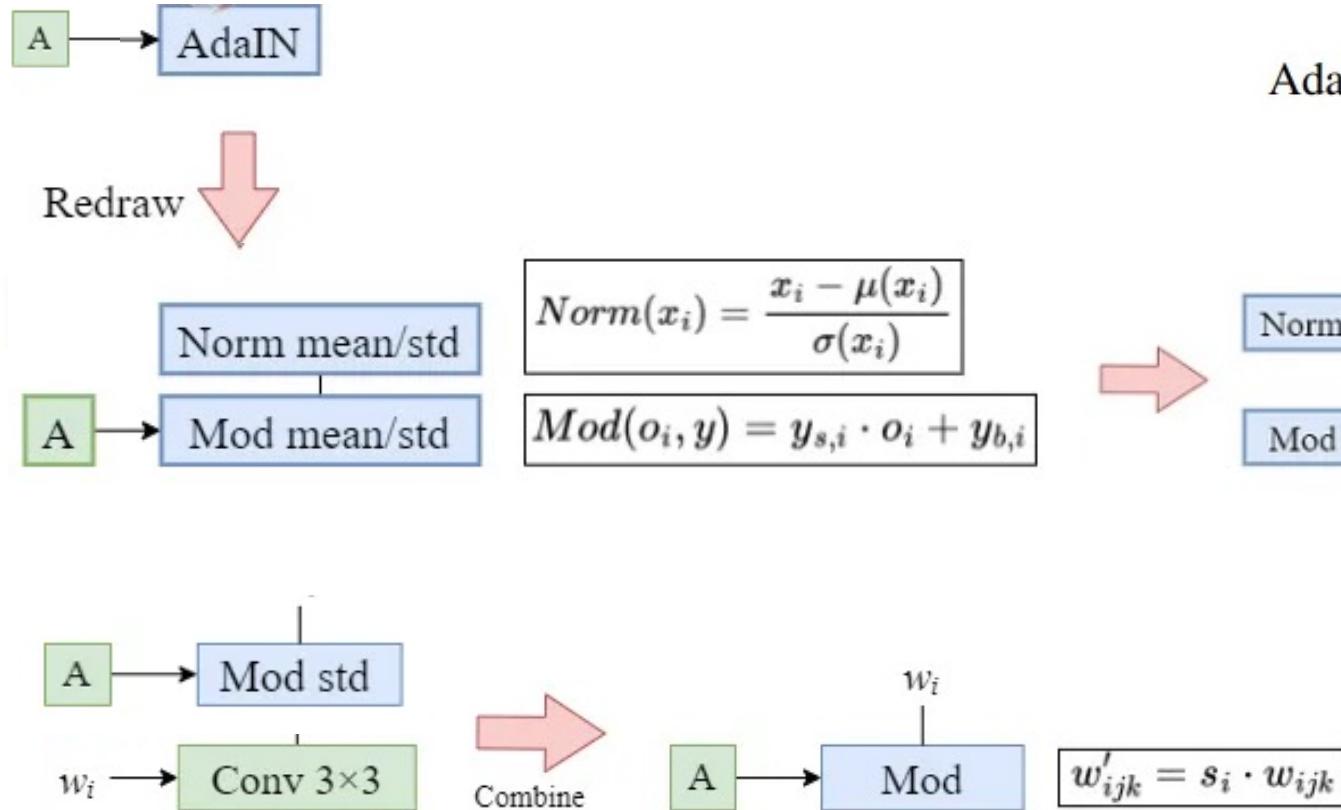
$$\text{Mod}(o_i, y) = y_{s,i} \cdot o_i$$

Vicky Kalogeiton

Lecture 1: CSC_52087_EP



StyleGAN2, Karras et al CVPR 2020



Normalization

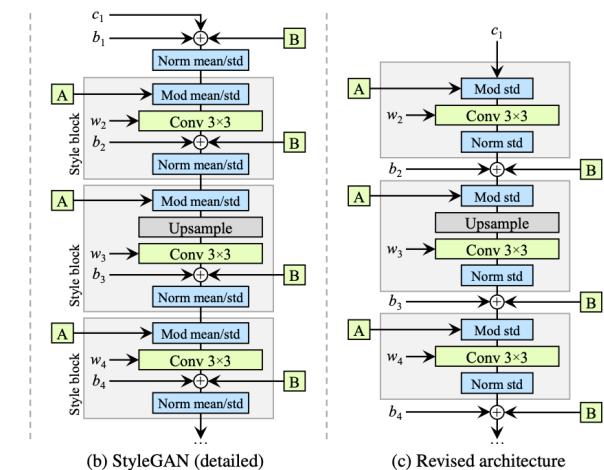
$$AdaIN(\mathbf{x}_i, \mathbf{y}) = \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{b,i}$$

Modulation

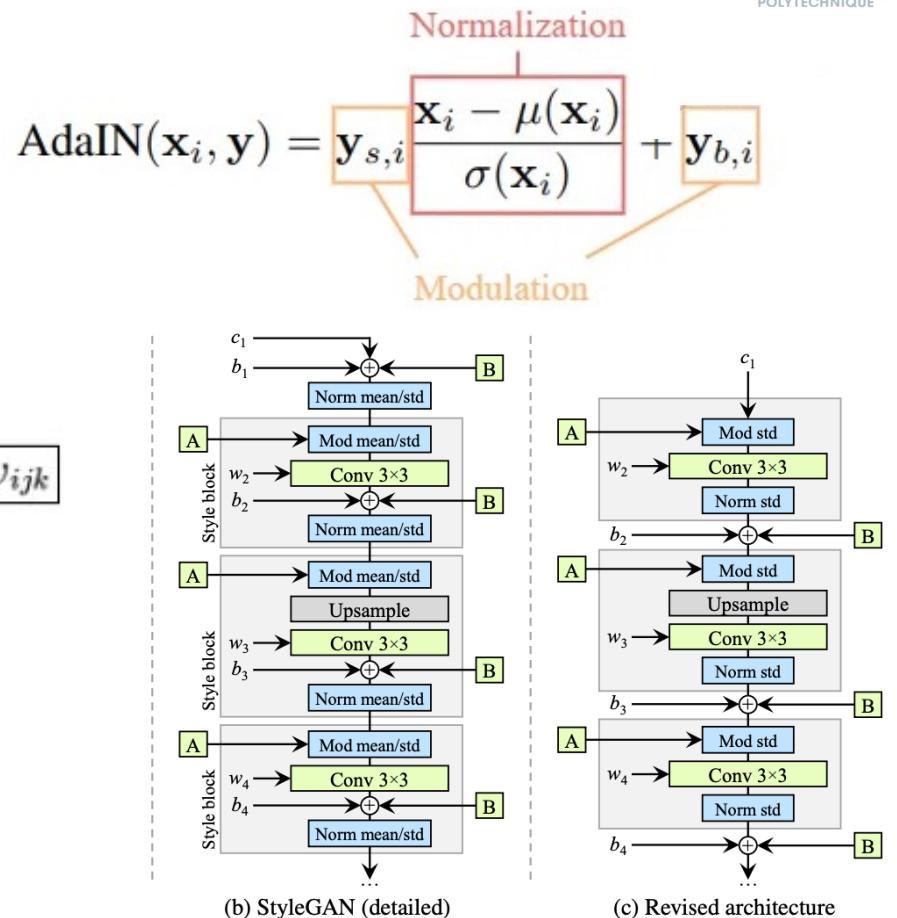
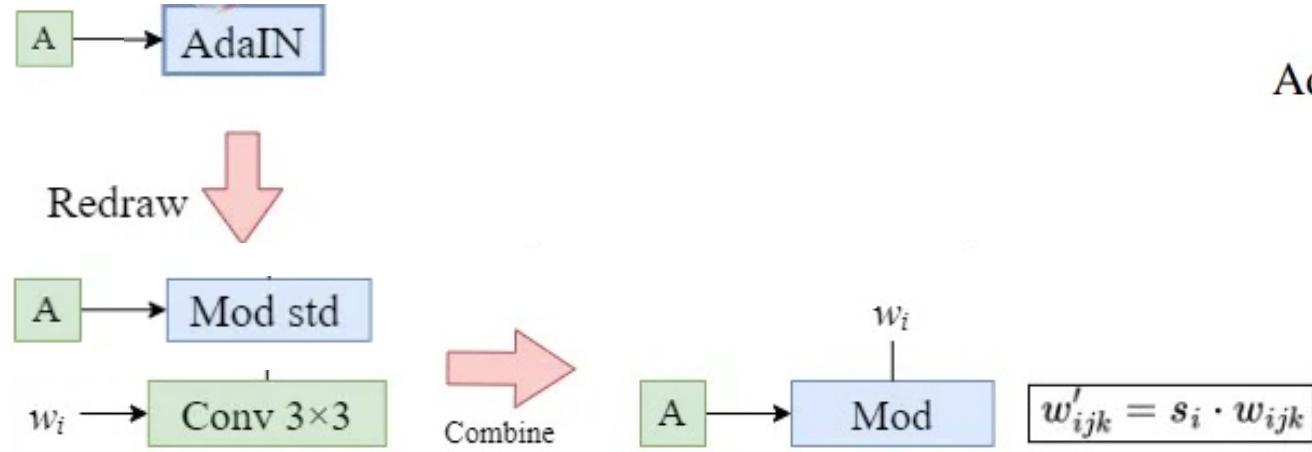
$$Mod(o_i, y) = y_{s,i} \cdot o_i + y_{b,i}$$

Vicky Kalogeiton

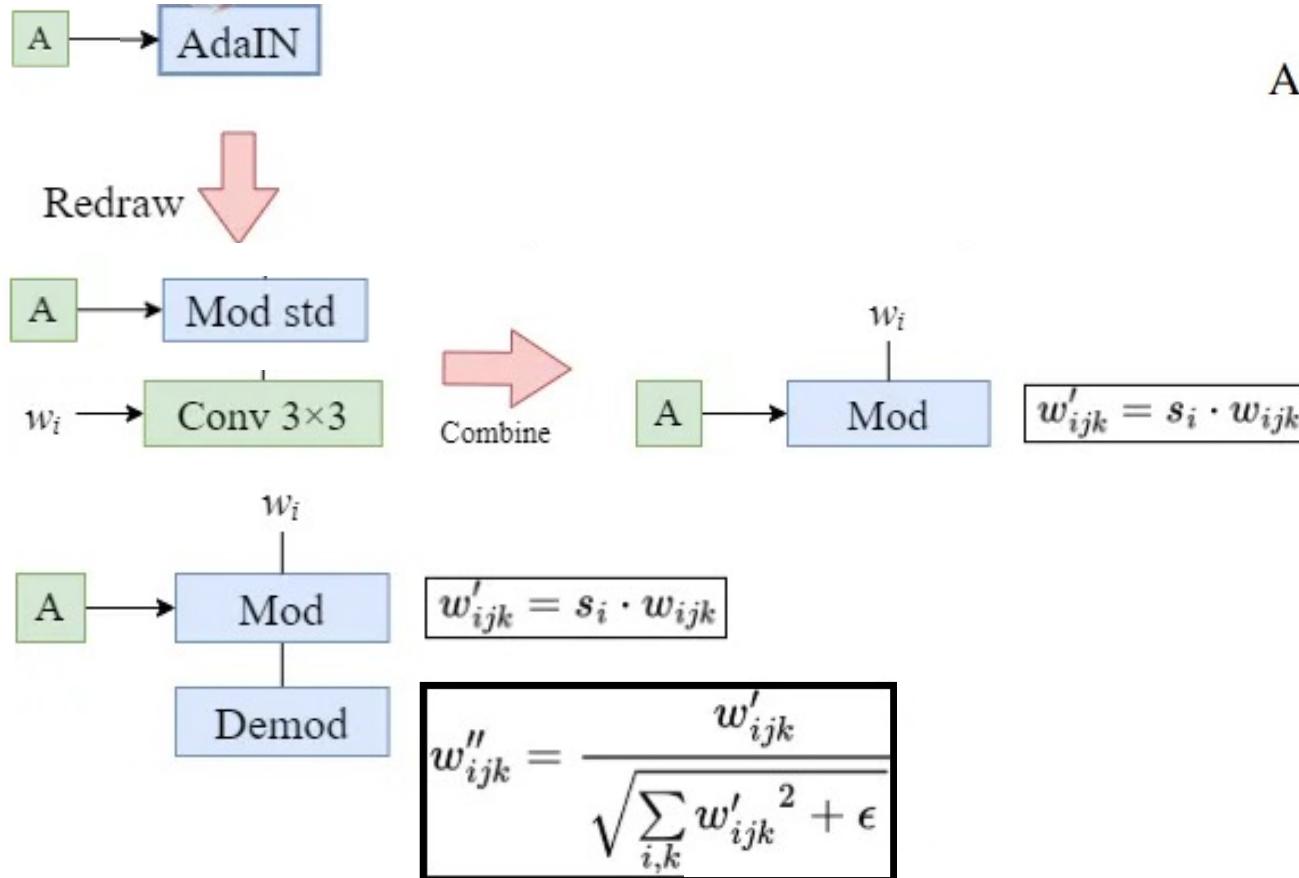
Lecture 1: CSC_52087_EP



StyleGAN2, Karras et al CVPR 2020



StyleGAN2, Karras et al CVPR 2020



StyleGAN2, Karras et al CVPR 2020

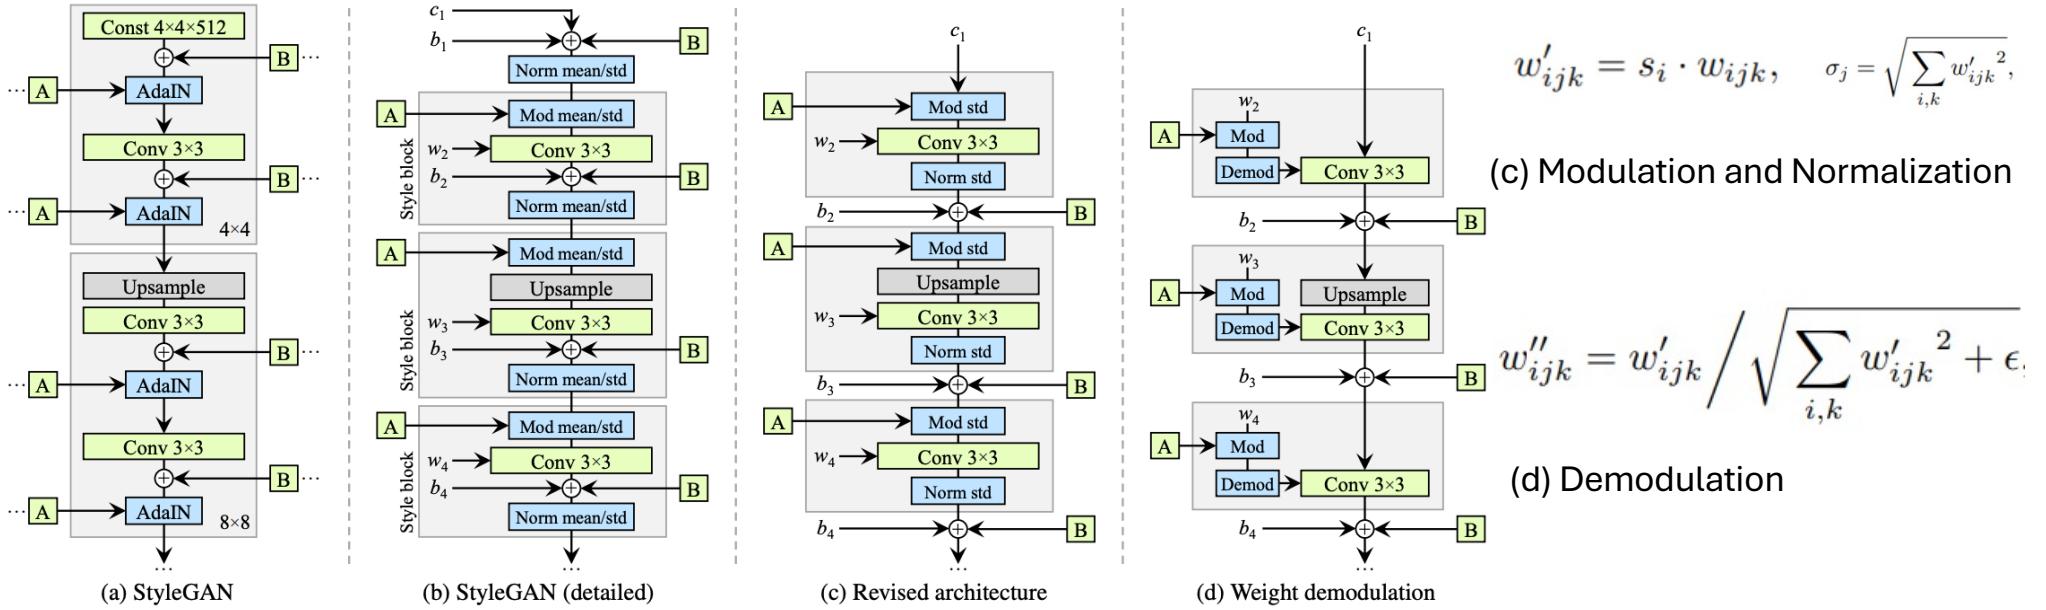


Figure 2. We redesign the architecture of the StyleGAN synthesis network. (a) The original StyleGAN, where \boxed{A} denotes a learned affine transform from \mathcal{W} that produces a style and \boxed{B} is a noise broadcast operation. (b) The same diagram with full detail. Here we have broken the AdaIN to explicit normalization followed by modulation, both operating on the mean and standard deviation per feature map. We have also annotated the learned weights (w), biases (b), and constant input (c), and redrawn the gray boxes so that one style is active per box. The activation function (leaky ReLU) is always applied right after adding the bias. (c) We make several changes to the original architecture that are justified in the main text. We remove some redundant operations at the beginning, move the addition of b and \boxed{B} to be outside active area of a style, and adjust only the standard deviation per feature map. (d) The revised architecture enables us to replace instance normalization with a “demodulation” operation, which we apply to the weights associated with each convolution layer.

StyleGAN2 - ADA, Karras et al NeurIPS 2020

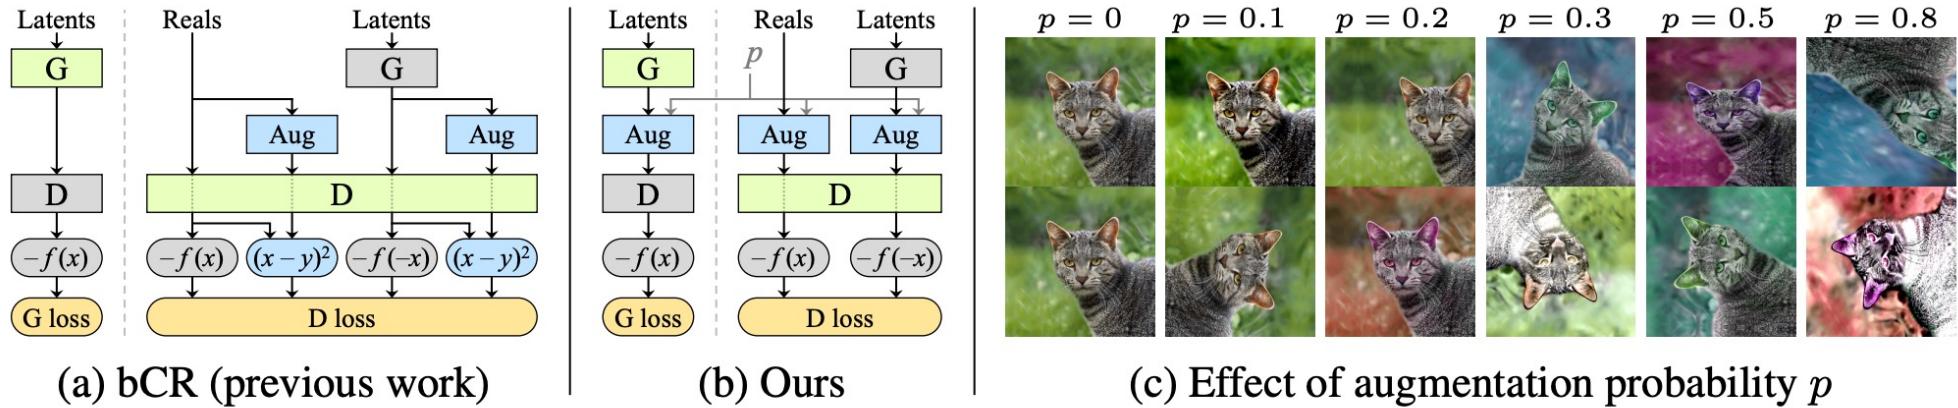


Figure 2: (a,b) Flowcharts for balanced consistency regularization (bCR) [43] and our stochastic discriminator augmentations. The blue elements highlight operations related to augmentations, while the rest implement standard GAN training with generator G and discriminator D [12]. The orange elements indicate the loss function and the green boxes mark the network being trained. We use the non-saturating logistic loss [12] $f(x) = \log(\text{sigmoid}(x))$. (c) We apply a diverse set of augmentations to every image that the discriminator sees, controlled by an augmentation probability p .

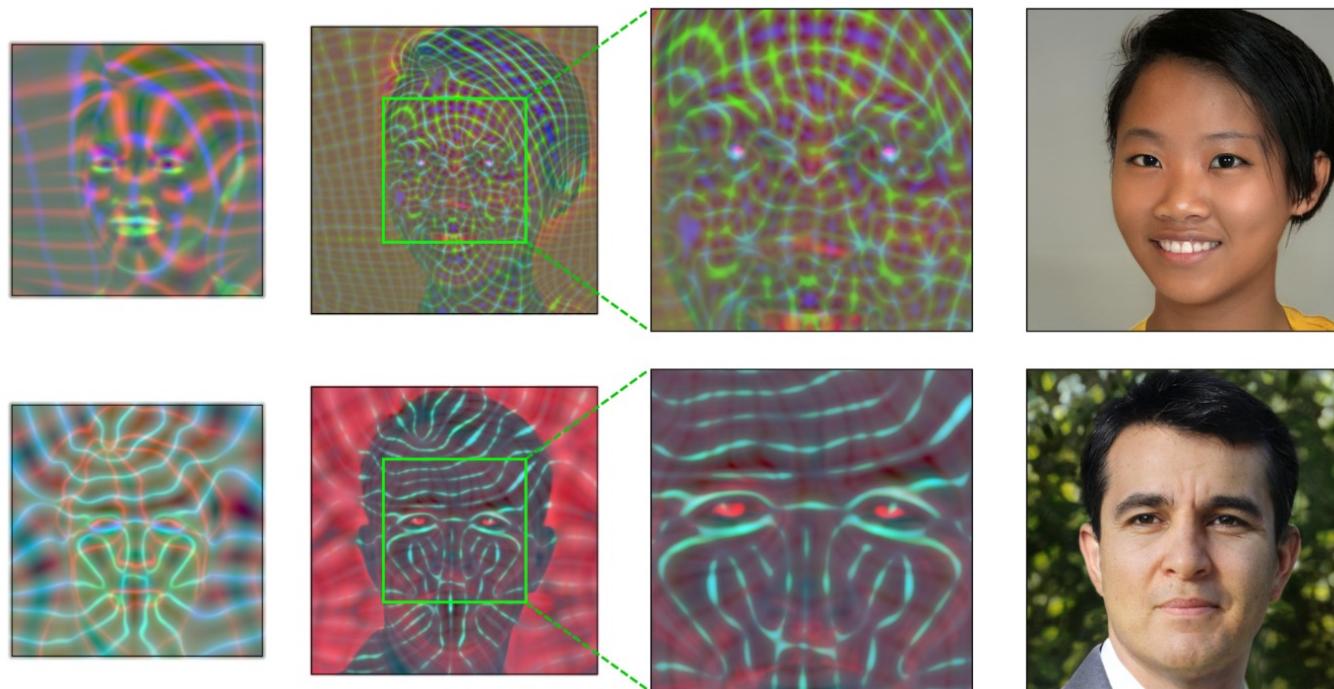
<https://github.com/NVlabs/stylegan2-ada>

GAN progress on face generation



Source: https://twitter.com/goodfellow_ian/status/1084973596236144640/
+StyleGANv2

StyleGAN3, Karras NeurIPS 2021



<https://nvlabs.github.io/stylegan3/>

Examples

- <http://www.whichfaceisreal.com>

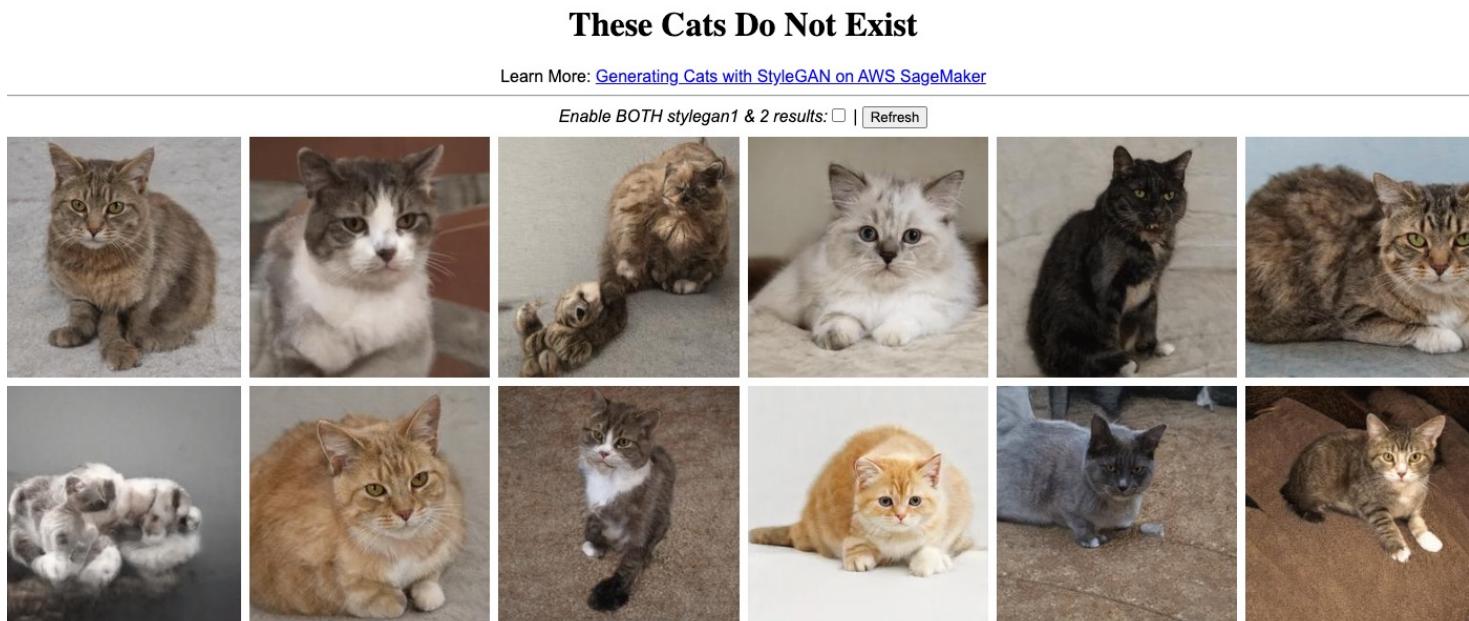
You are **correct**. The image on the left is real.

[Play again.](#)



Examples

- <https://thiscatdoesnotexist.com/>



Which face is real?

- Cool website
 - <http://www.whichfaceisreal.com>
- Others
 - <https://thisrentaldoesnotexist.com/>
 - <https://thiscatdoesnotexist.com/>
 - <https://thishorsedoesnotexist.com/>
 - <https://www.thiswaifudoesnotexist.net/>



Metrics

Data augmentation in GANs

- Data augmentation is helpful during GAN training, if done to **real AND fake** images
 - Need differentiable augmentations
 - <https://github.com/mit-han-lab/data-efficient-gans>



[Zhao Zhengli et al. "Image Augmentations for GAN Training." arXiv 2006.02595, 2020]

FID - Frechet Inception Distance

- A metric attempting to quantitatively evaluate the quality of generated images
- Main Idea
 - Use features learned for classification, and see how similar they are for real and generated images
- How to:
 - Take the 2048 dimensional output of the global pooling layer in Inception v3
 - Compute the mean and covariance matrix of the features for real and fake images
 - Calculate the Frechet distance between the real and generated images

$$\|\mathbf{m} - \mathbf{m}_w\|_2^2 + \text{Tr}(\mathbf{C} + \mathbf{C}_w - 2(\mathbf{C}\mathbf{C}_w)^{1/2})$$

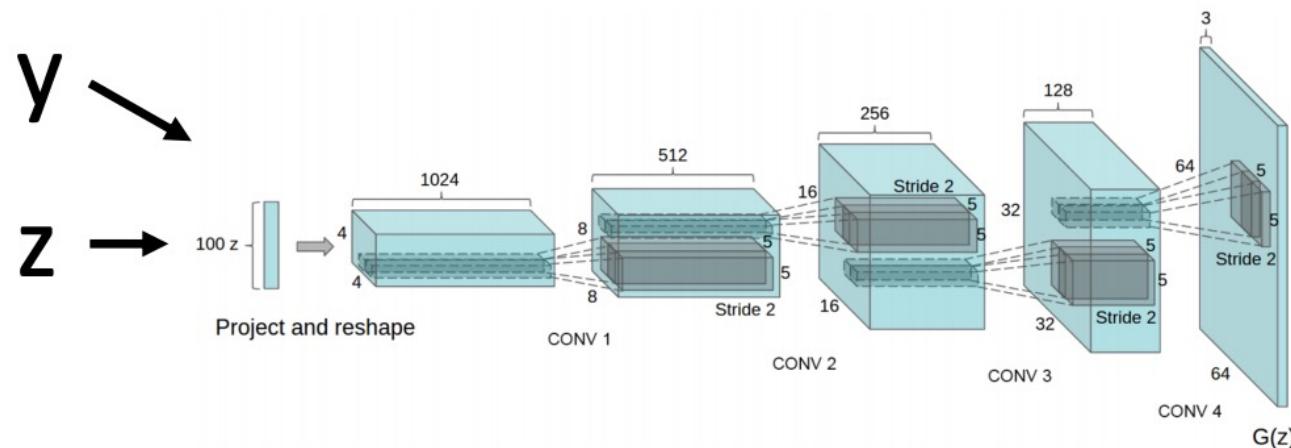
[Merceil. Martin, de al. "Gans trained by a two time-scale update rule converge to a local nash equilibrium". In NeurIPS 2017]



Conditional GANs

Conditional GANs

- Conditional GAN: learn $p(x | y)$ instead of $p(x)$
 - Make generator and discriminator both take label y as an additional input



Recall: Batch Normalization

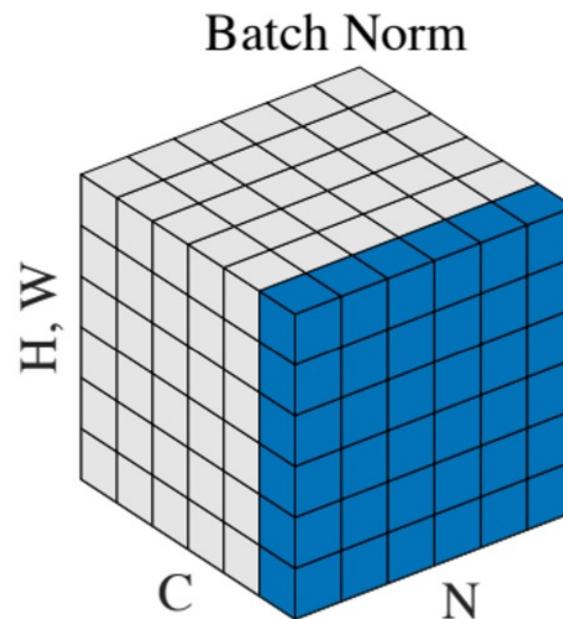
- A way to keep the weights in the gaussian range
- Idea: “Normalize” the outputs of a layer so they have zero mean and unit variance
 - Helps reduce “internal covariate shift”
 - improves optimization
- Before non-linearity – run batch normalization empirically on the specific batch

$$\hat{x} = \frac{x - E[x]}{\sqrt{Var[x]}}$$

This is a differentiable function, so we can use it as an operator in our networks and backprop through it!

Batch Normalization: Accelerating Deep Network Training
by
[Reducing Internal Covariate Shift \[Ioffe and Szegedy 2015\]](#)

Recall: Batch Normalization



[Wu and He, “Group Normalization”, ECCV
2018]

Vicky Kalogeiton

Lecture 1: CSC_52087_EP

107

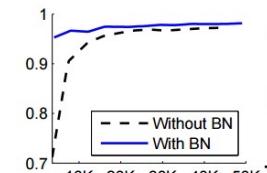
Recall: Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;
 Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

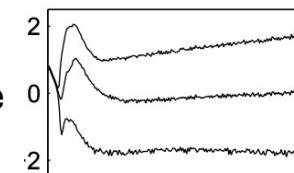
$$\begin{aligned}\mu_{\mathcal{B}} &\leftarrow \frac{1}{m} \sum_{i=1}^m x_i && // \text{mini-batch mean} \\ \sigma_{\mathcal{B}}^2 &\leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 && // \text{mini-batch variance} \\ \hat{x}_i &\leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} && // \text{normalize} \\ y_i &\leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) && // \text{scale and shift}\end{aligned}$$

Test accuracy

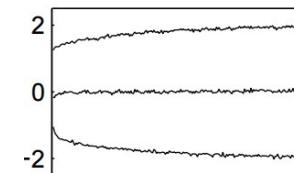


(a)

Responses of different layers over iterations



(b) Without BN



(c) With BN

Batch Normalization: Accelerating Deep Network Training
 by
 Reducing Internal Covariate Shift [[Ioffe and Szegedy 2015](#)]

Conditional GANs

Batch Normalization

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$



Learn a separate scale and shift for each different label y

Conditional Batch Normalization

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

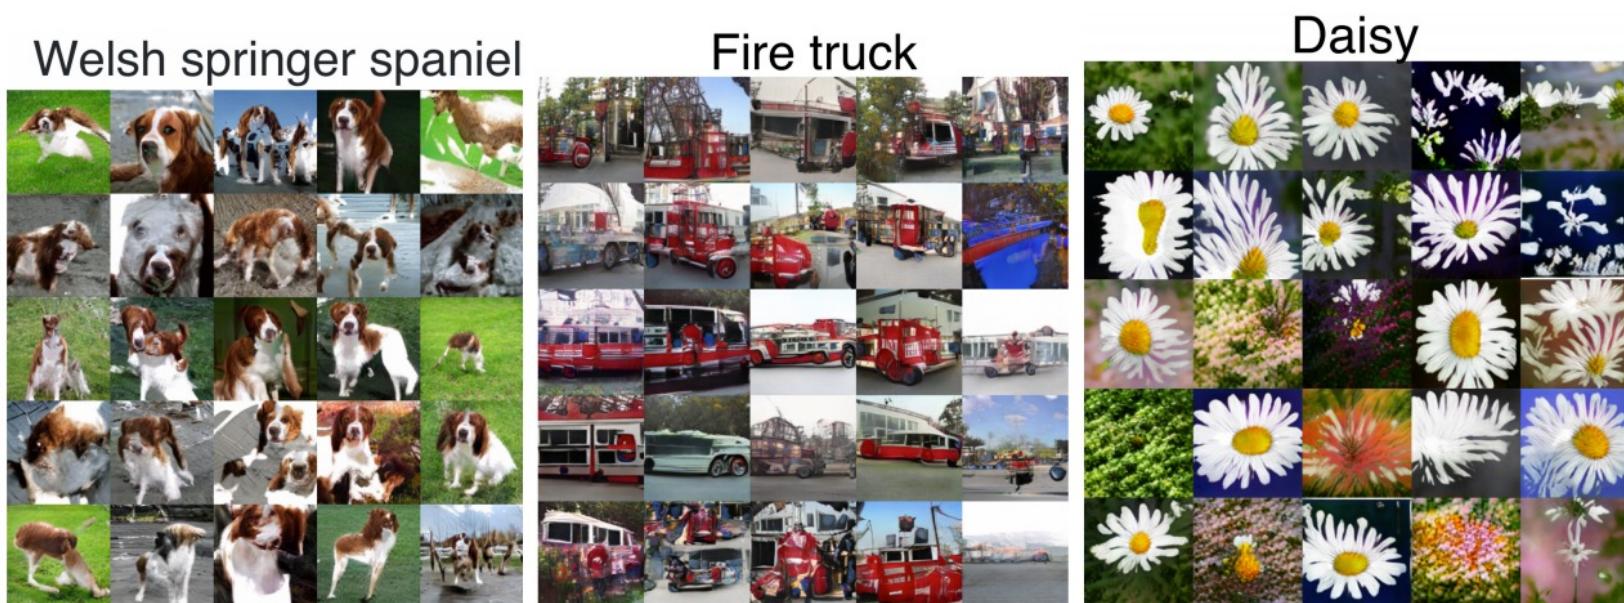
$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

$$y_{i,j} = \gamma_j^y \hat{x}_{i,j} + \beta_j^y$$

Dumoulin et al, “A learned representation for artistic style”, ICLR 2017

Conditional GANs



Miyato et al, “Spectral Normalization for Generative Adversarial Networks”, ICLR 2018

Conditional GANs

BigGAN



Brock et al, "Large Scale GAN Training for High Fidelity Natural Image Synthesis", ICLR 2019

512x512 images on ImageNet

Conditional GANs

More than labels: Text to Image



This bird is red and brown in color, with a stubby beak



The bird is short and stubby with yellow on its body



A bird with a medium orange bill white body gray wings and webbed feet



This small black bird has a short, slightly curved bill and long legs



A picture of a very clean living room



A group of people on skis stand in the snow



Eggs fruit candy nuts and meat served on white dish



A street sign on a stoplight pole in the middle of a day



Zhang et al, "StackGAN++: Realistic Image Synthesis with Stacked Generative Adversarial Networks.", TPAMI 2018

Zhang et al, "StackGAN: Text to Photo-realistic Image Synthesis with Stacked Generative Adversarial Networks.", ICCV 2017

Reed et al, "Generative Adversarial Text-to-Image Synthesis", ICML 2016



Image Translation

Image translation: Pix2Pix

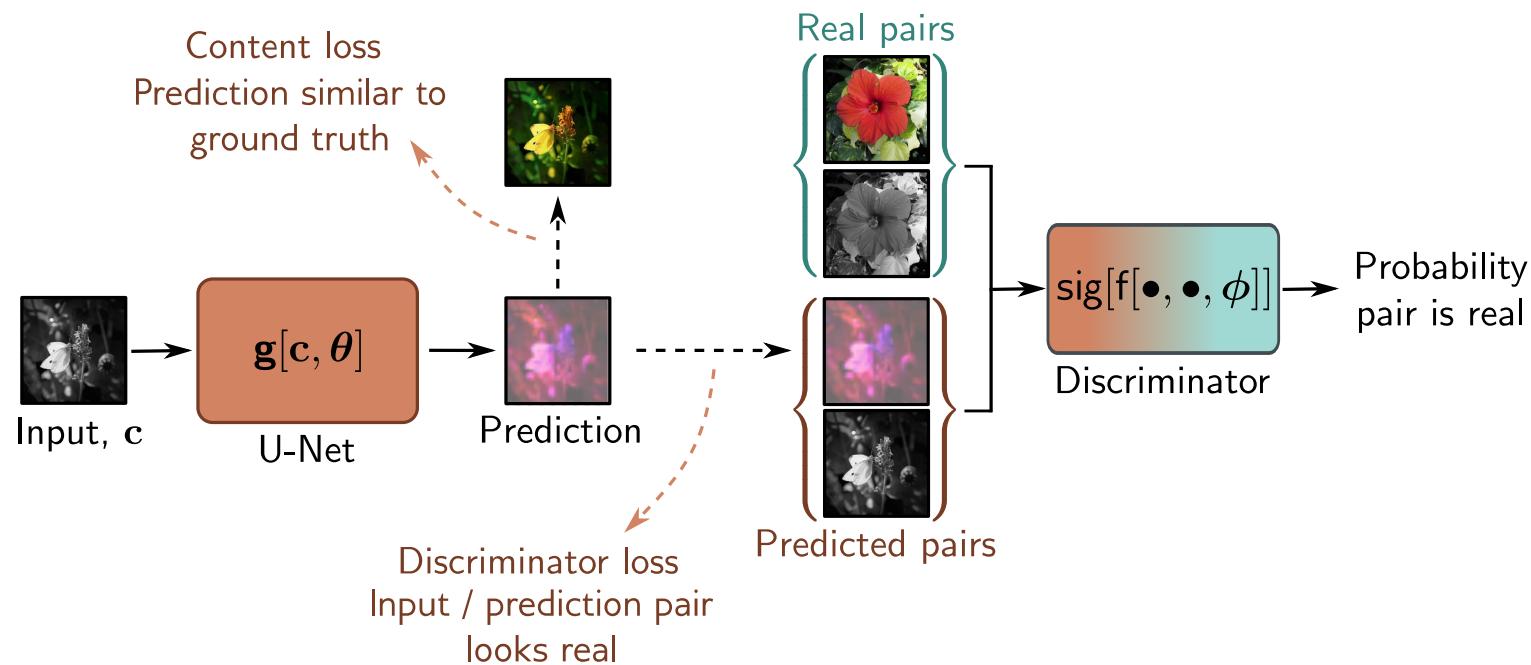


Image translation: Pix2Pix

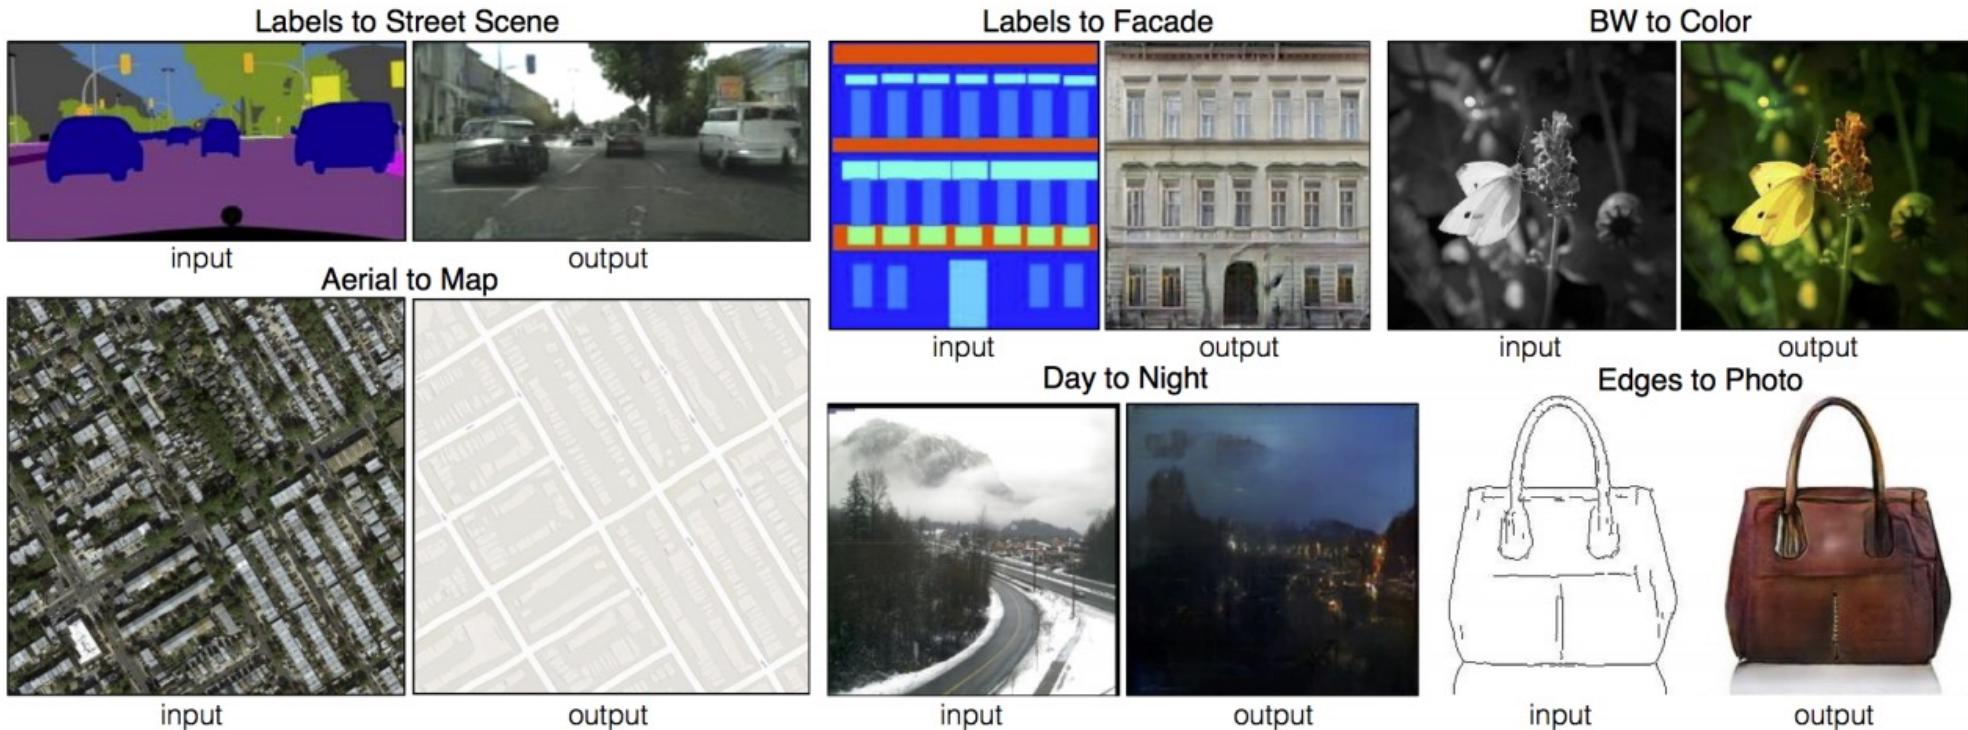


Image translation: SRGAN

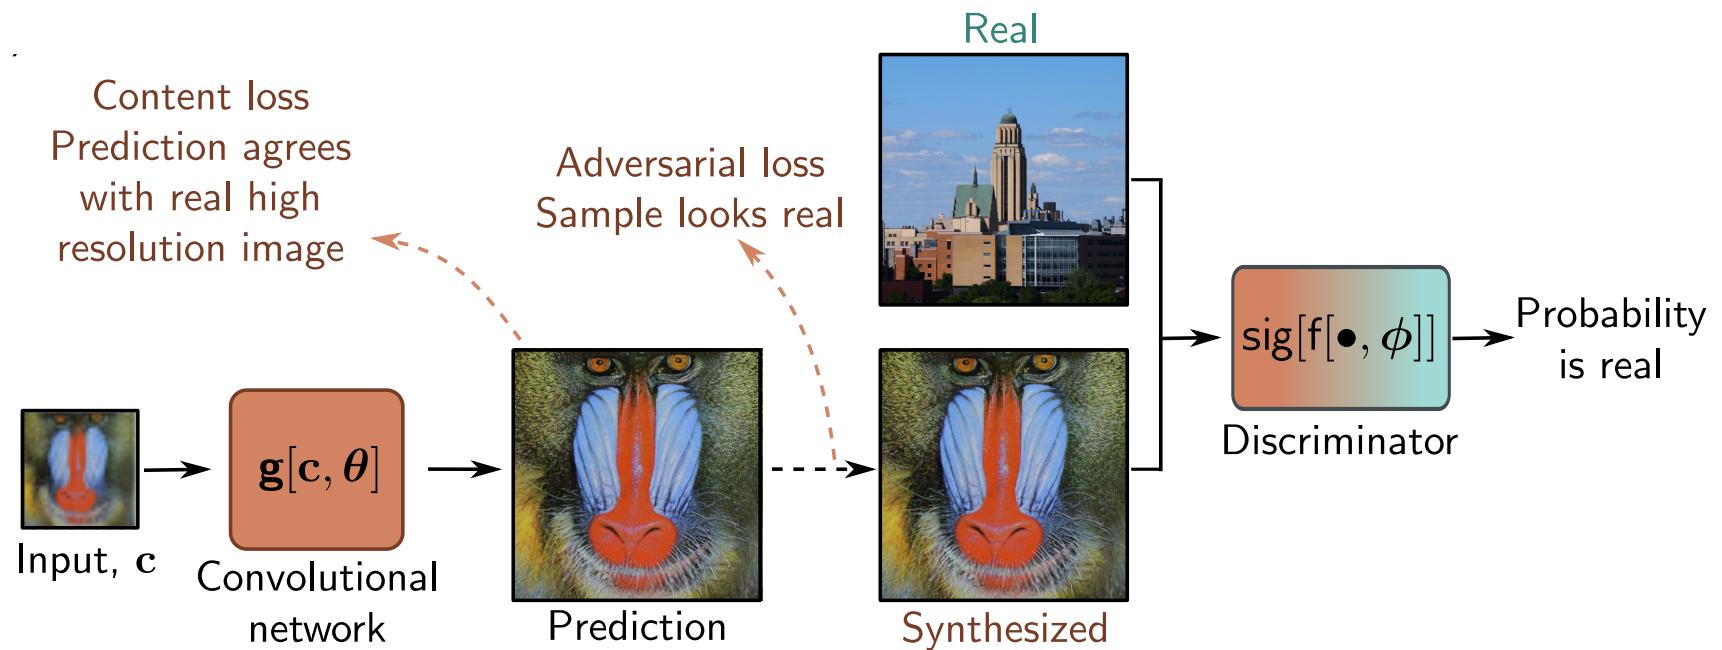
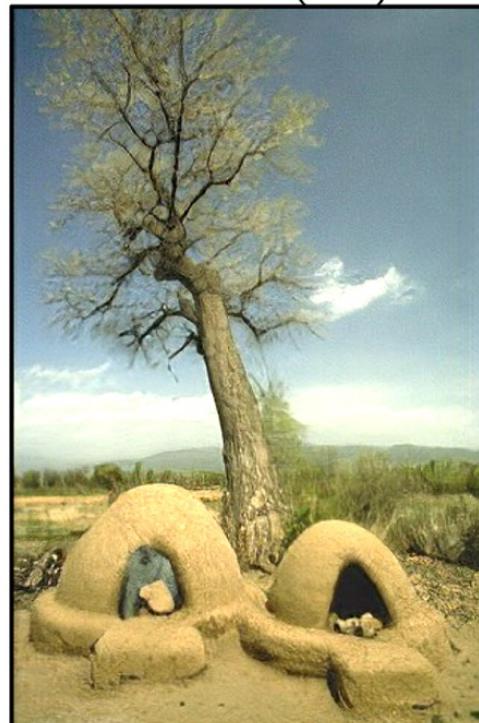


Image translation: SRGAN

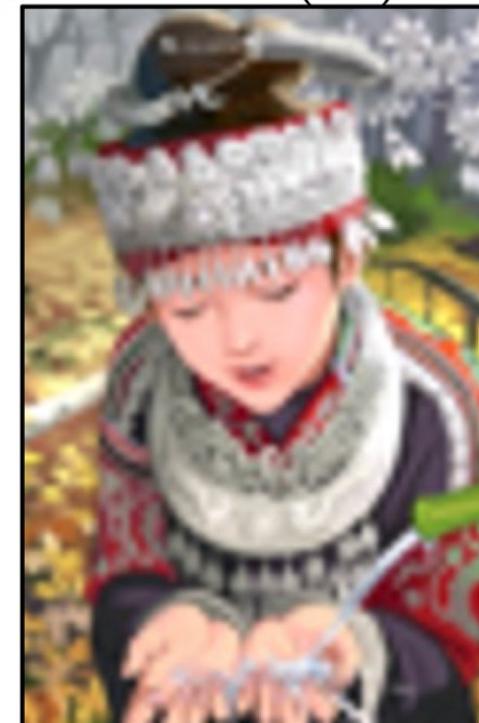
b) Bicubic (4 \times)



c) SRGAN (4 \times)



d) Bicubic (4 \times)



e) SRGAN (4 \times)

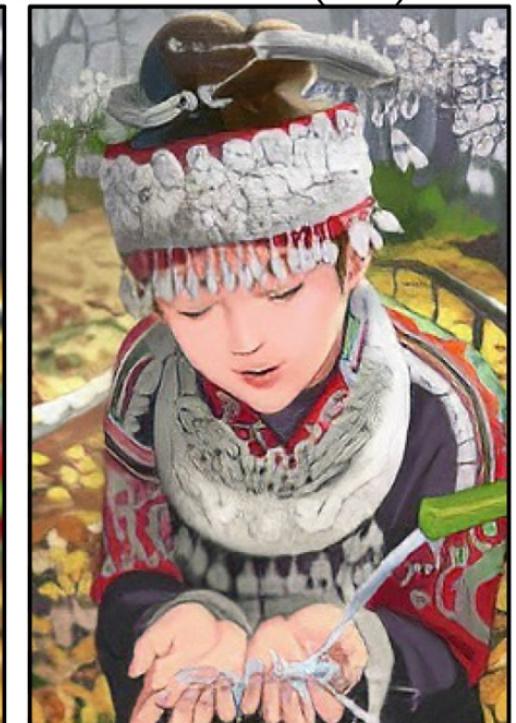


Image translation: CycleGAN

a)

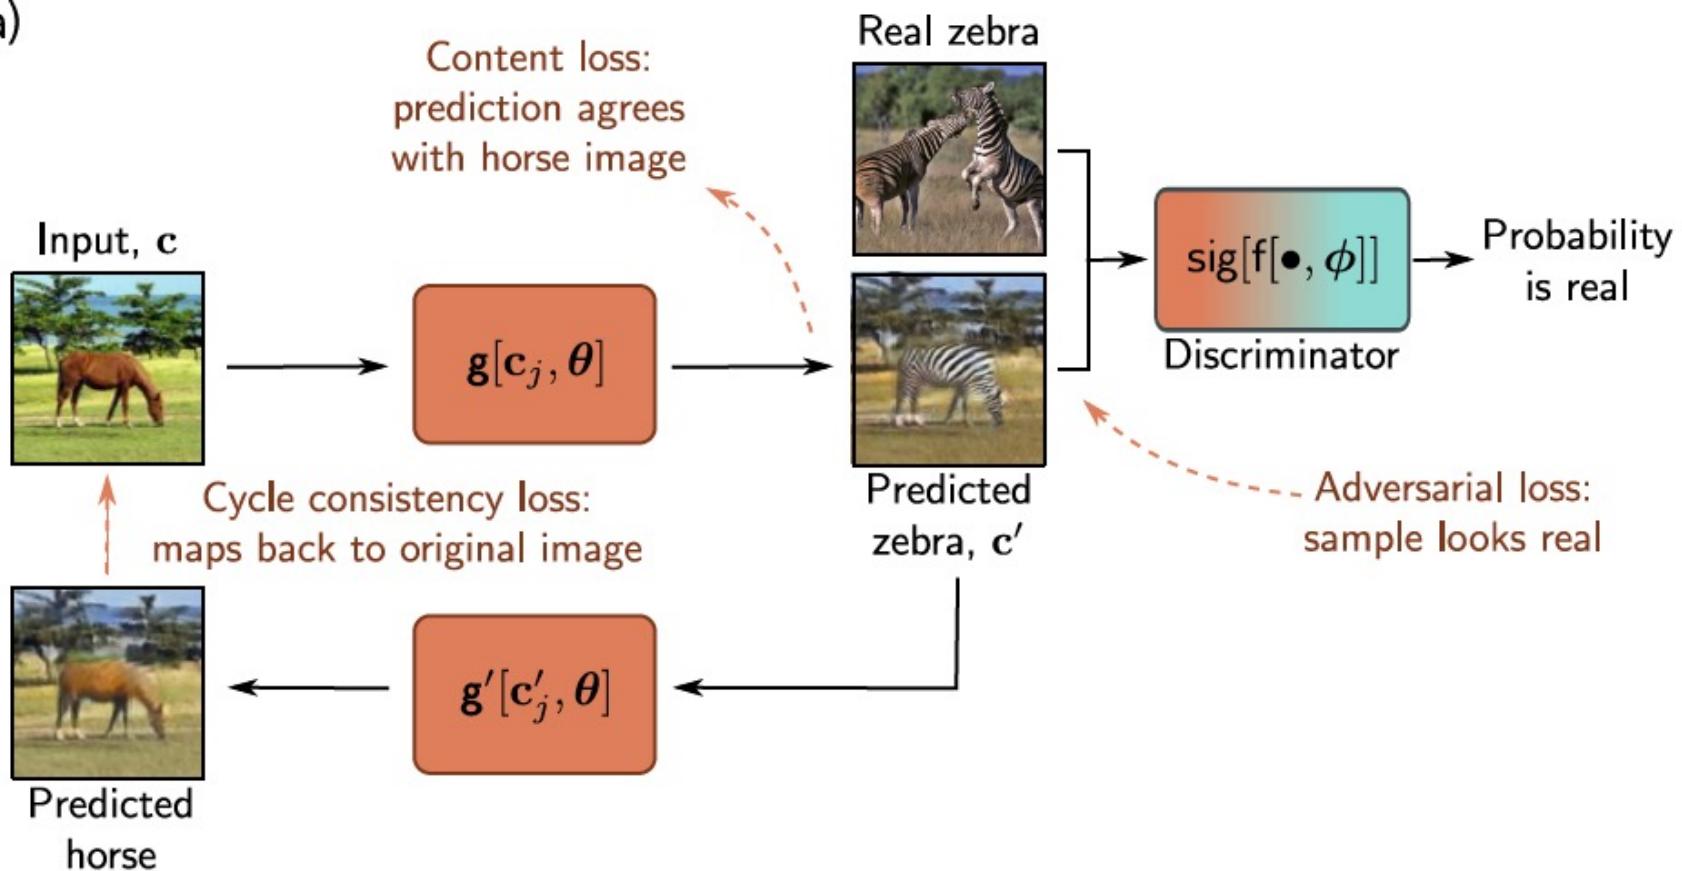


Image translation: CycleGAN

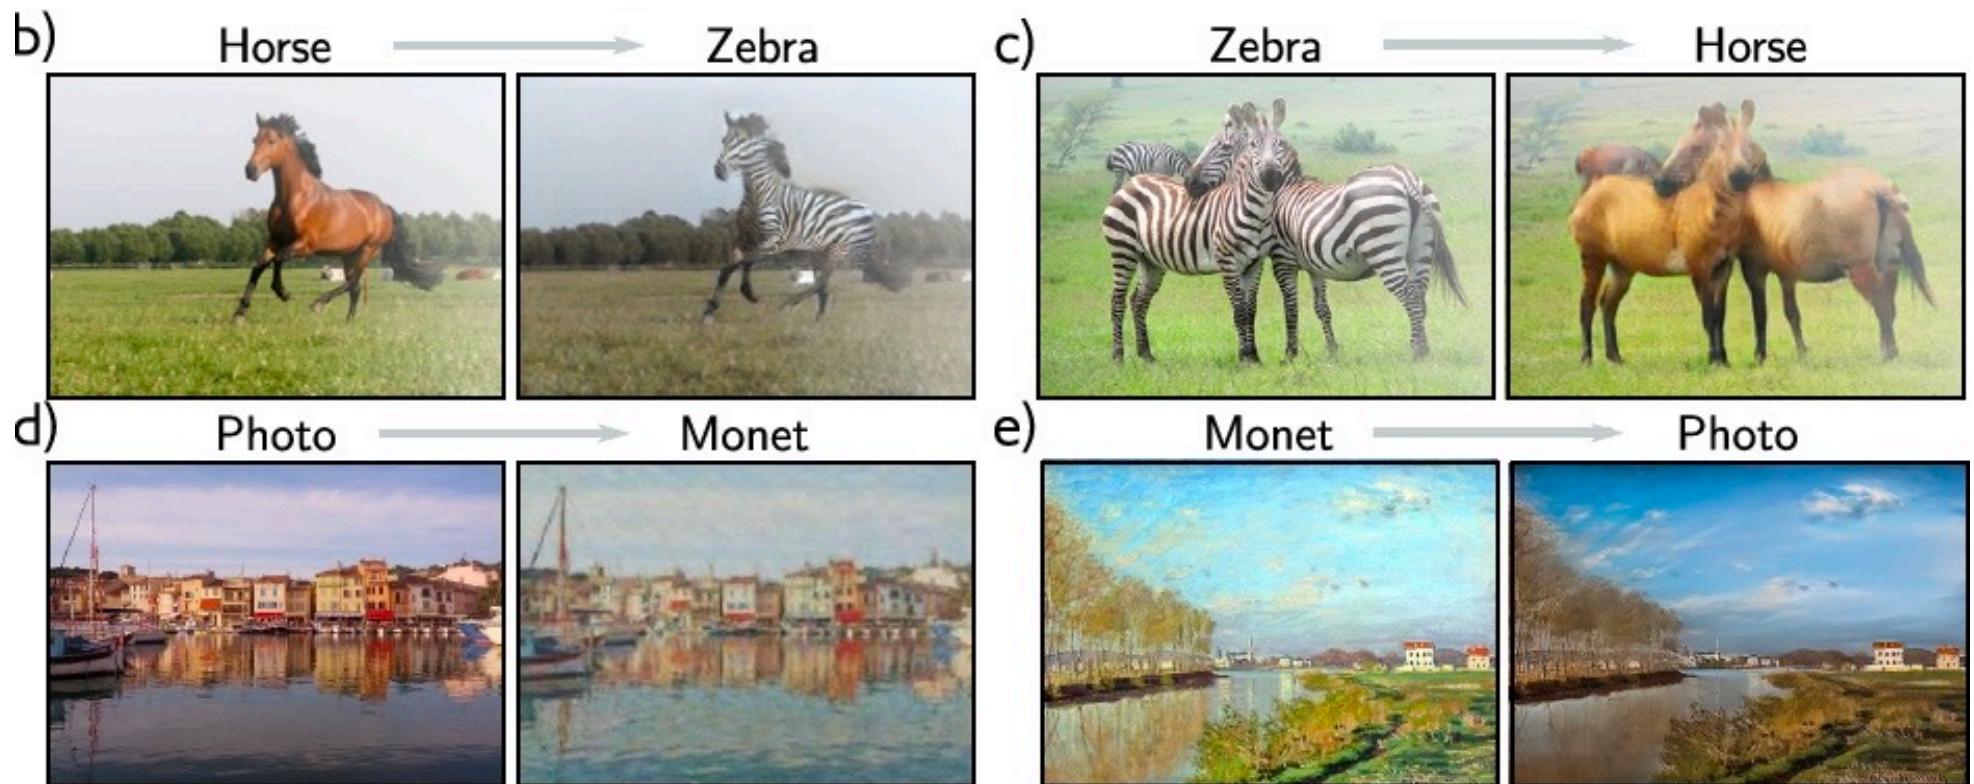
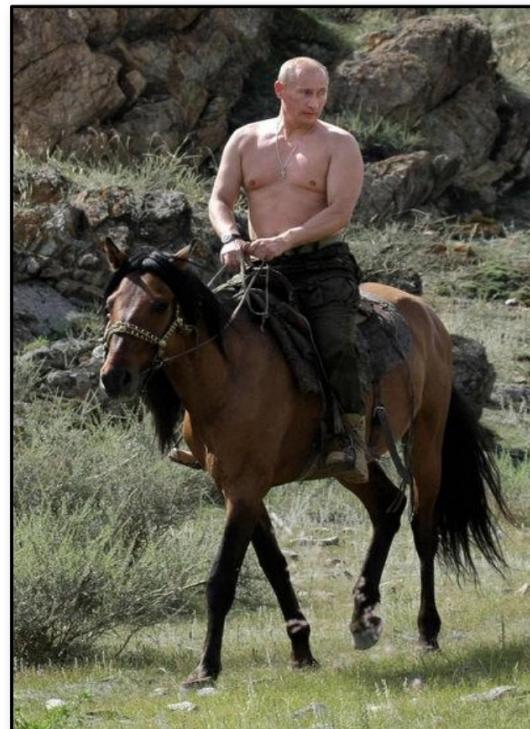


Image translation: CycleGAN

Input



Output



horse → zebra

Conditional GANs

Unpaired Image-to-Image Translation: CycleGAN



CycleGAN

Jun-Yan Zhu, Taesung Park et. al. ICCV 2017

<https://www.youtube.com/watch?v=9reHvktoWLY>

Vicky Kalogeiton

Lecture 1: CSC_52087_EP

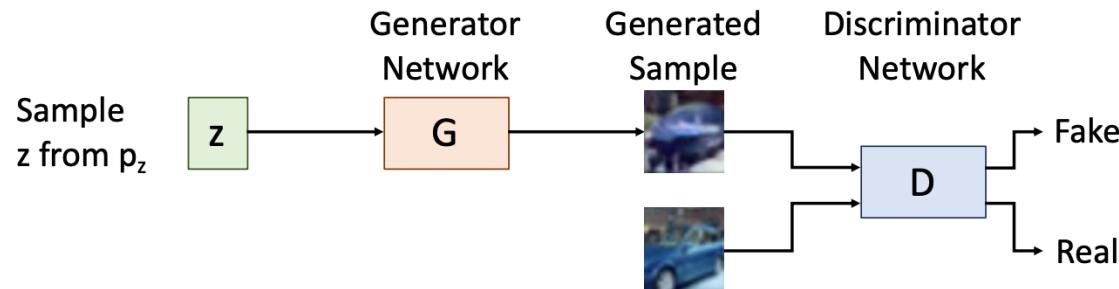
122

GAN summary

Jointly train two networks:

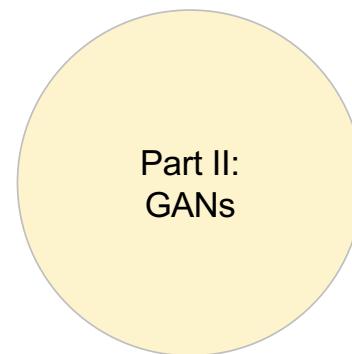
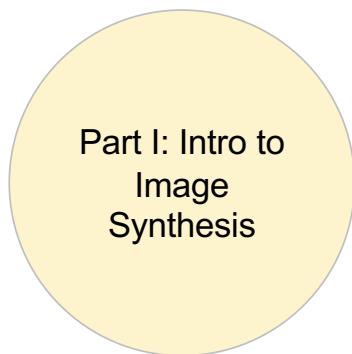
Discriminator: Classify data as real or fake

Generator: Generate data that fools the discriminator



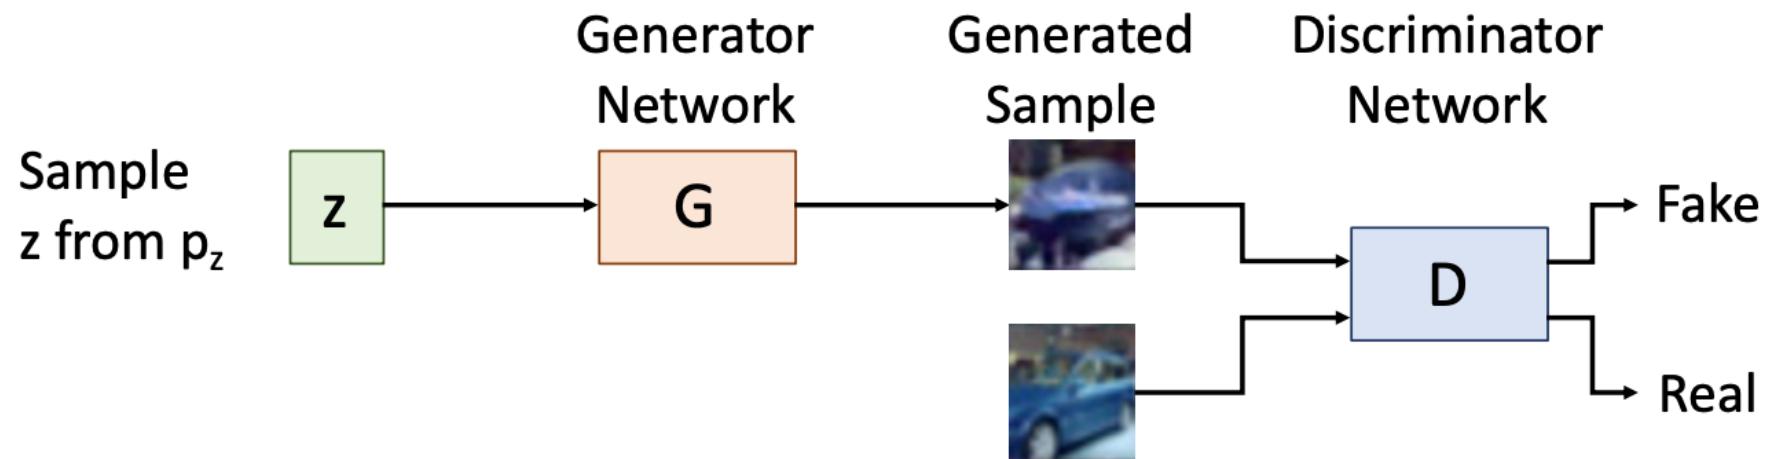
Under some assumptions, generator converges to true data distribution
Many applications! Very active area of research!

Today's lecture



Slides adapted from many resources:
[Fei-Fei Li & Andrej Karpathy & Justin Johnson & Ross Girshick & VGG]

Latent space of a GAN



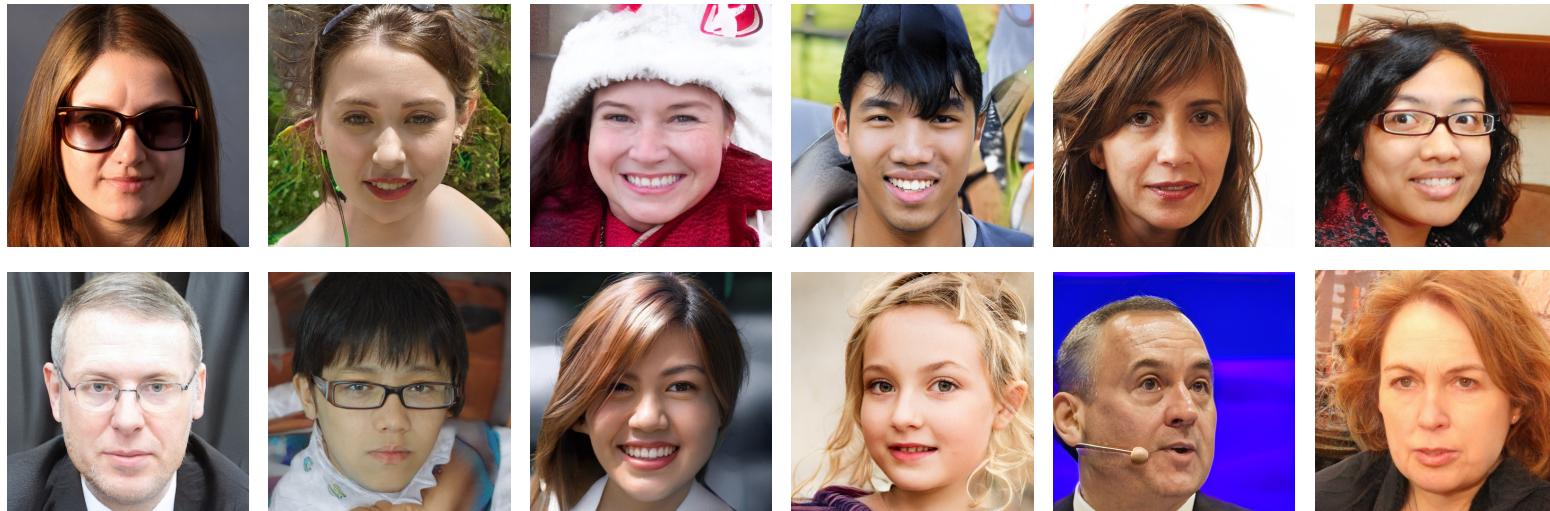
Part III: GAN editing

- Interpolation in latent space
- GAN Math
- GAN inversion
- Learn and apply latent directions
- CLIP + StyleGAN

Part III: GAN editing

- **Interpolation in latent space**
- GAN Math
- GAN inversion
- Learn and apply latent directions
- CLIP + StyleGAN

Generate random images



<https://github.com/NVlabs/stylegan2-ada>

```
python generate.py --outdir=out --trunc=1 --seeds=666-700 --
network=https://nvlabs-fi-cdn.nvidia.com/stylegan2-ada-pytorch/pretrained/ffhq.pkl
```

GAN Interpolation

Interpolating
between
points in
latent z
space



GAN Interpolation

Interpolating
between
points in
latent z
space



[Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks.", ICLR 2016]

Latent space of a GAN



z_1

$a^*z_1 + (1-a)^*z_2$

z_2

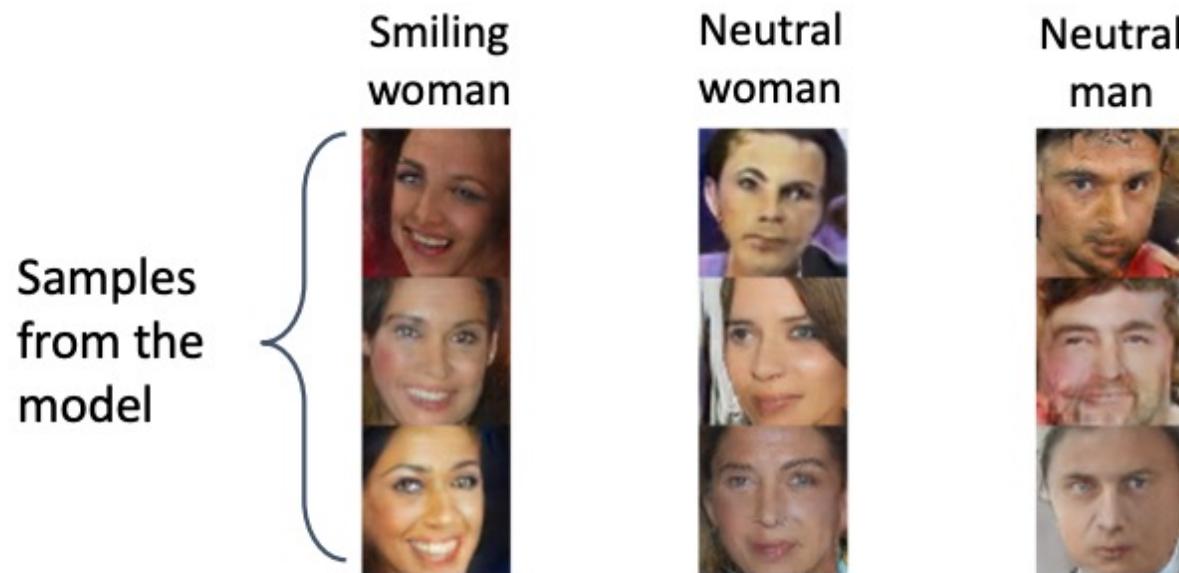
<https://github.com/NVlabs/stylegan2-ada>

```
python generate.py --outdir=out --dlatents=out/dlatents.npz --network=https://nvlabs-fi-cdn.nvidia.com/stylegan2-ada/pretrained/ffhq.pkl
```

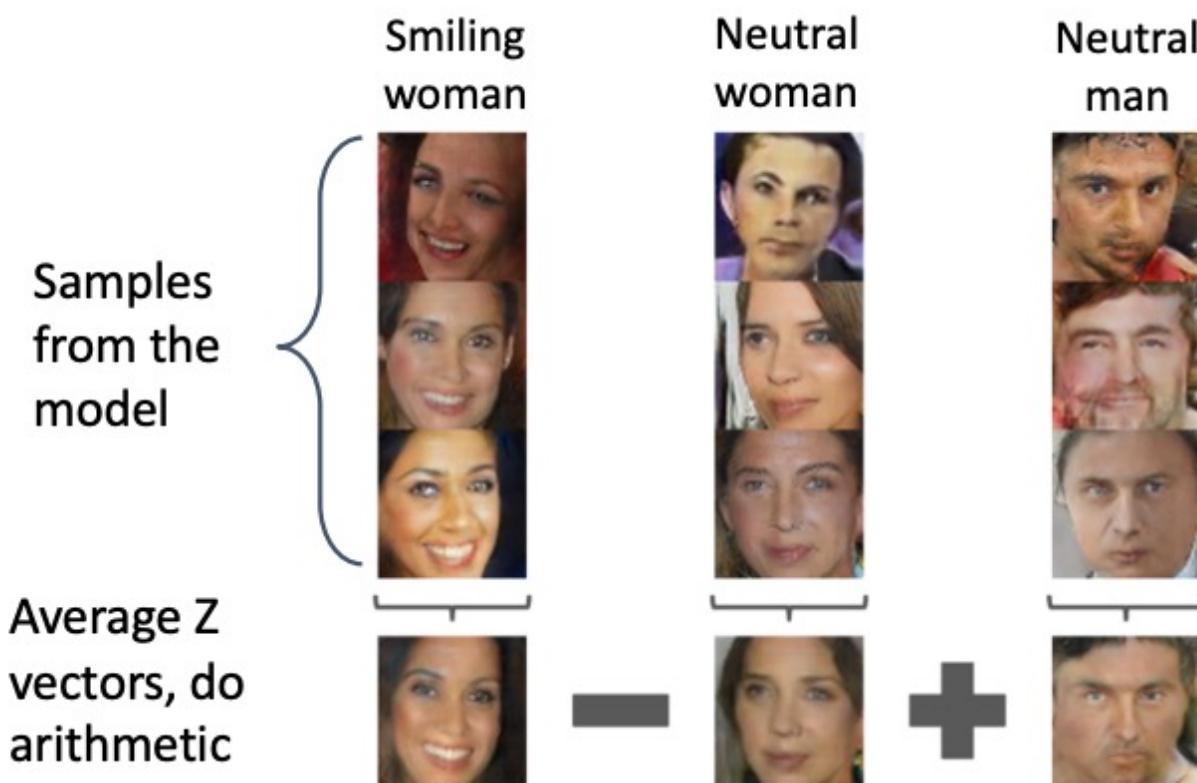
Part III: GAN editing

- Interpolation in latent space
- **GAN Math**
- GAN inversion
- Learn and apply latent directions
- CLIP + StyleGAN

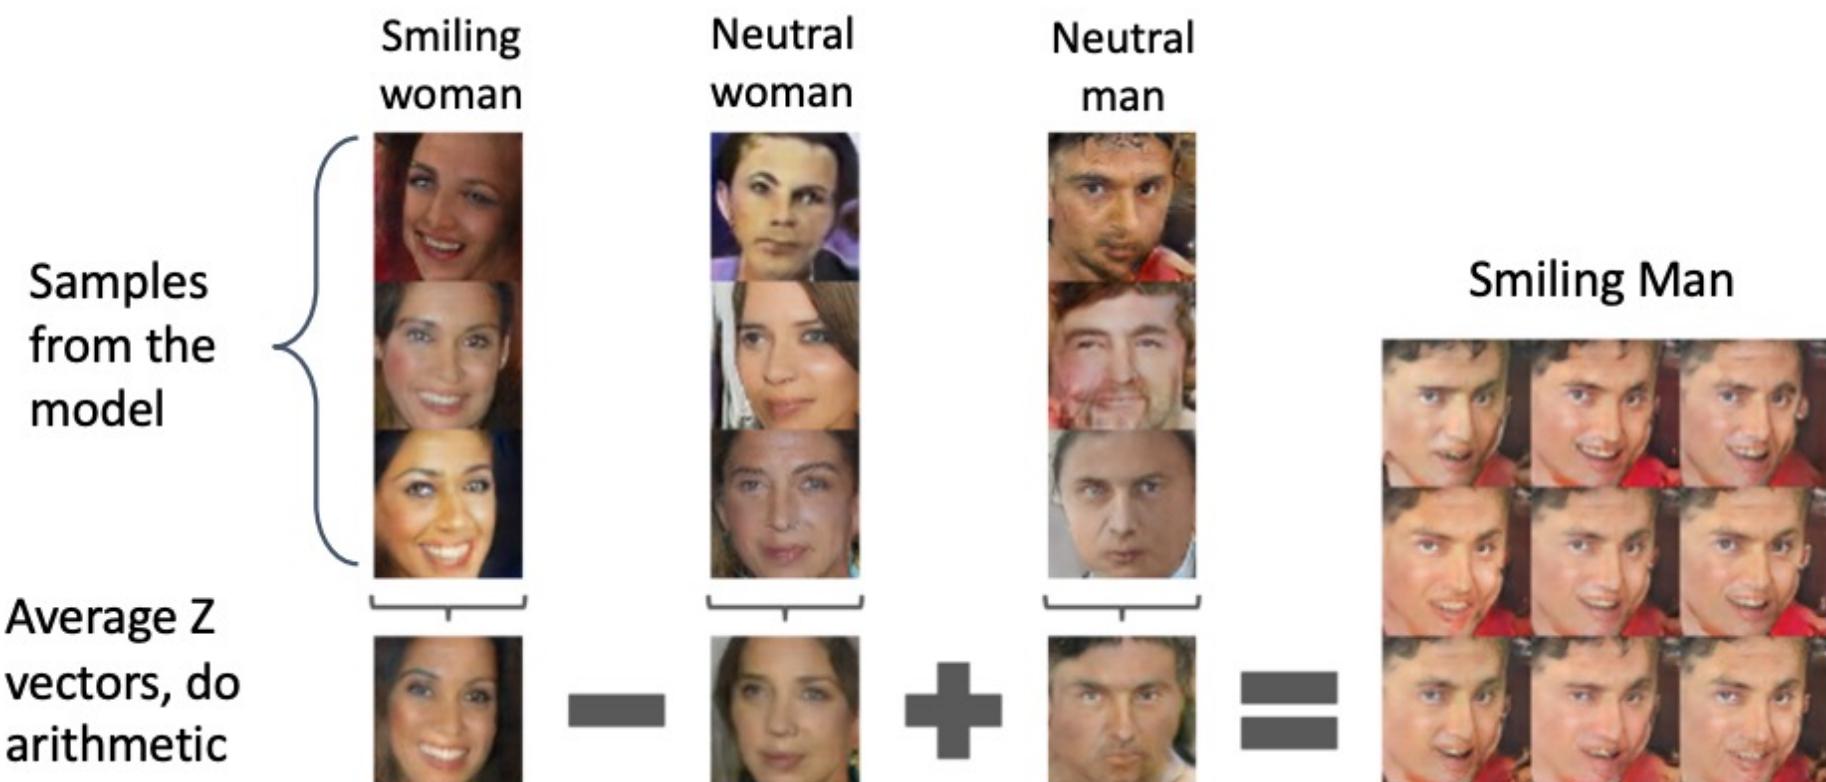
GAN Vector Math



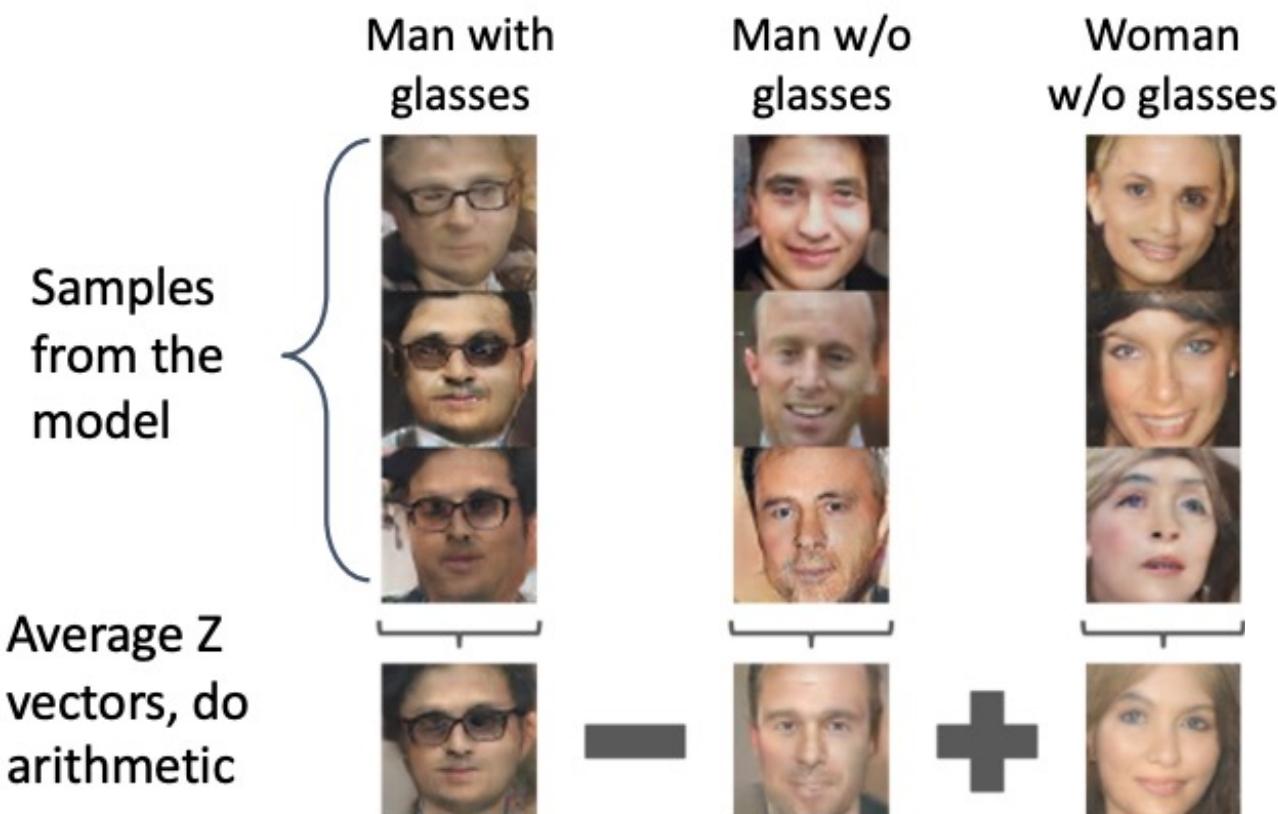
GAN Vector Math



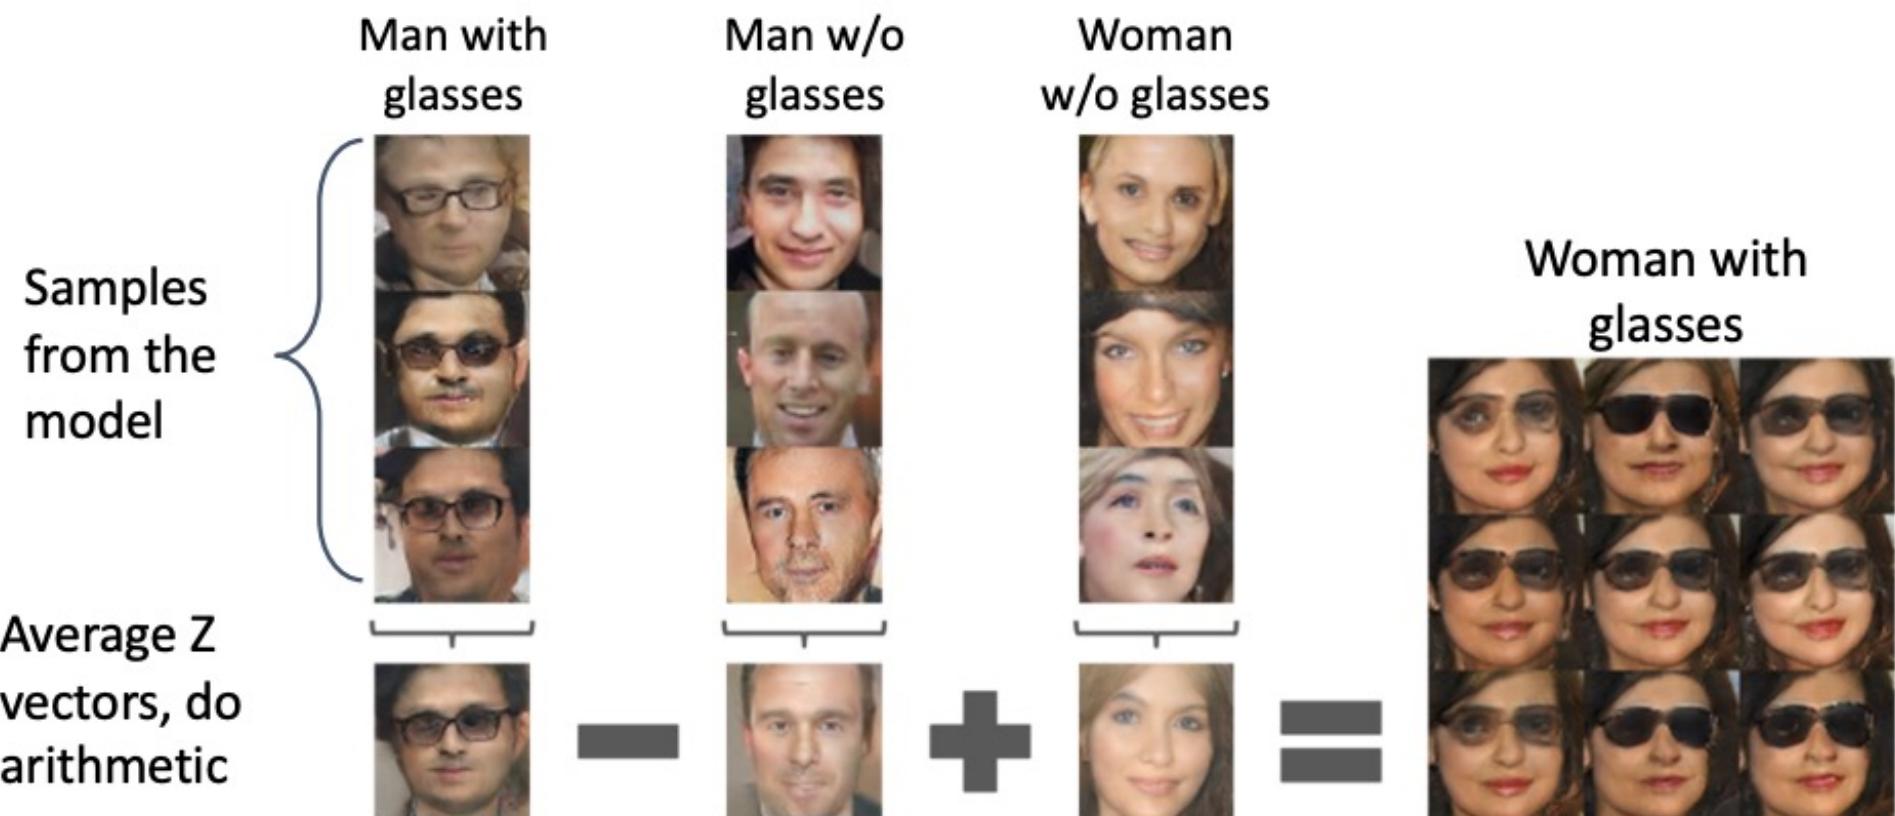
GAN Vector Math



GAN Vector Math



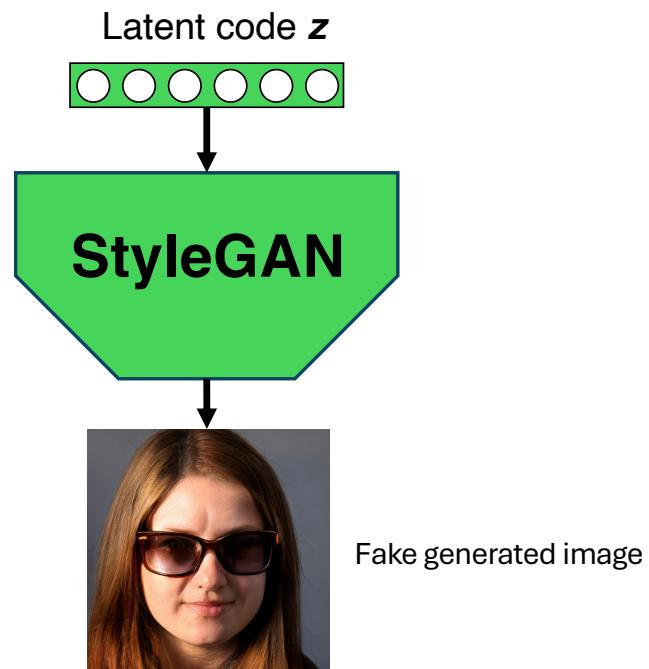
GAN Vector Math



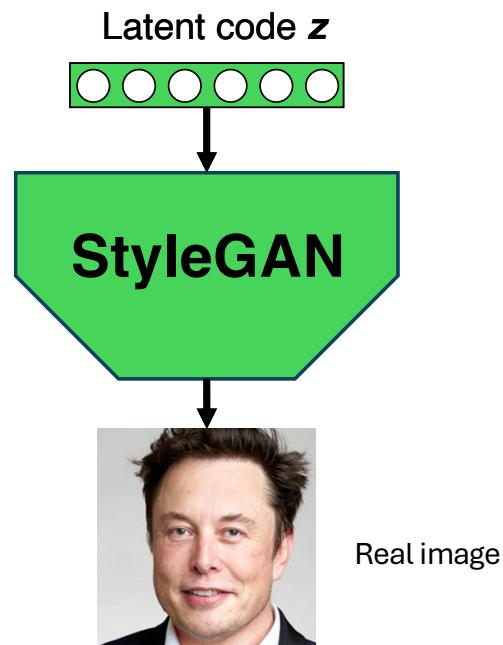
Part III: GAN editing

- Interpolation in latent space
- GAN Math
- **GAN inversion**
- Learn and apply latent directions
- CLIP + StyleGAN

Real images: GAN inversion

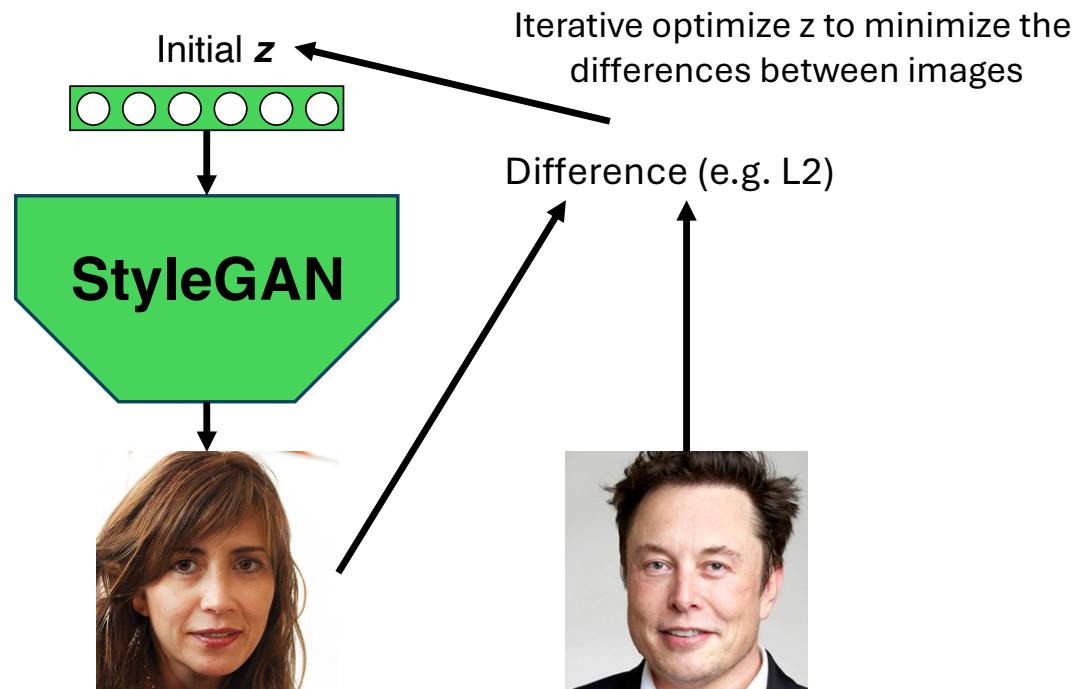


Real images: GAN inversion

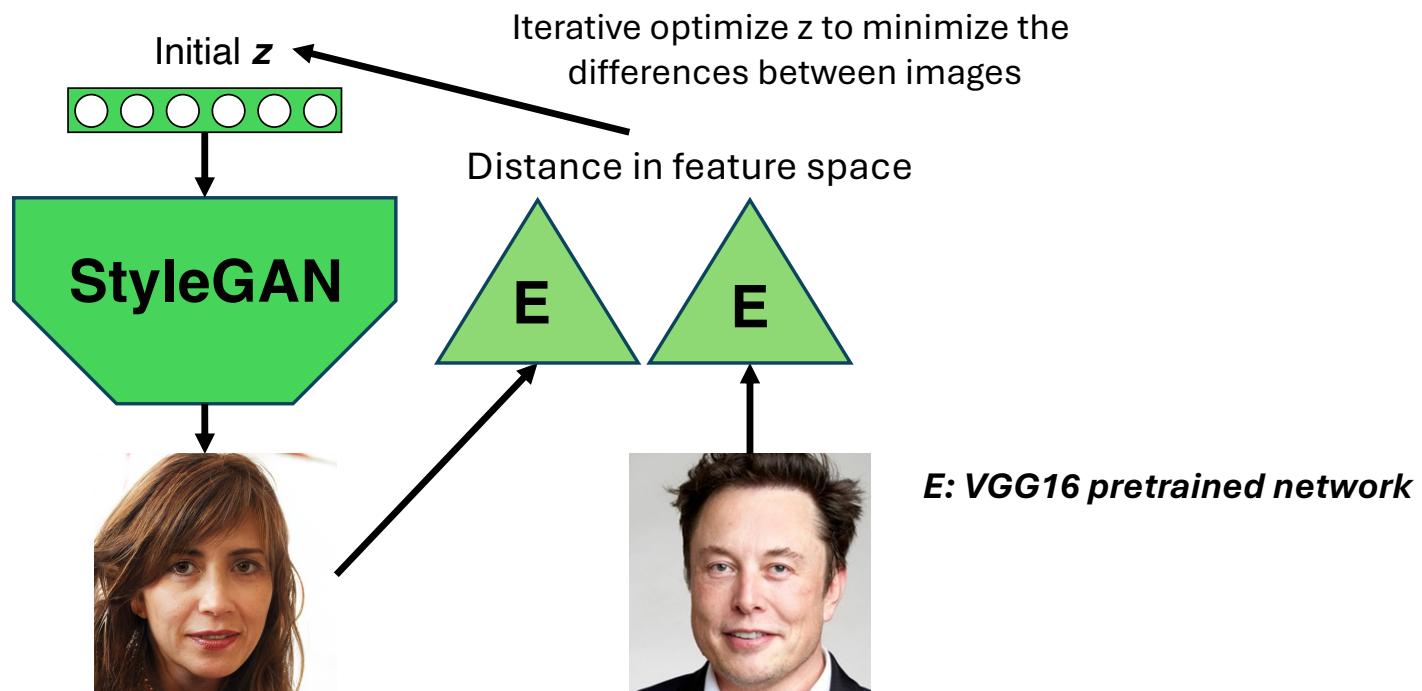


How to find z_o that generates Elon Musk?

Real images: GAN inversion



Real images: GAN inversion



Real images: GAN inversion



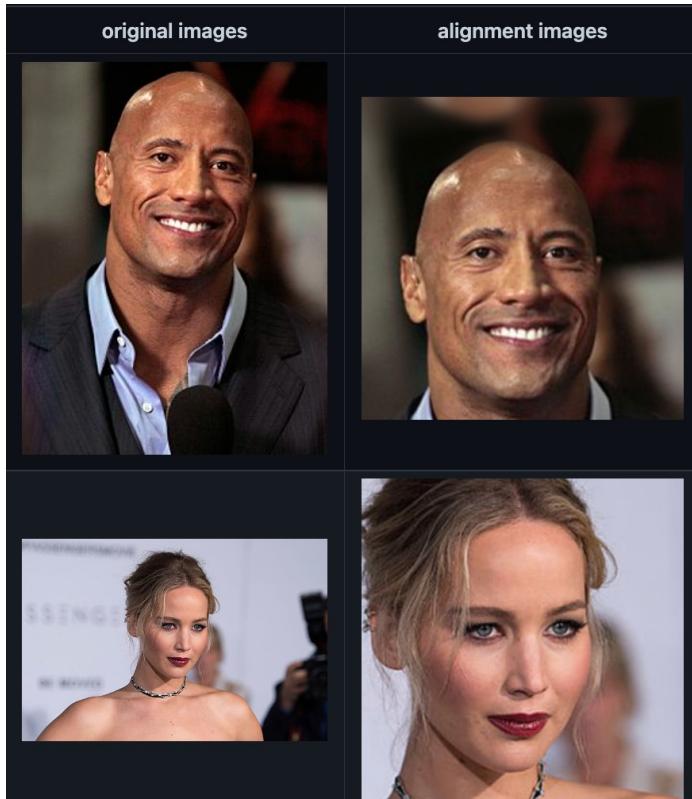
```
python projector.py --outdir=out --target=targetimg.png \ --  
network=https://nvlabs-fi-cdn.nvidia.com/stylegan2-ada/pretrained/ffhq.pkl
```

Real images: GAN inversion



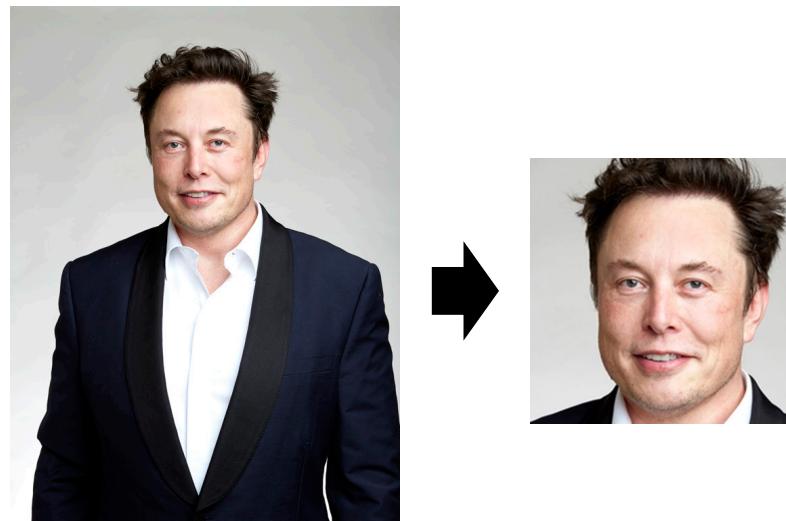
```
python projector.py --outdir=out --target=targetimg.png \ --  
network=https://nvlabs-fi-cdn.nvidia.com/stylegan2-ada/pretrained/ffhq.pkl
```

Align Images with FFHQ



<https://github.com/happy-jihye/FFHQ-Alignment>

Align Images with FFHQ

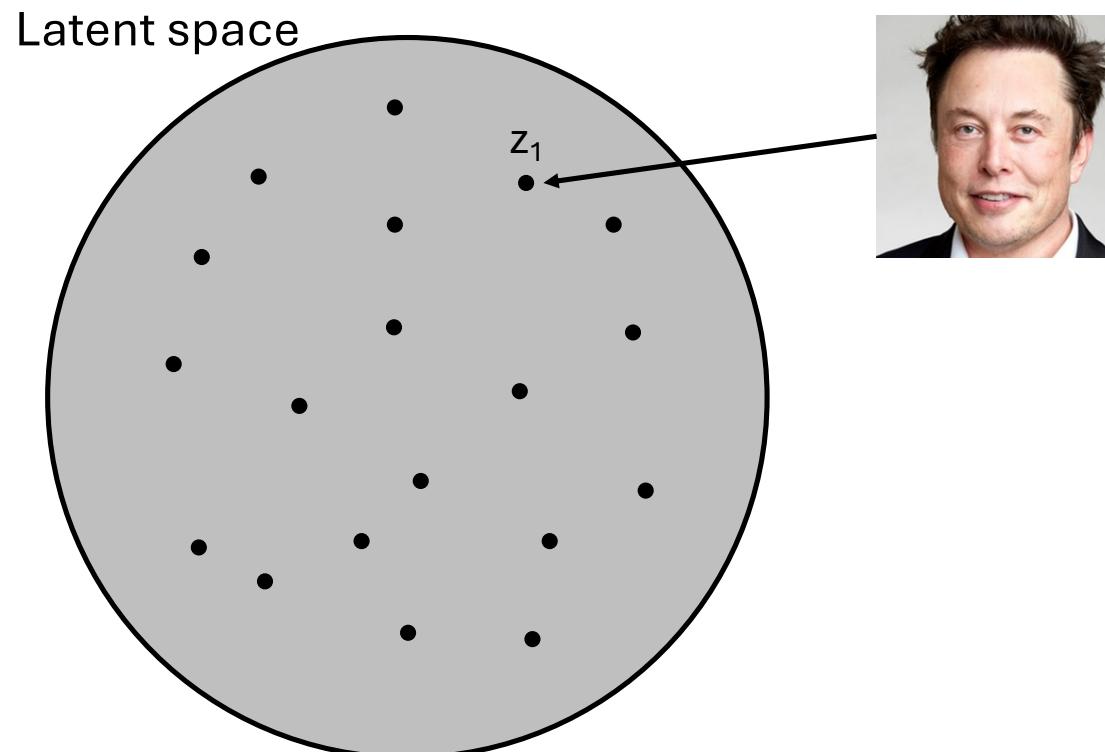


<https://github.com/happy-jihye/FFHQ-Alignment>

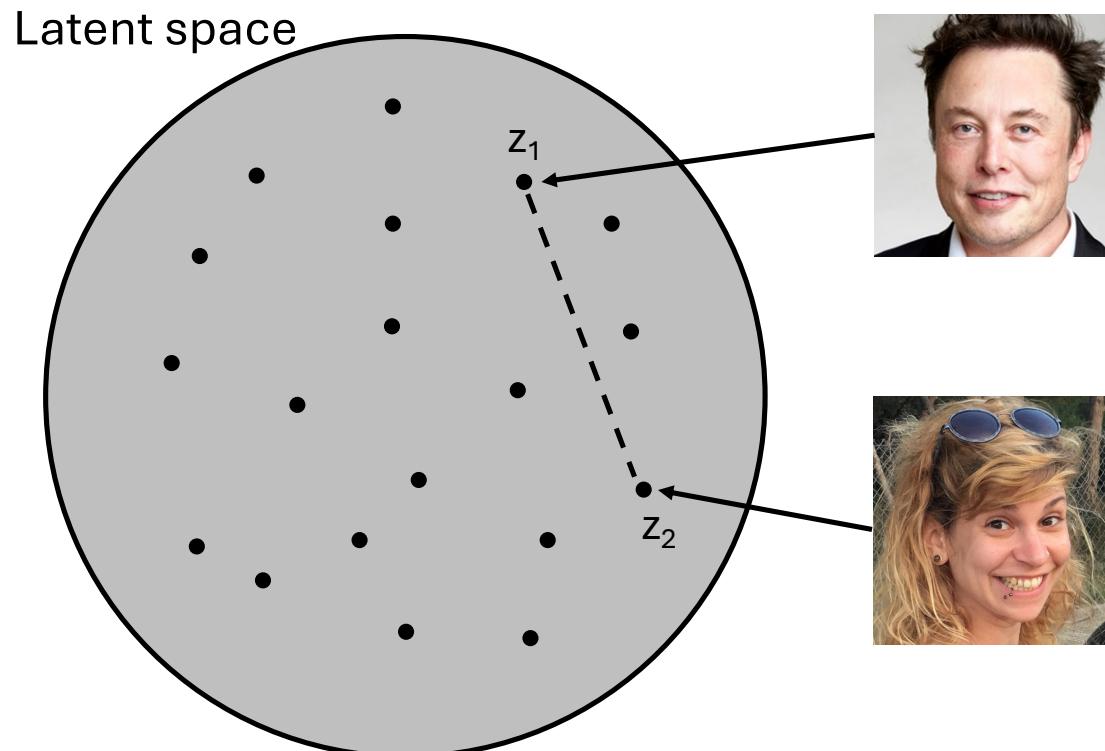
Part III: GAN editing

- Interpolation in latent space
- GAN Math
- GAN inversion
- **Learn and apply latent directions**
- CLIP + StyleGAN

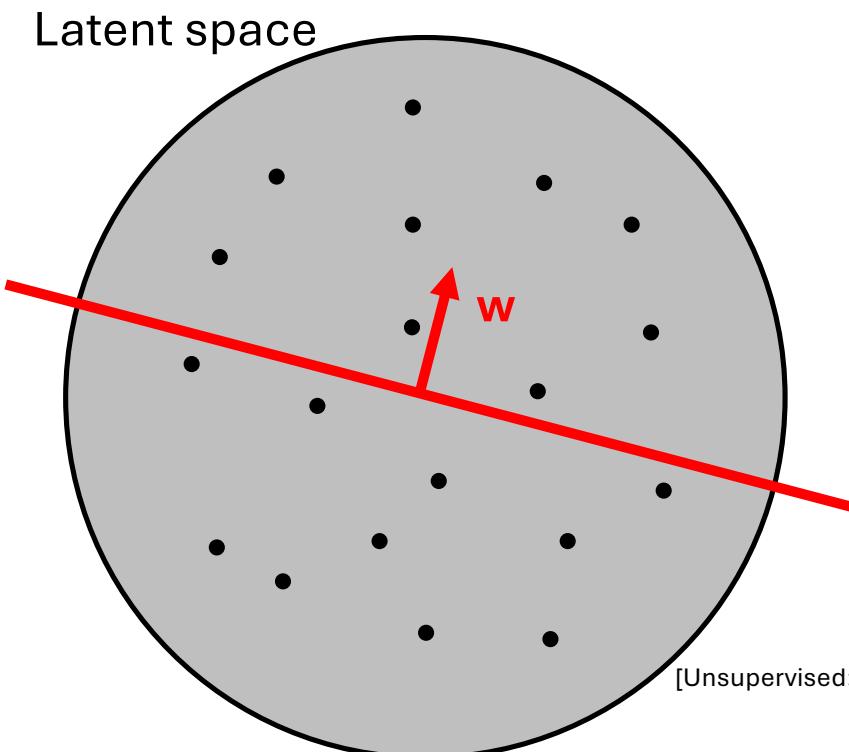
Latent directions



Latent directions



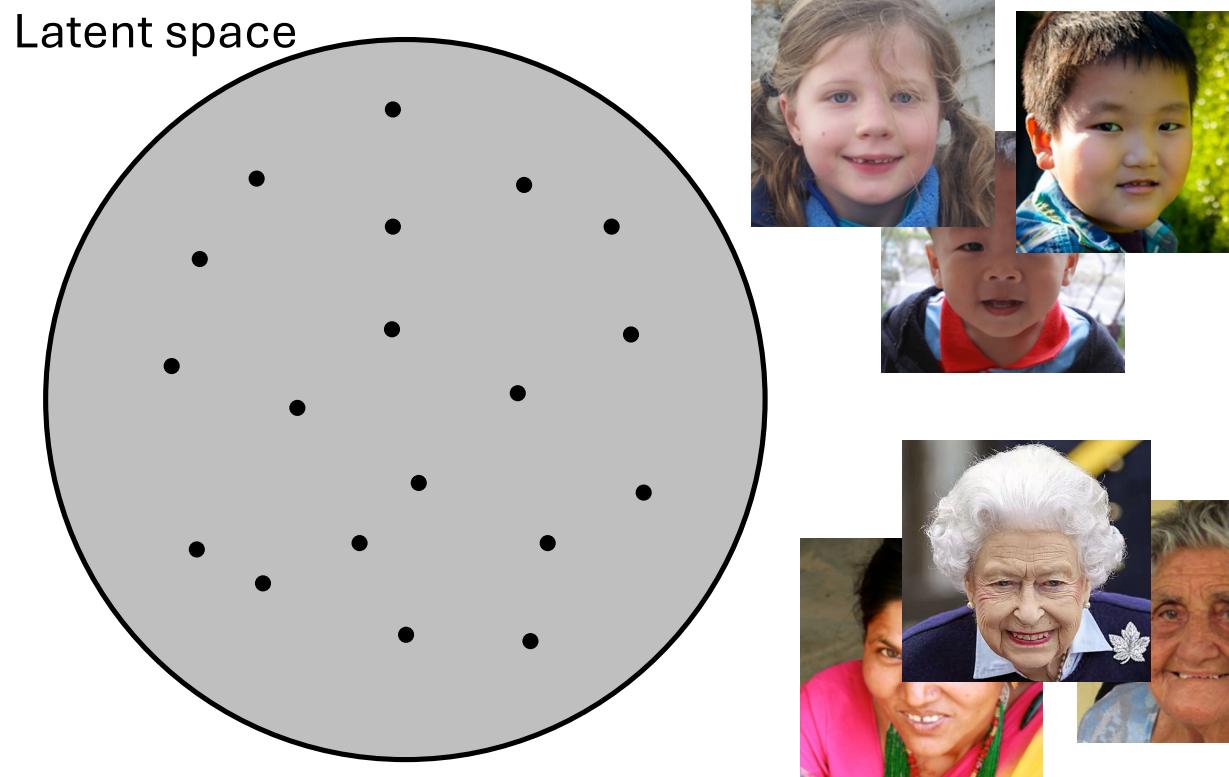
Latent directions



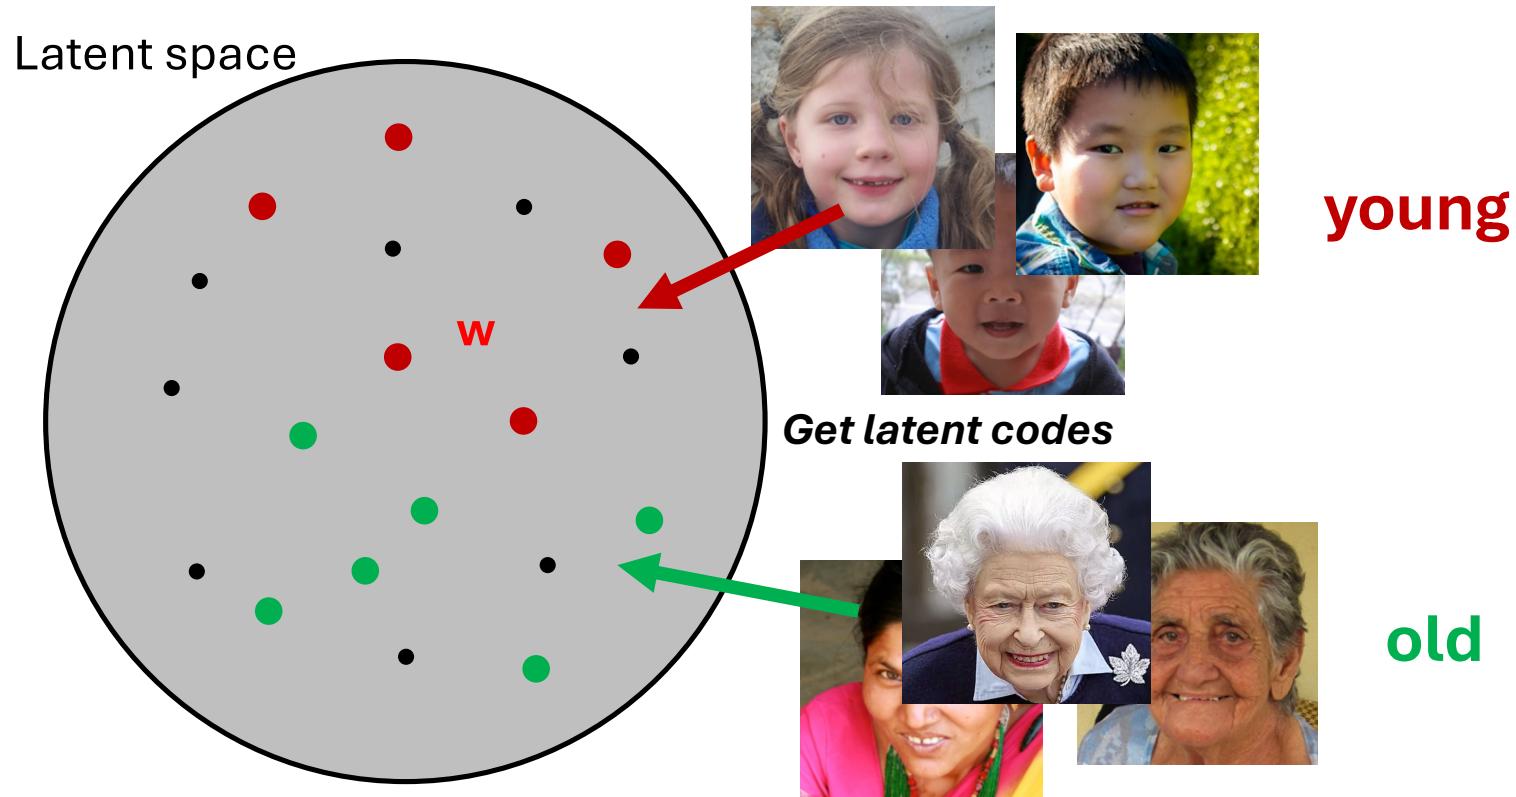
- Learn meaningful latent directions
- Either supervised or unsupervised
- Apply these directions in any image
- Example of directions in faces:
 - Aging
 - Smiling
 - Hair or skin color
 - Gender

[Unsupervised: Voynov and Babenko. "Unsupervised discovery of interpretable directions in the gan latent space." In ICML 2020]

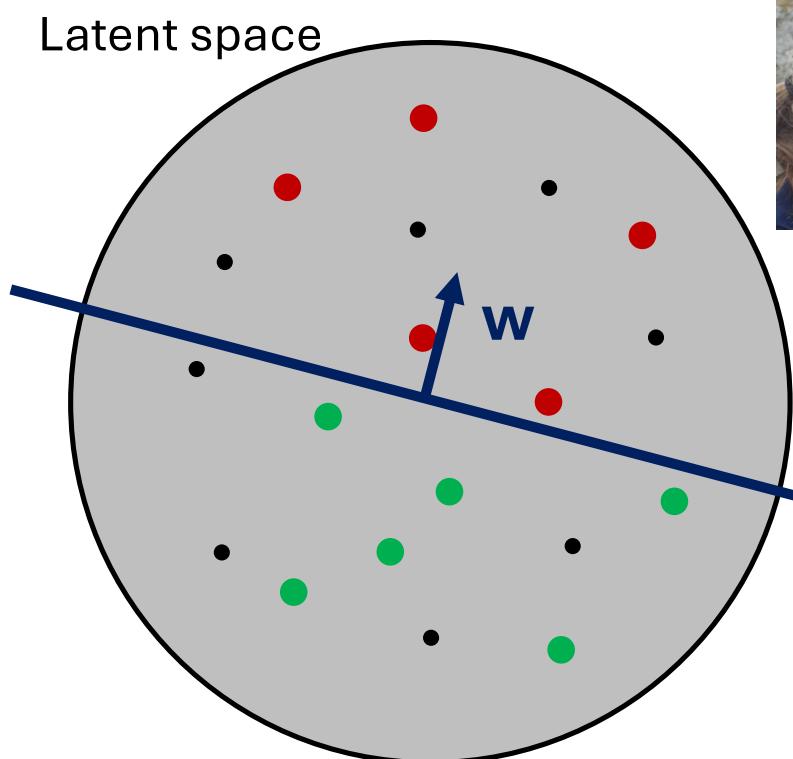
Supervised Learning of Latent directions



Supervised Learning of Latent directions



Supervised Learning of Latent directions

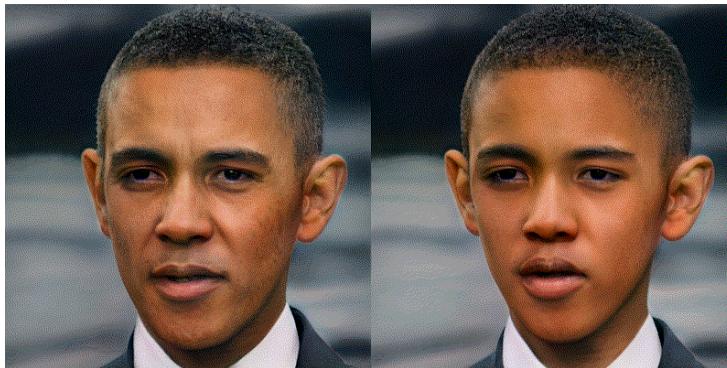


Train a linear classifier on the latent codes

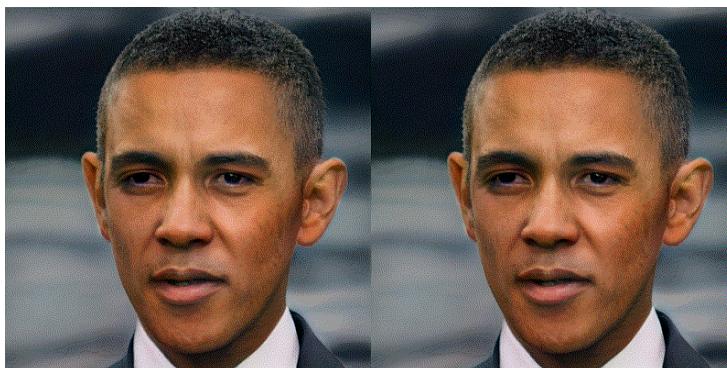


Supervised Learning of Latent directions

aging

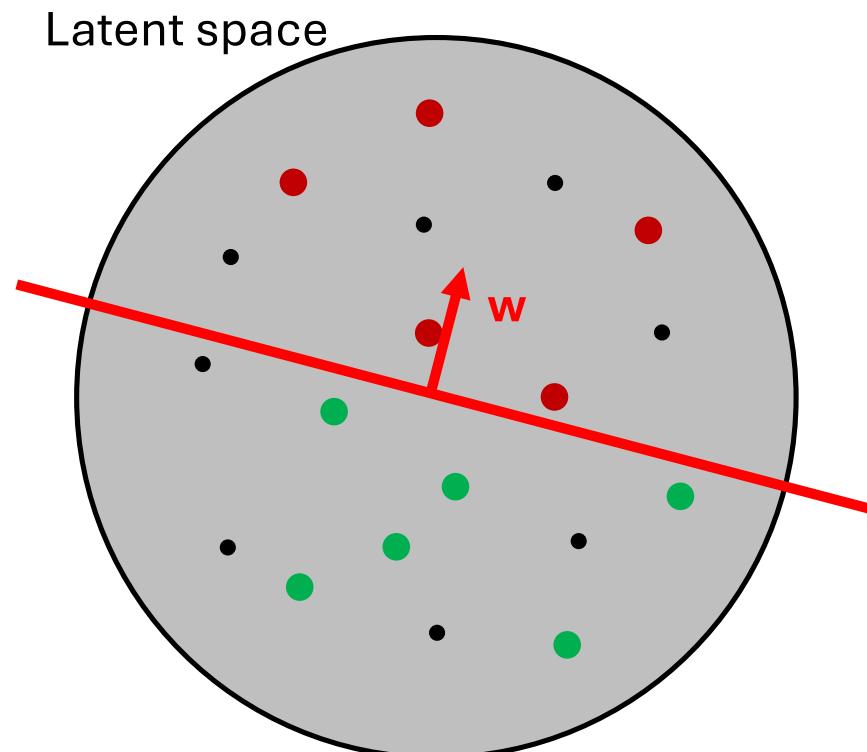


smiling



$$z_{\text{new}} = z + w^*m$$

Supervised Learning of Latent directions



$$z_{\text{new}} = z + w^* m$$

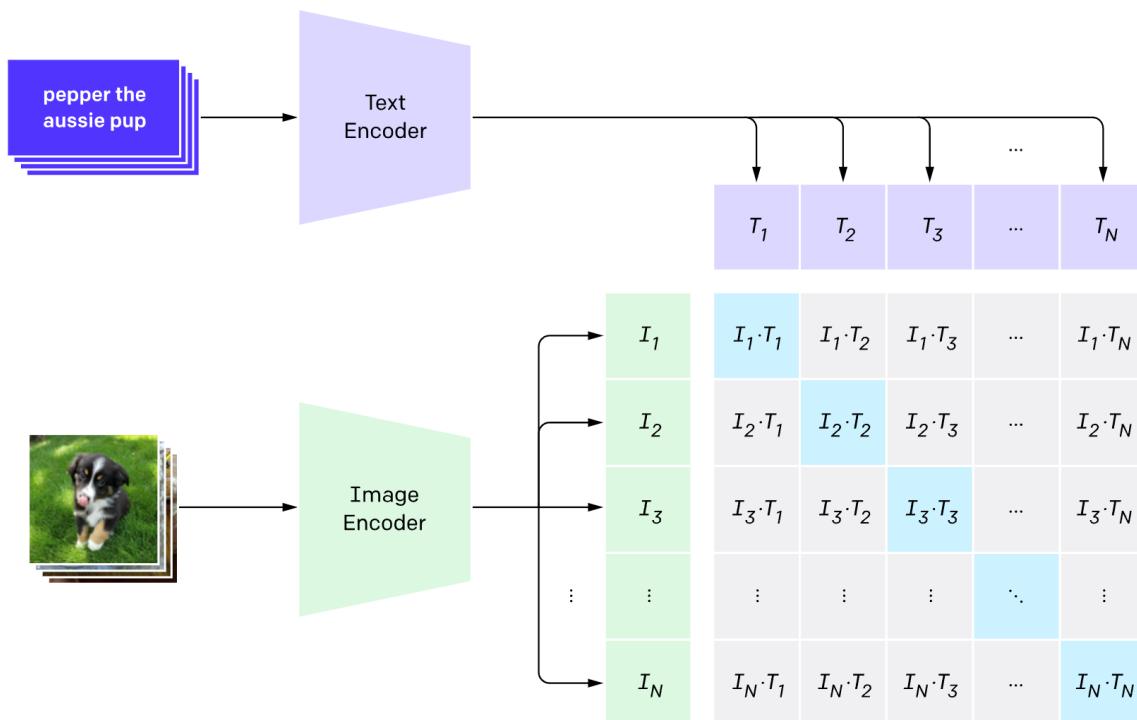
m:magnitude

Part III: GAN editing

- Interpolation in latent space
- GAN Math
- GAN inversion
- Learn and apply latent directions
- **CLIP + StyleGAN**

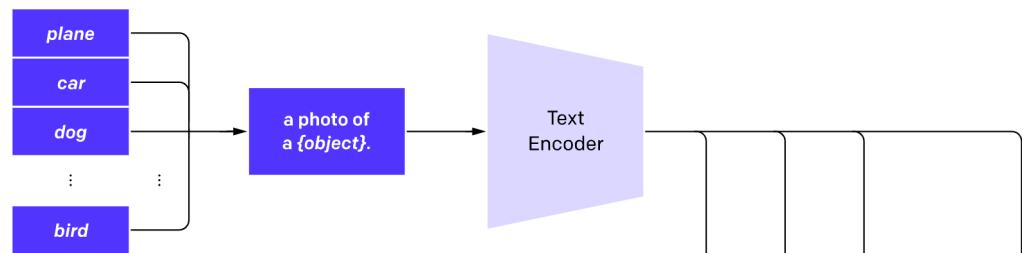
CLIP

1. Contrastive pre-training



CLIP

2. Create dataset classifier from label text



3. Use for zero-shot prediction

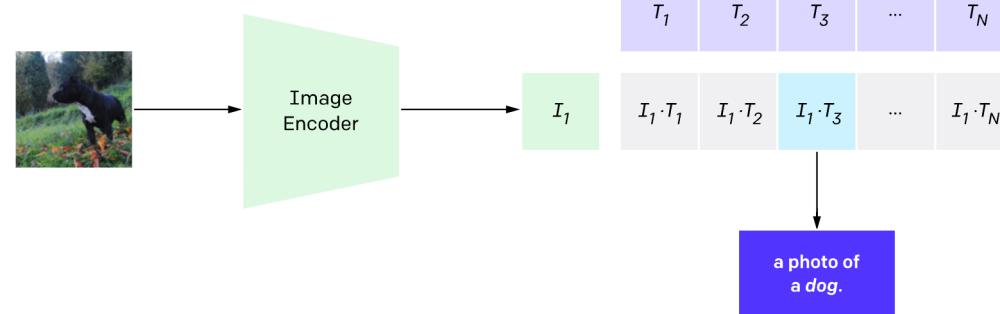


Image generation from text

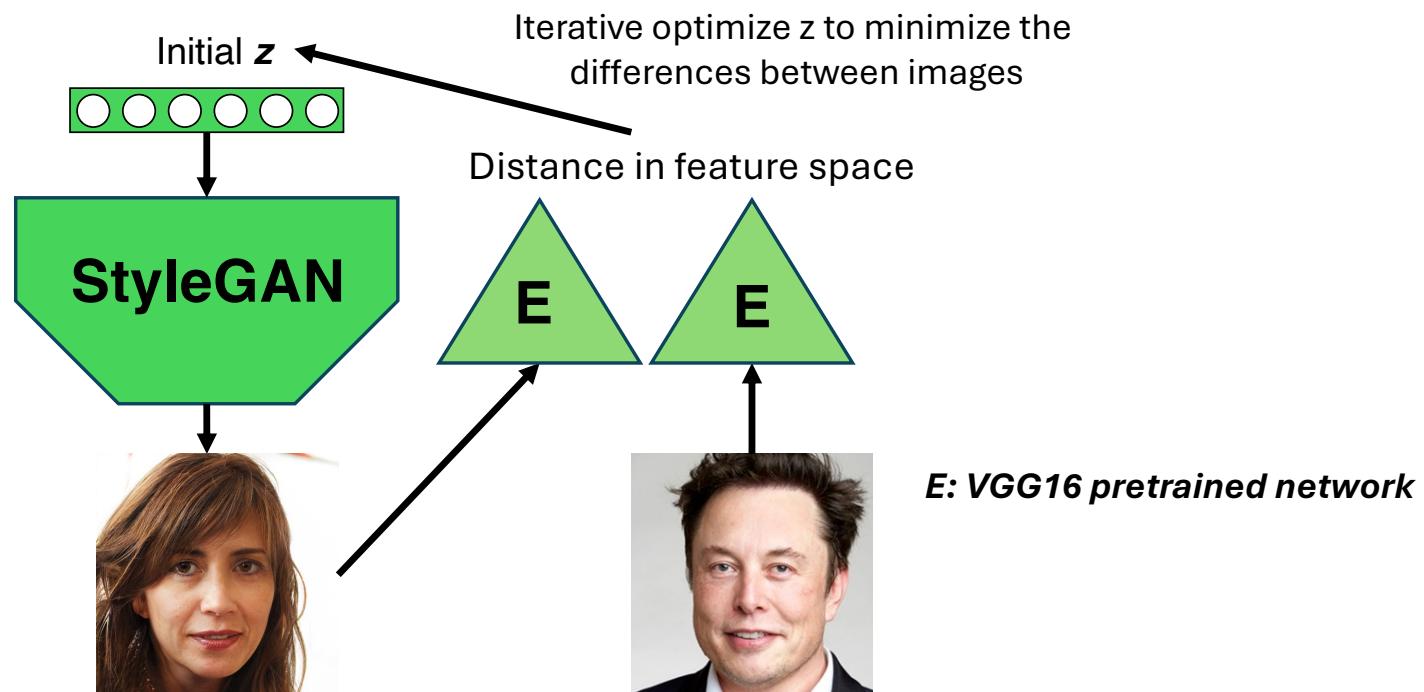


Image generation from text

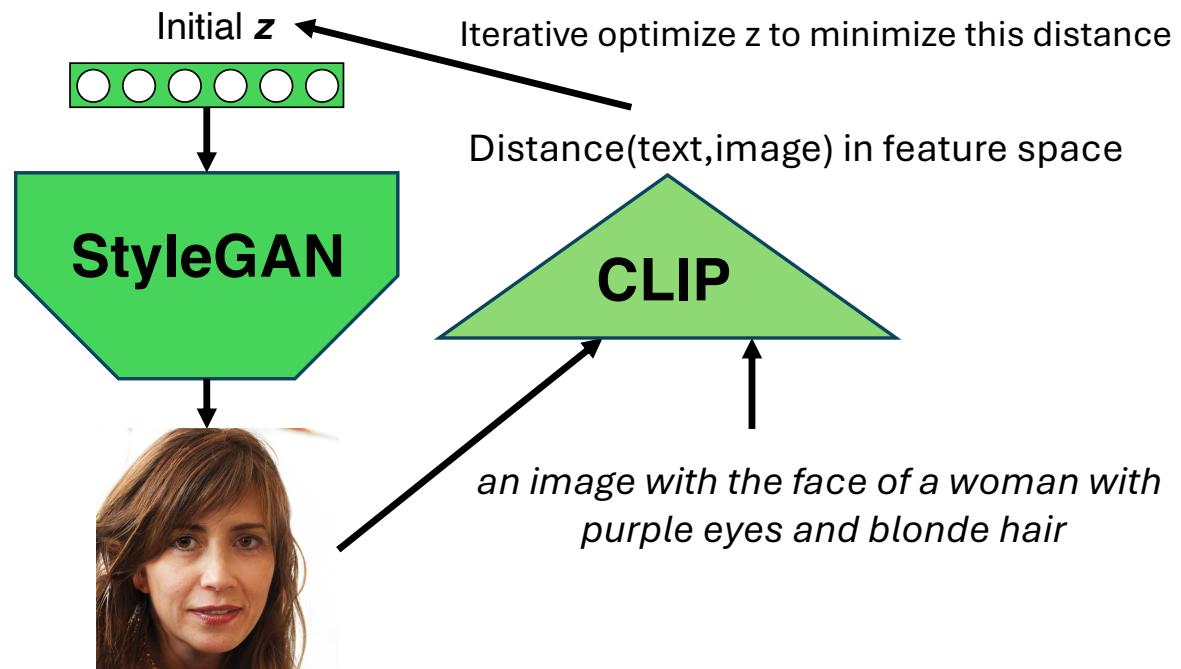


Image generation from text

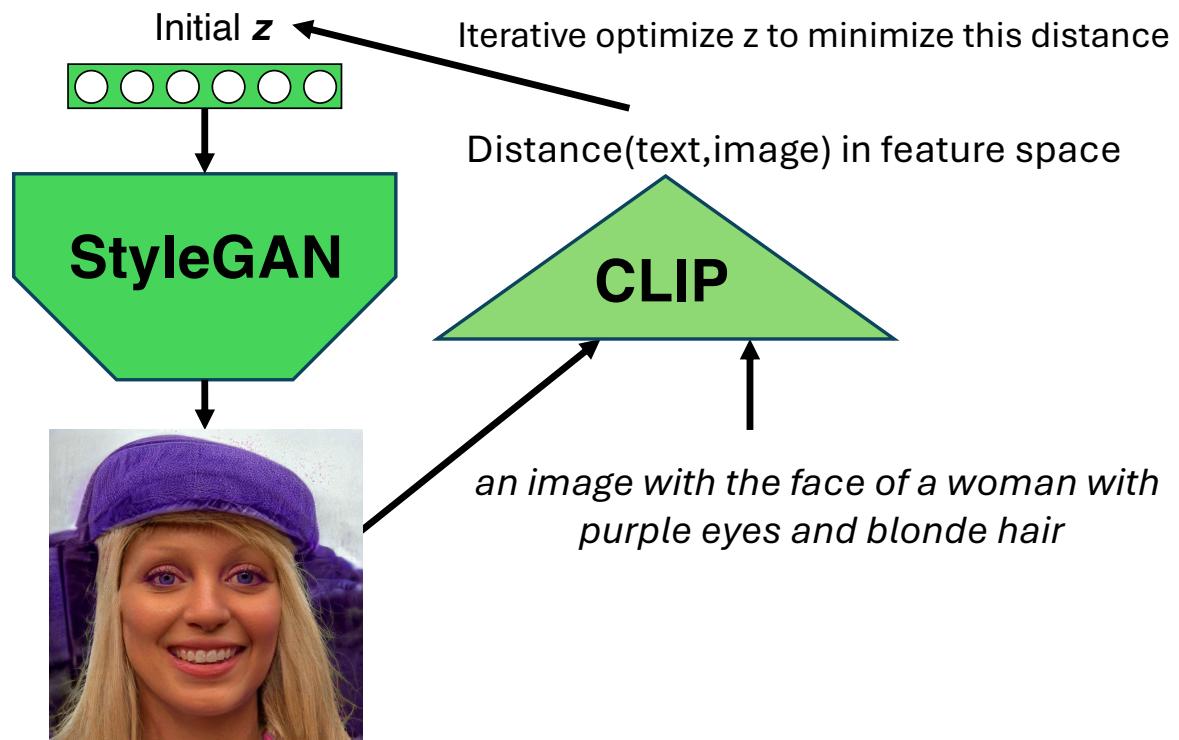
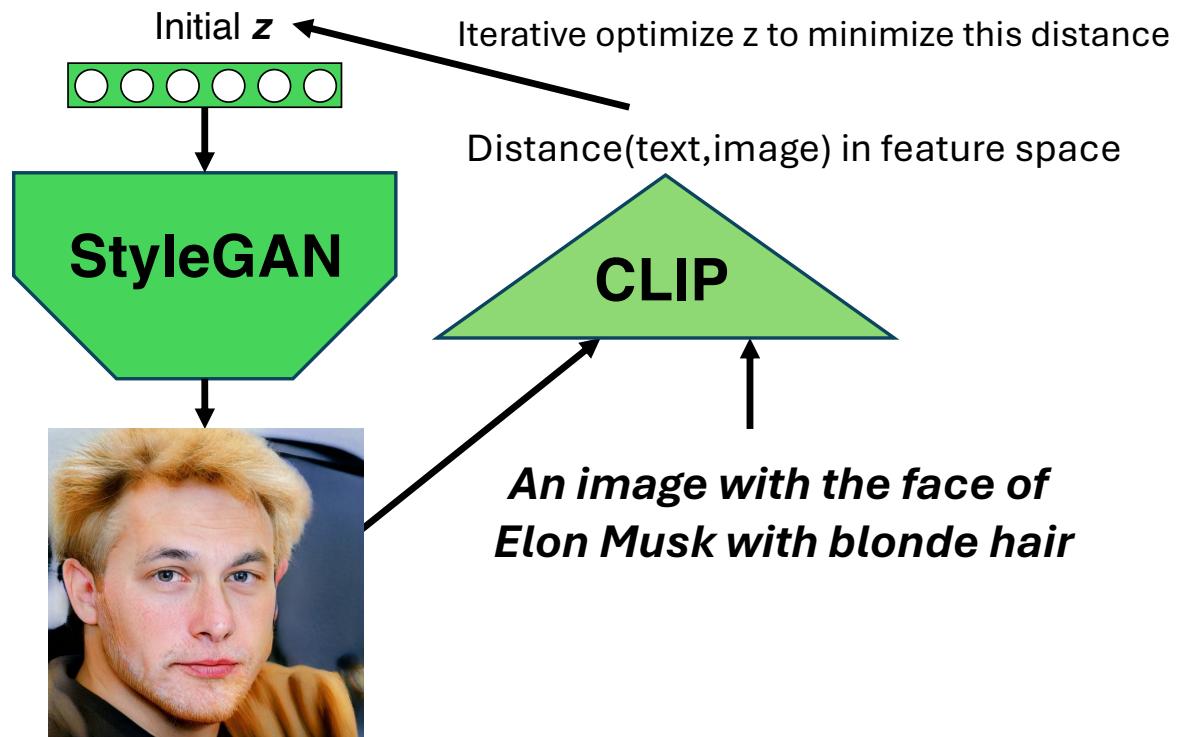


Image generation from text





Thank you